

CS614: Advanced Compilers

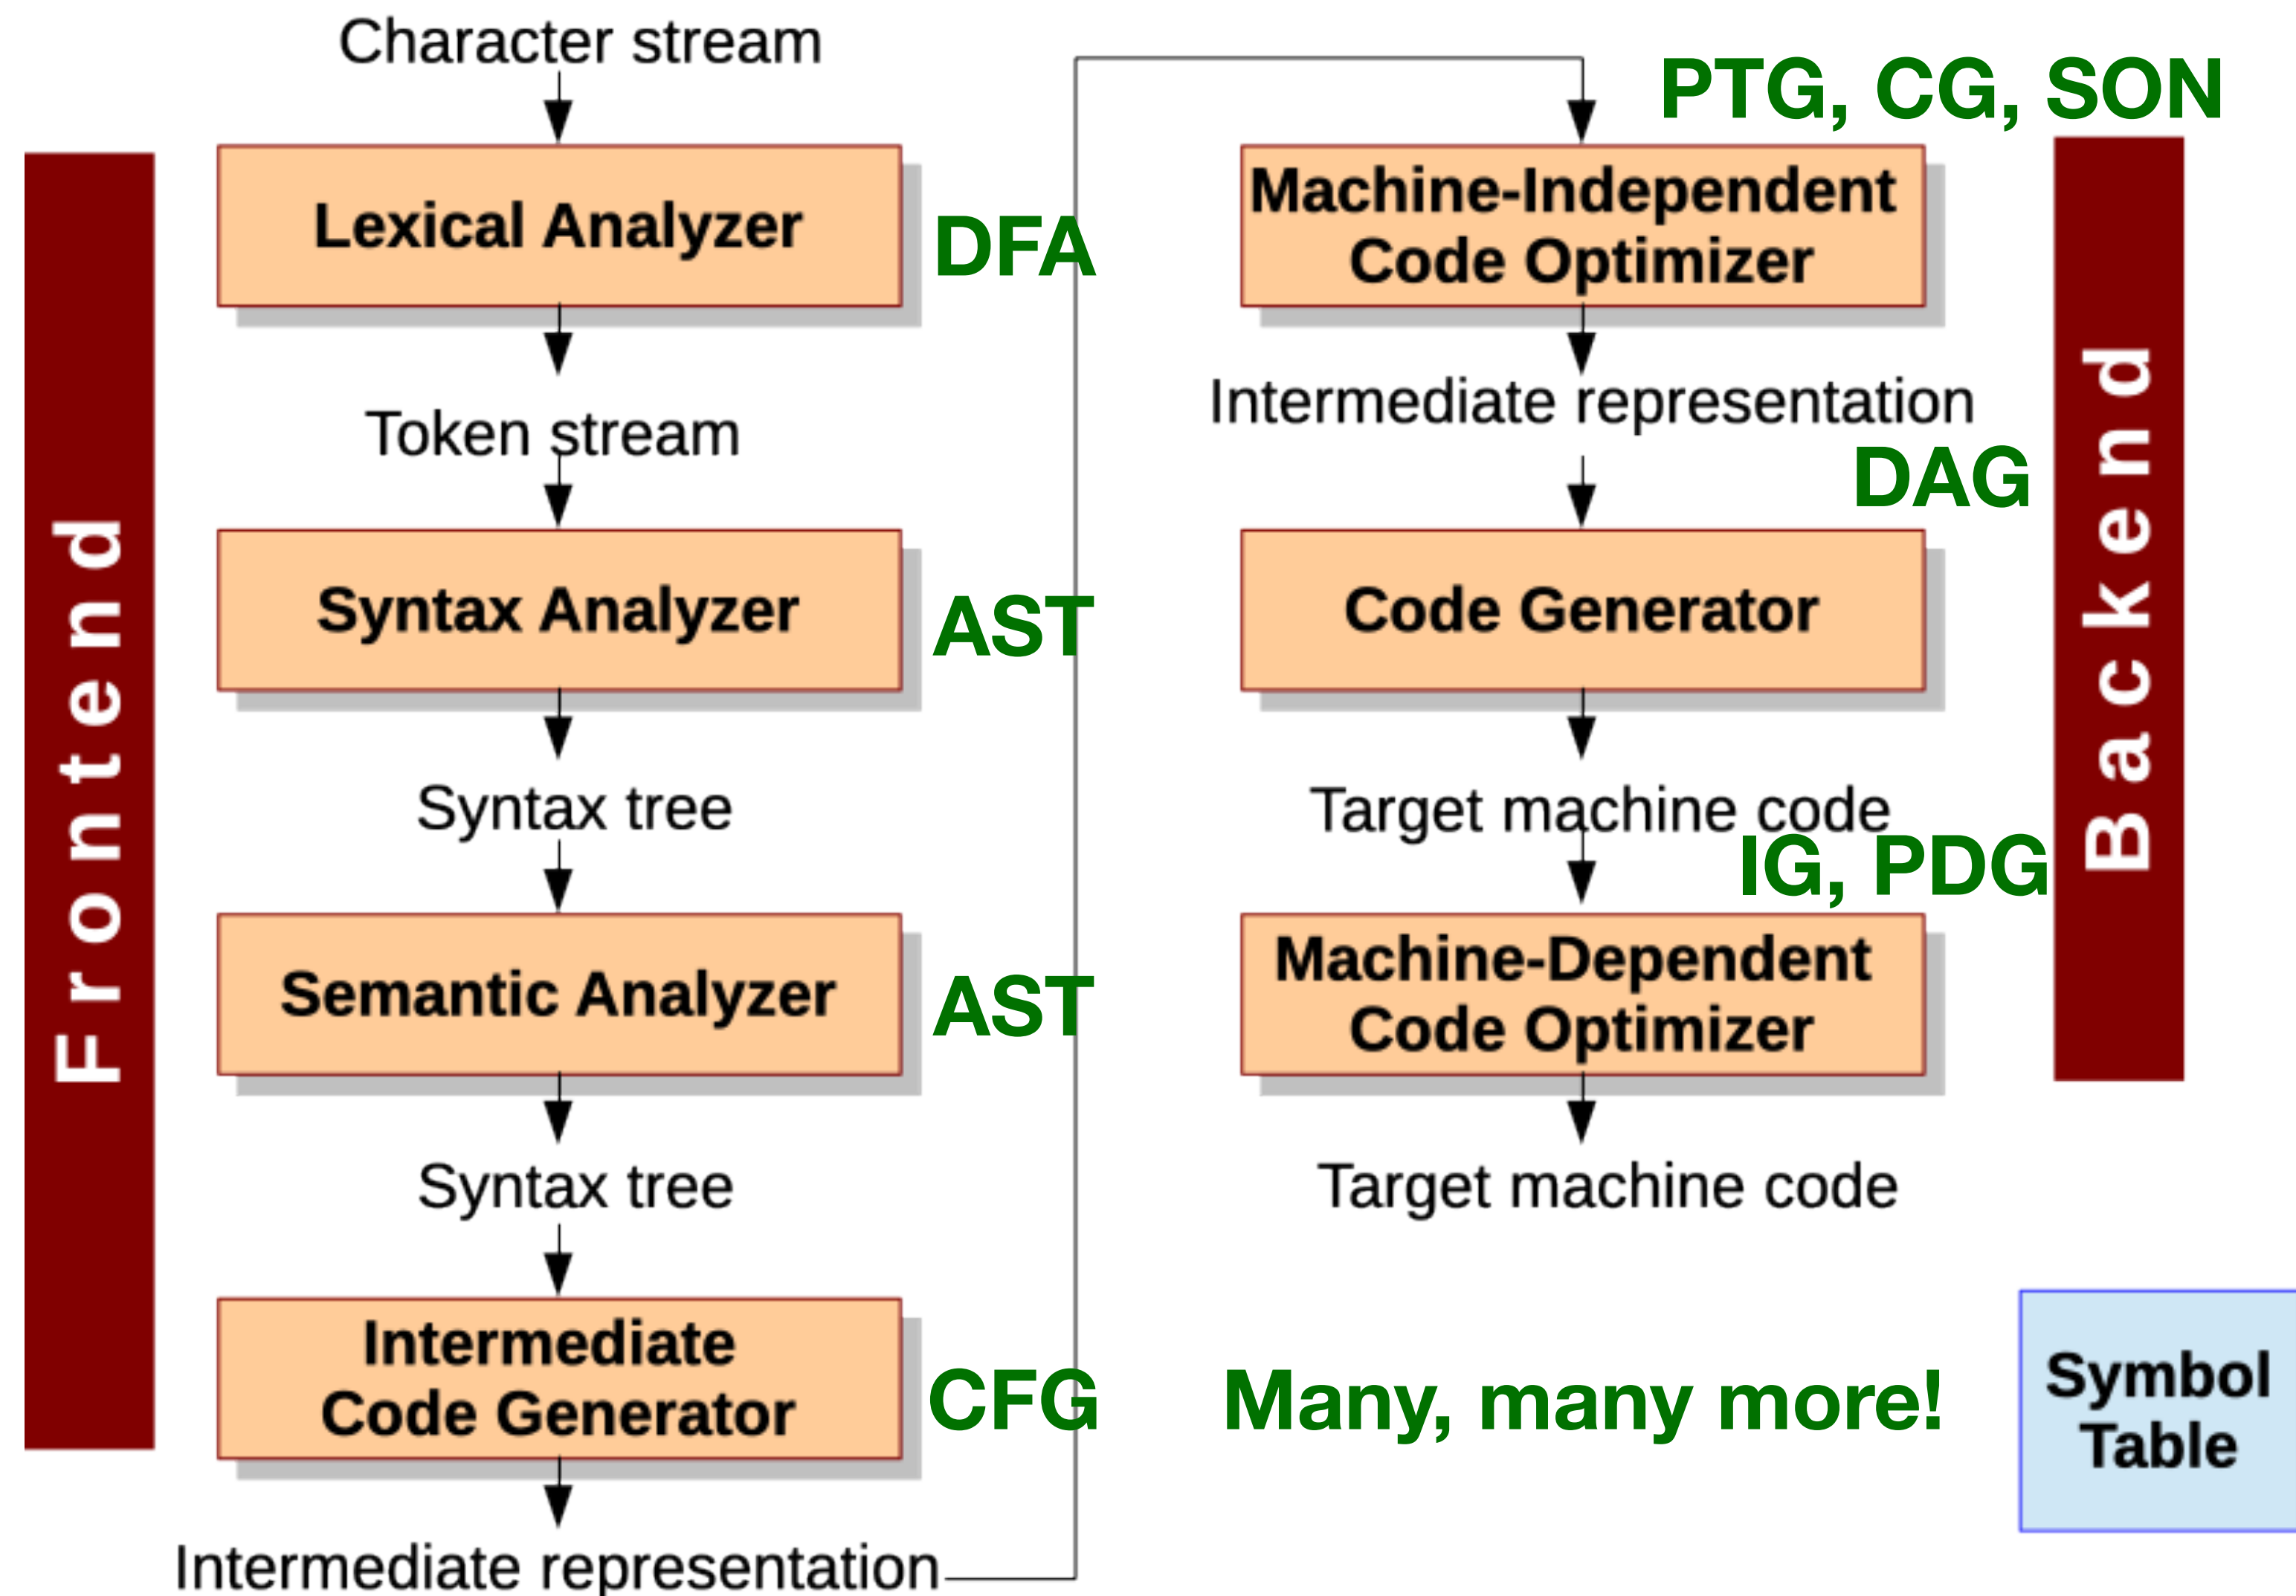
Intro to Program Optimization

Manas Thakur
CSE, IIT Bombay



Spring 2025

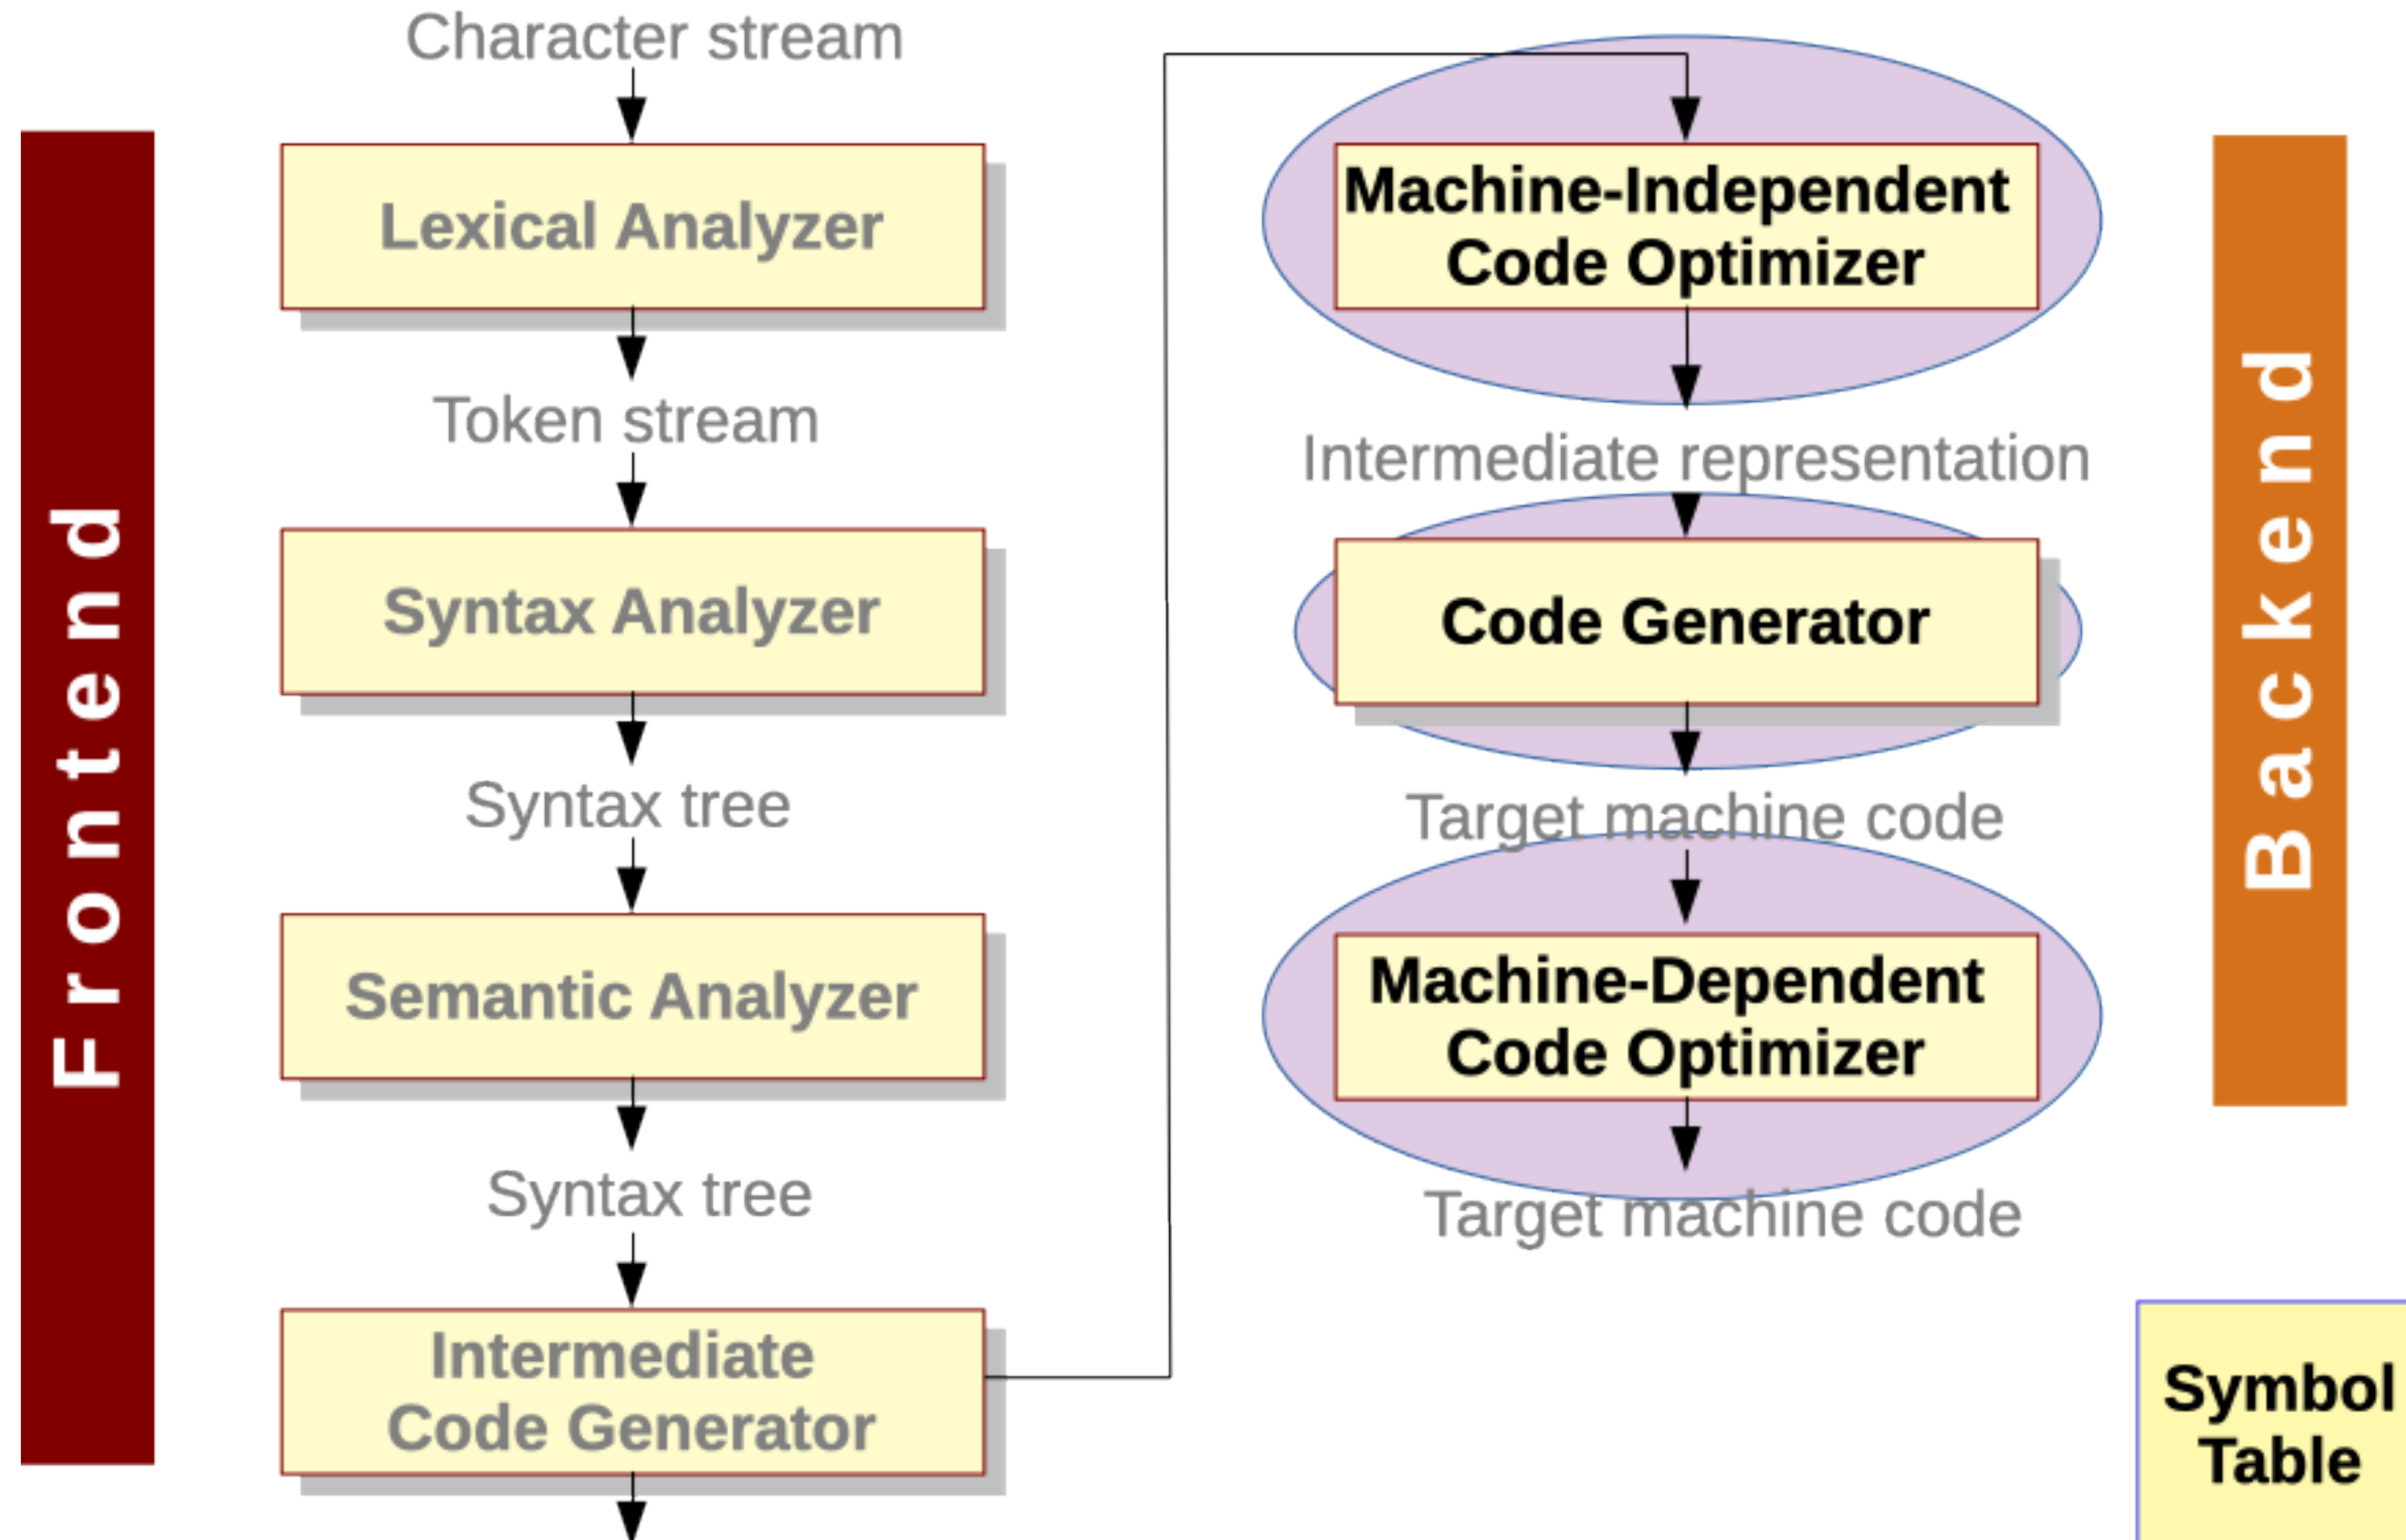
Graphs in Compilers



We will miss you profoundly...



Now we are moving to the back-end



Things we have already learnt

➤ Lowering

- From language-level constructs to simple limited number of constructs
- At this point we could generate machine code!
 - Map from lower-level IR to actual ISA
 - Maybe some register management
 - Maybe some instruction selection
 - Pass on to assembler
 - Why not generate machine code directly?
 - *Kucch to raham karo janaab!*



But first...

- The compiler “understands” the program
 - IR captures program semantics
 - Lowering: semantics-preserving transformation
 - Why not do others?
- Compiler optimizations
 - Oh great, now my code will be optimal!
 - Sorry, it’s a misnomer
 - What is an “optimization”?



Code Optimization: **An Intro**

- **Goal:** Generate optimized code
- **Metrics:**
 - **Code size**
 - Memory requirements
 - Say opcodes take 1 byte, each operand another byte
 - **Number of registers**
 - Along with constraints on their usage
 - **Estimated cost**
 - How fast is the code?
 - Say instructions with single operand take 1 cycle, with two operands take 2 cycles, and those involving memory take 4 cycles
 - **Sometimes**
 - power, energy, platform (in)dependence, ...



Code Optimization: **An Intro**

A simple one-to-one mapping:

```
1: a = 0
2: b = a + 1
3: c = c + b
4: a = b * 2
```



```
MOV R1 0
STA a R1
```

```
MOV R1 1
LDA R2 a
ADD R2 R1
STA b R2
```

```
LDA R1 b
LDA R2 c
ADD R2 R1
STA c R2
```

```
MOV R1 2
LDA R2 b
MUL R2 R1
STA a R2
```

Cost:

Registers: 2

Space: 42 bytes

Time: 44 cycles

Can we do better?

Code Optimization: An Intro

Better register usage:

```
MOV R1 0
STA a R1
```

```
MOV R1 1
LDA R2 a
ADD R2 R1
STA b R2
```

```
LDA R1 b
LDA R2 c
ADD R2 R1
STA c R2
```

```
MOV R1 2
LDA R2 b
MUL R2 R1
STA a R2
```

Cost:

Registers: 2
Space: 42 bytes
Time: 44 cycles

```
MOV R1 0
STA a R1
```

```
MOV R2 1
ADD R1 R2
STA b R1
```

```
LDA R2 c
ADD R2 R1
STA c R2
```

```
MOV R2 2
MUL R1 R2
STA a R1
```

Cost:

Registers: 2
Space: 33 bytes
Time: 32 cycles

```
1: a = 0
2: b = a + 1
3: c = c + b
4: a = b * 2
```

Can we do better?



Code Optimization: **An Intro**

Remove **redundant** store to a:

```
MOV R1 0
STA a R1
```

```
MOV R2 1
ADD R1 R2
STA b R1
```

```
LDA R2 c
ADD R2 R1
STA c R2
```

Cost:
Registers: 2
Space: 33 bytes
Time: 32 cycles

```
MOV R2 2
MUL R1 R2
STA a R1
```



```
MOV R1 0
```

```
MOV R2 1
ADD R1 R2
STA b R1
```

```
LDA R2 c
ADD R2 R1
STA c R2
```

Cost:
Registers: 2
Space: 30 bytes
Time: 28 cycles

```
MOV R2 2
MUL R1 R2
STA a R1
```

```
1: a = 0
2: b = a + 1
3: c = c + b
4: a = b * 2
```

Can we do better?

Code Optimization: An Intro

Select specialized instructions:

```
MOV R1 0
```

```
MOV R2 1  
ADD R1 R2  
STA b R1
```

```
LDA R2 c  
ADD R2 R1  
STA c R2
```

Cost:

Registers: 2
Space: 30 bytes
Time: 28 cycles

```
MOV R2 2  
MUL R1 R2  
STA a R1
```



Cost:

Registers: 2
Space: 21 bytes
Time: 21 cycles

```
CLR R1  
INC R1  
STA b R1  
LDA R2 c  
ADD R2 R1  
STA c R2  
SHL R1  
STA a R1
```

```
1: a = 0  
2: b = a + 1  
3: c = c + b  
4: a = b * 2
```

Can we do better?

Code Optimization: An Intro

Propagate constant 0:

```
1: a = 0
2: b = a + 1
3: c = c + b
4: a = b * 2
```

Cost:

Registers: 2
Space: 21 bytes
Time: 21 cycles

```
CLR R1
INC R1
STA b R1
LDA R2 c
ADD R2 R1
STA c R2
SHL R1
STA a R1
```



```
MOV R1 1
STA b R1
LDA R2 c
ADD R2 R1
STA c R2
SHL R1
STA a R1
```

Cost:

Registers: 2
Space: 20 bytes
Time: 21 cycles

Can we do better?



Code Optimization: **An Intro**

Assuming b is not used in future,
propagate constant 1:

```
1: a = 0
2: b = a + 1
3: c = c + b
4: a = b * 2
```

```
MOV R1 1
STA b R1
LDA R2 c
ADD R2 R1
STA c R2
SHL R1
STA a R1
```

Cost:

Registers: 2
Space: 20 bytes
Time: 21 cycles



```
LDA R1 c
INC R1
STA c R1
MOV R1 2
STA a R1
```

Cost:

Registers: 1
Space: 14 bytes
Time: 15 cycles

Can we do better?

Code Optimization: An Intro

Assuming the availability of a store-immediate instruction:

```
1: a = 0
2: b = a + 1
3: c = c + b
4: a = b * 2
```

```
LDA R1 c
INC R1
STA c R1
MOV R1 2
STA a R1
```

Cost:

Registers: 1
Space: 14 bytes
Time: 15 cycles



```
LDA R1 c
INC R1
STA c R1
STI a 2
```

Cost:

Registers: 1
Space: 11 bytes
Time: 13 cycles

Assuming no other special instruction and no new knowledge about past or future instructions:

2 PCs if you can make it even better!

Hint (-1): Think about the execution at architectural level.

Code Optimization: **An Intro**

Reordering stores to optimize memory accesses:

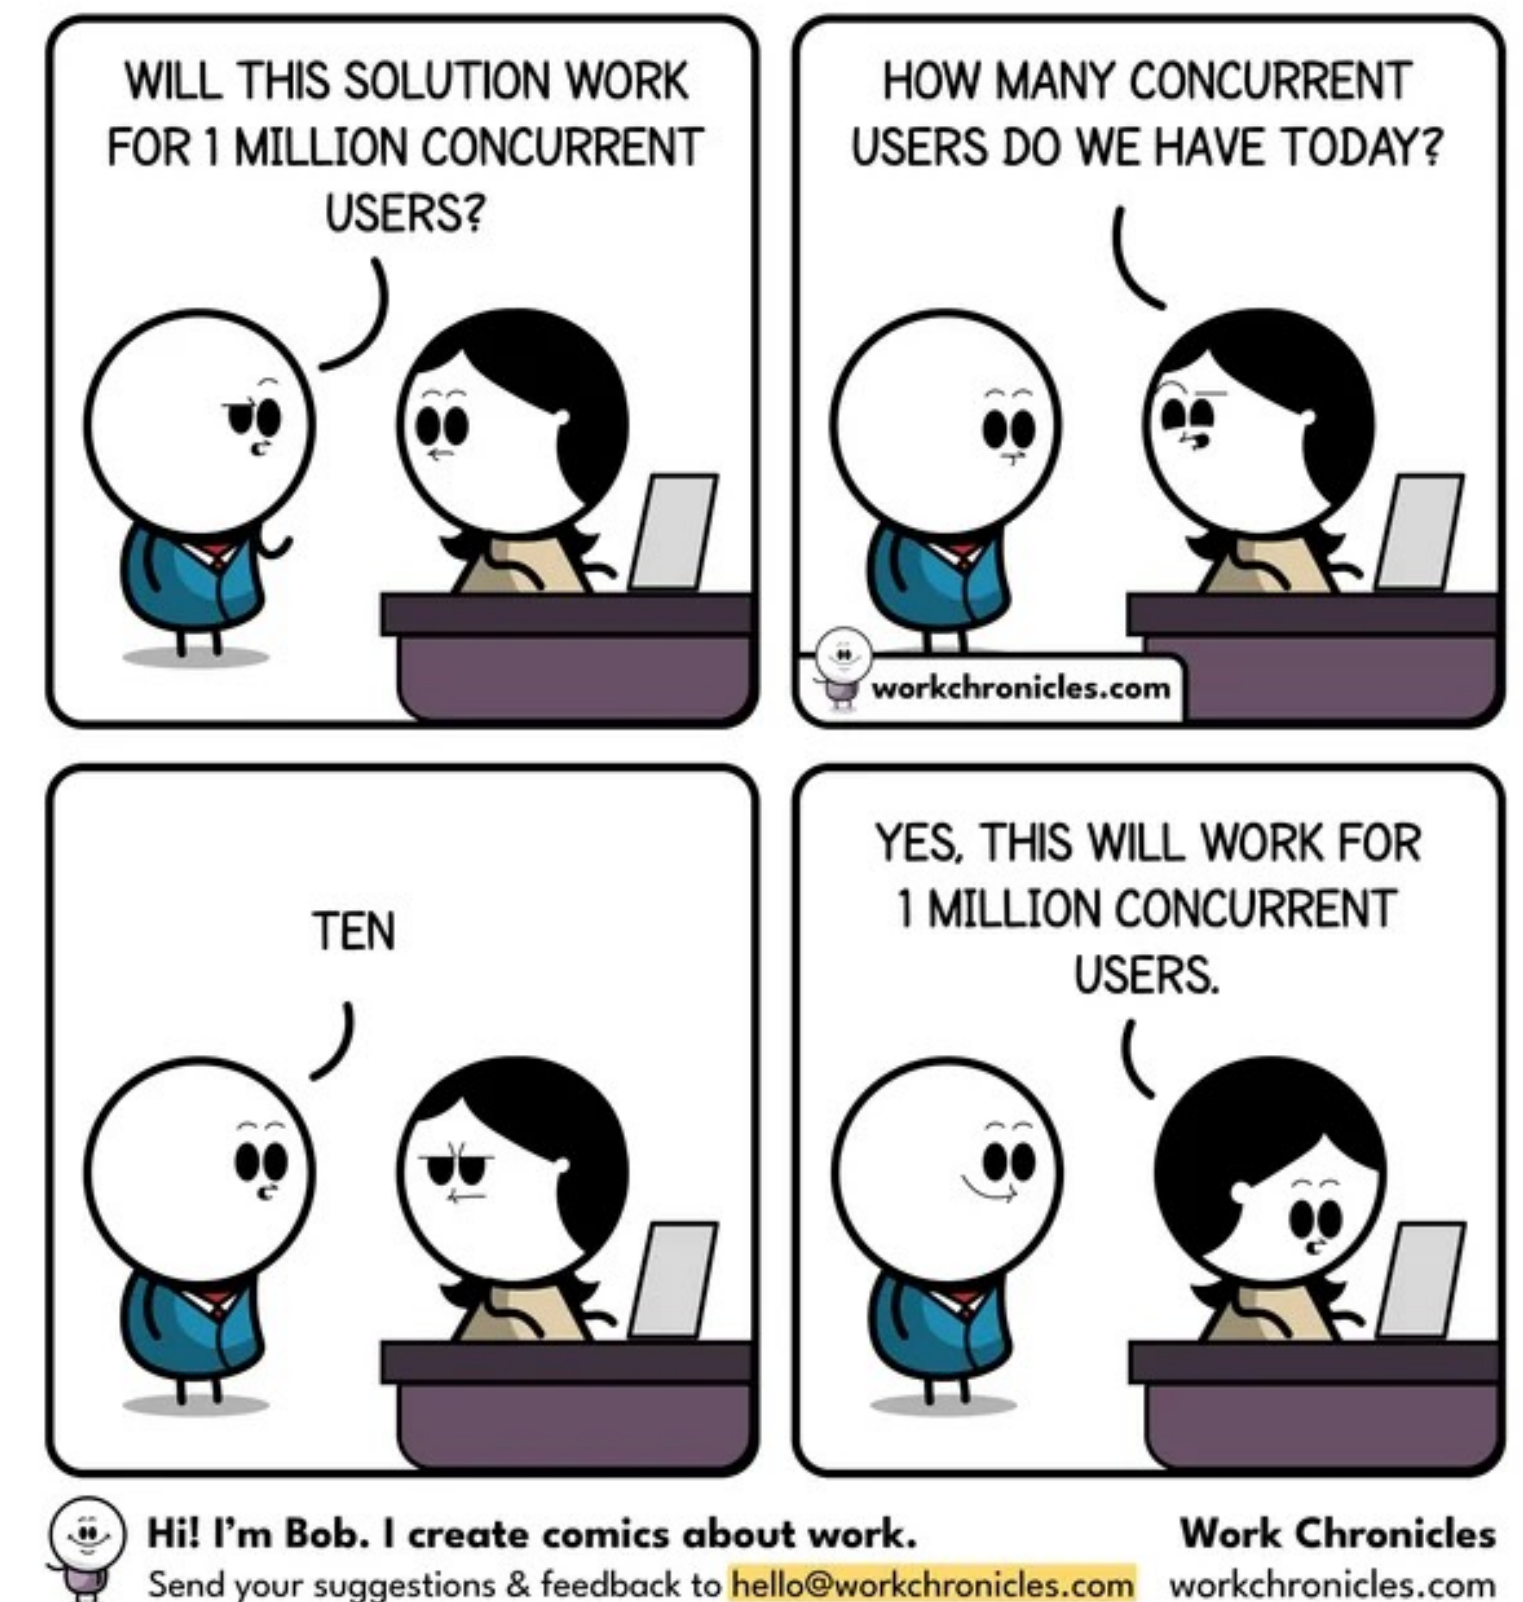
```
LDA R1 c
INC R1
STA c R1
STI a 2
```



```
STI a 2
LDA R1 c
INC R1
STA c R1
```

Foor for thought: Impact of multithreaded programs on possible/allowed reorderings.

```
1: a = 0
2: b = a + 1
3: c = c + b
4: a = b * 2
```



This is where we are headed now!


```
MOV R1 0
STA a R1
```

```
MOV R1 1
LDA R2 a
ADD R2 R1
STA b R2
```

```
LDA R1 b
LDA R2 c
ADD R2 R1
STA c R2
```

Cost:
Registers: 2
Space: 42 bytes
Time: 44 cycles

```
MOV R1 2
LDA R2 b
MUL R2 R1
STA a R2
```



```
STI a 2
LDA R1 c
INC R1
STA c R1
```

Cost:
Registers: 1
Space: 11 bytes
Time: 13 cycles
+Reordered



Full employment theorem for compiler writers

- **Statement:** There is no *fully optimizing* compiler.
- Assume it exists:
 - such that it transforms a program P to the smallest program $\text{Opt}(P)$ that has the same behaviour as P .
 - Halting problem comes to the rescue:
 - Smallest program that never halts:
`L1: goto L1`
 - Thus, a fully optimizing compiler could solve the halting problem by checking if a given program is
`L1: goto L1!`
 - But HP is an undecidable problem.
 - Hence, a fully optimizing compiler can't exist!
- Therefore we talk just about an *optimizing compiler*,
and keep working without worrying about future prospects!



How to perform optimizations?

➤ Analysis

- Go over the program
- Identify some properties
 - Potentially useful properties

➤ Transformation

- Use the information computed by the analysis to transform the program
 - without affecting the semantics

➤ Example:

- Compute liveness information
- Delete assignments to variables that are dead



Many many optimizations

- *Constant folding, constant propagation, tail-call elimination, redundancy elimination, dead-code elimination, loop-invariant code motion, loop splitting, loop fusion, strength reduction, inlining, scalarization, synchronization elision, cloning, data prefetching, parallelization . . . etc.*
- How do they interact?
 - **Optimist:** we get the sum of all improvements.
 - **Realist:** many are in direct opposition.
- Let us *study* some of them!



Constant propagation

- **Idea:** If the value of a variable is known to be a constant at compile-time, replace the use of the variable with the constant.

```
n = 10;  
c = 2;  
for (i=0; i<n; ++i)  
    s = s + i * c;
```



```
n = 10;  
c = 2;  
for (i=0; i<10; ++i)  
    s = s + i * 2;
```

- Usually a very helpful optimization
 - e.g., Can we now *unroll* the loop?
 - Why is it good?
 - Why could it be bad?
- When can we eliminate `n` and `c` themselves?
- Now you know how well different optimizations might interact.



Constant folding

- Idea: If operands are known at compile-time, evaluate expression at compile-time.

```
r = 3.141 * 10;
```




```
r = 31.41;
```

- What if the code was?

```
PI = 3.141;  
r = PI * 10;
```

- And what now?

```
PI = 3.141;  
r = PI * 10;  
d = 2 * r;
```



Called **partial evaluation**

Common sub-expression elimination


- **Idea:** If a program computes the same value multiple times, reuse the value.



- Subexpressions can be reused until operands are redefined.

Copy propagation

- Idea: After an assignment $x = y$, replace the uses of x with y .



- Can only apply up to another assignment to x , or ... another assignment to y !
- What if there was an assignment $y = z$ earlier?
 - Apply transitively to all assignments.

Dead-code elimination

- **Idea:** If the result of a computation is never used, remove the computation.



- Remove code that assigns to dead variables.
 - Liveness analysis done before would help!
- This may, in turn, create more dead code.
 - Dead-code elimination usually works transitively.

Unreachable-code elimination

- **Idea:** Eliminate code that can never be executed

```
#define DEBUG 0  
if (DEBUG)  
    print("Current value = ", v);
```



- High-level: look for `if (false)` or `while (false)`
 - perhaps after constant propagation!
- Low-level: more difficult
 - Code is just labels and gotos
 - Traverse the program (as per its flow), marking reachable statements

Observations

- Some optimizations can be done again after doing them once.
- Some optimizations may get enabled after performing other optimizations.
- Almost all optimizations require information from some pre-performed analysis.
- Many analyses require information connecting variable **uses to their definitions**.
- Compilers require information from many such “**dataflow analyses**” to optimize code *soundly* and *precisely*.
 - Improvements in the underlying analyses may significantly impact other compilation passes.
- One of the important program representations that simplifies many tasks of a compiler is the **SSA form**.

