

CS614: Advanced Compilers

Function Inlining and Devirtualization

Manas Thakur
CSE, IIT Bombay



Spring 2025

How do we call and return from functions?

➤ In the caller:

- Save state of current procedure
 - Program counter (where to resume)
 - Registers (holding current computations)
- Store arguments in a callee-accessible location
- Transfer control-flow

➤ In the callee:

- Collect parameters
- Declare variables
- Perform computations (perhaps in temporaries)

➤ Return to caller:

- Store return value in a caller-accessible location

```
void foo() {  
    ...  
    x = 20;  
    r = bar(p, q);  
    y = x + r;  
    ...  
}  
int bar(int x, int y) {  
    return x + y;  
}
```

Function calls are quite expensive.



Every analysis/optimization
we have seen till now is
intraprocedural.



How do we optimize across function calls?

- Control Flow Graphs cannot be used to optimize across function boundaries.
- Without interprocedural constant propagation, the safest “correct” value of x and y (in `foo`) is bottom!
 - x must not be modified in `bar`.
 - y ’s value needs to be determined by analyzing `bar`.
- Same story for all optimizations across function boundaries.

```
void foo() {  
    ...  
    x = 20;  
    p = 10;  
    q = 30;  
    r = bar(p, q);  
    y = x + r;  
    // constant propagation?  
}  
int bar(int x, int y) {  
    return x + y;  
}
```

Function calls
inhibit optimization.

Compilers try to
inline functions!



Function Inlining

➤ Idea:

- Replace a function call with the body of the callee

➤ Benefits:

- Eliminate call/return overhead
- Increase the scope of performing optimizations
 - More constant propagation
 - More partial evaluation
 - More everything else!
- Hardware:
 - Eliminate two jumps
 - Keep the pipeline filled



How to inline a function?

- Attempt 1: Just copy-paste
 - With assignments from arguments to parameters
 - And of the return value

```
void foo() {  
    ...  
    x = 20;  
    r = bar(p, q);  
    y = x + r;  
    ...  
}  
int bar(int x, int y) {  
    return x + y;  
}
```

```
void foo() {  
    ...  
    x = 20;  
    x = p;  
    y = q;  
    r = x + y;  
    y = x + r;  
    ...  
}  
int bar(int x, int y) {  
    return x + y;  
}
```

- Problem: Final y was earlier $20+p+q$; now it is $2p+q$.



How to inline a function?

- Rename variables (uniquely):
 - With assignments from arguments to parameters
 - And of the return value

```
void foo() {  
    ...  
    x = 20;  
    r = bar(p, q);  
    y = x + r;  
    ...  
}  
int bar(int x, int y) {  
    return x + y;  
}
```

```
void foo() {  
    ...  
    x = 20;  
    bar_x = p;  
    bar_y = q;  
    r = bar_x + bar_y;  
    y = x + r;  
    ...  
}  
int bar(int x, int y) {  
    return x + y;  
}
```

- Final y is back to $20+p+q$:-)



Our inlining algorithm

A. Rename variables in the function being inlined

- Follow use-def chains

B. Add assignments from arguments to parameters

C. Copy-paste code from callee to caller

D. Replace return statements with assignments to collecting variable

- What if there are multiple return statements in the callee?

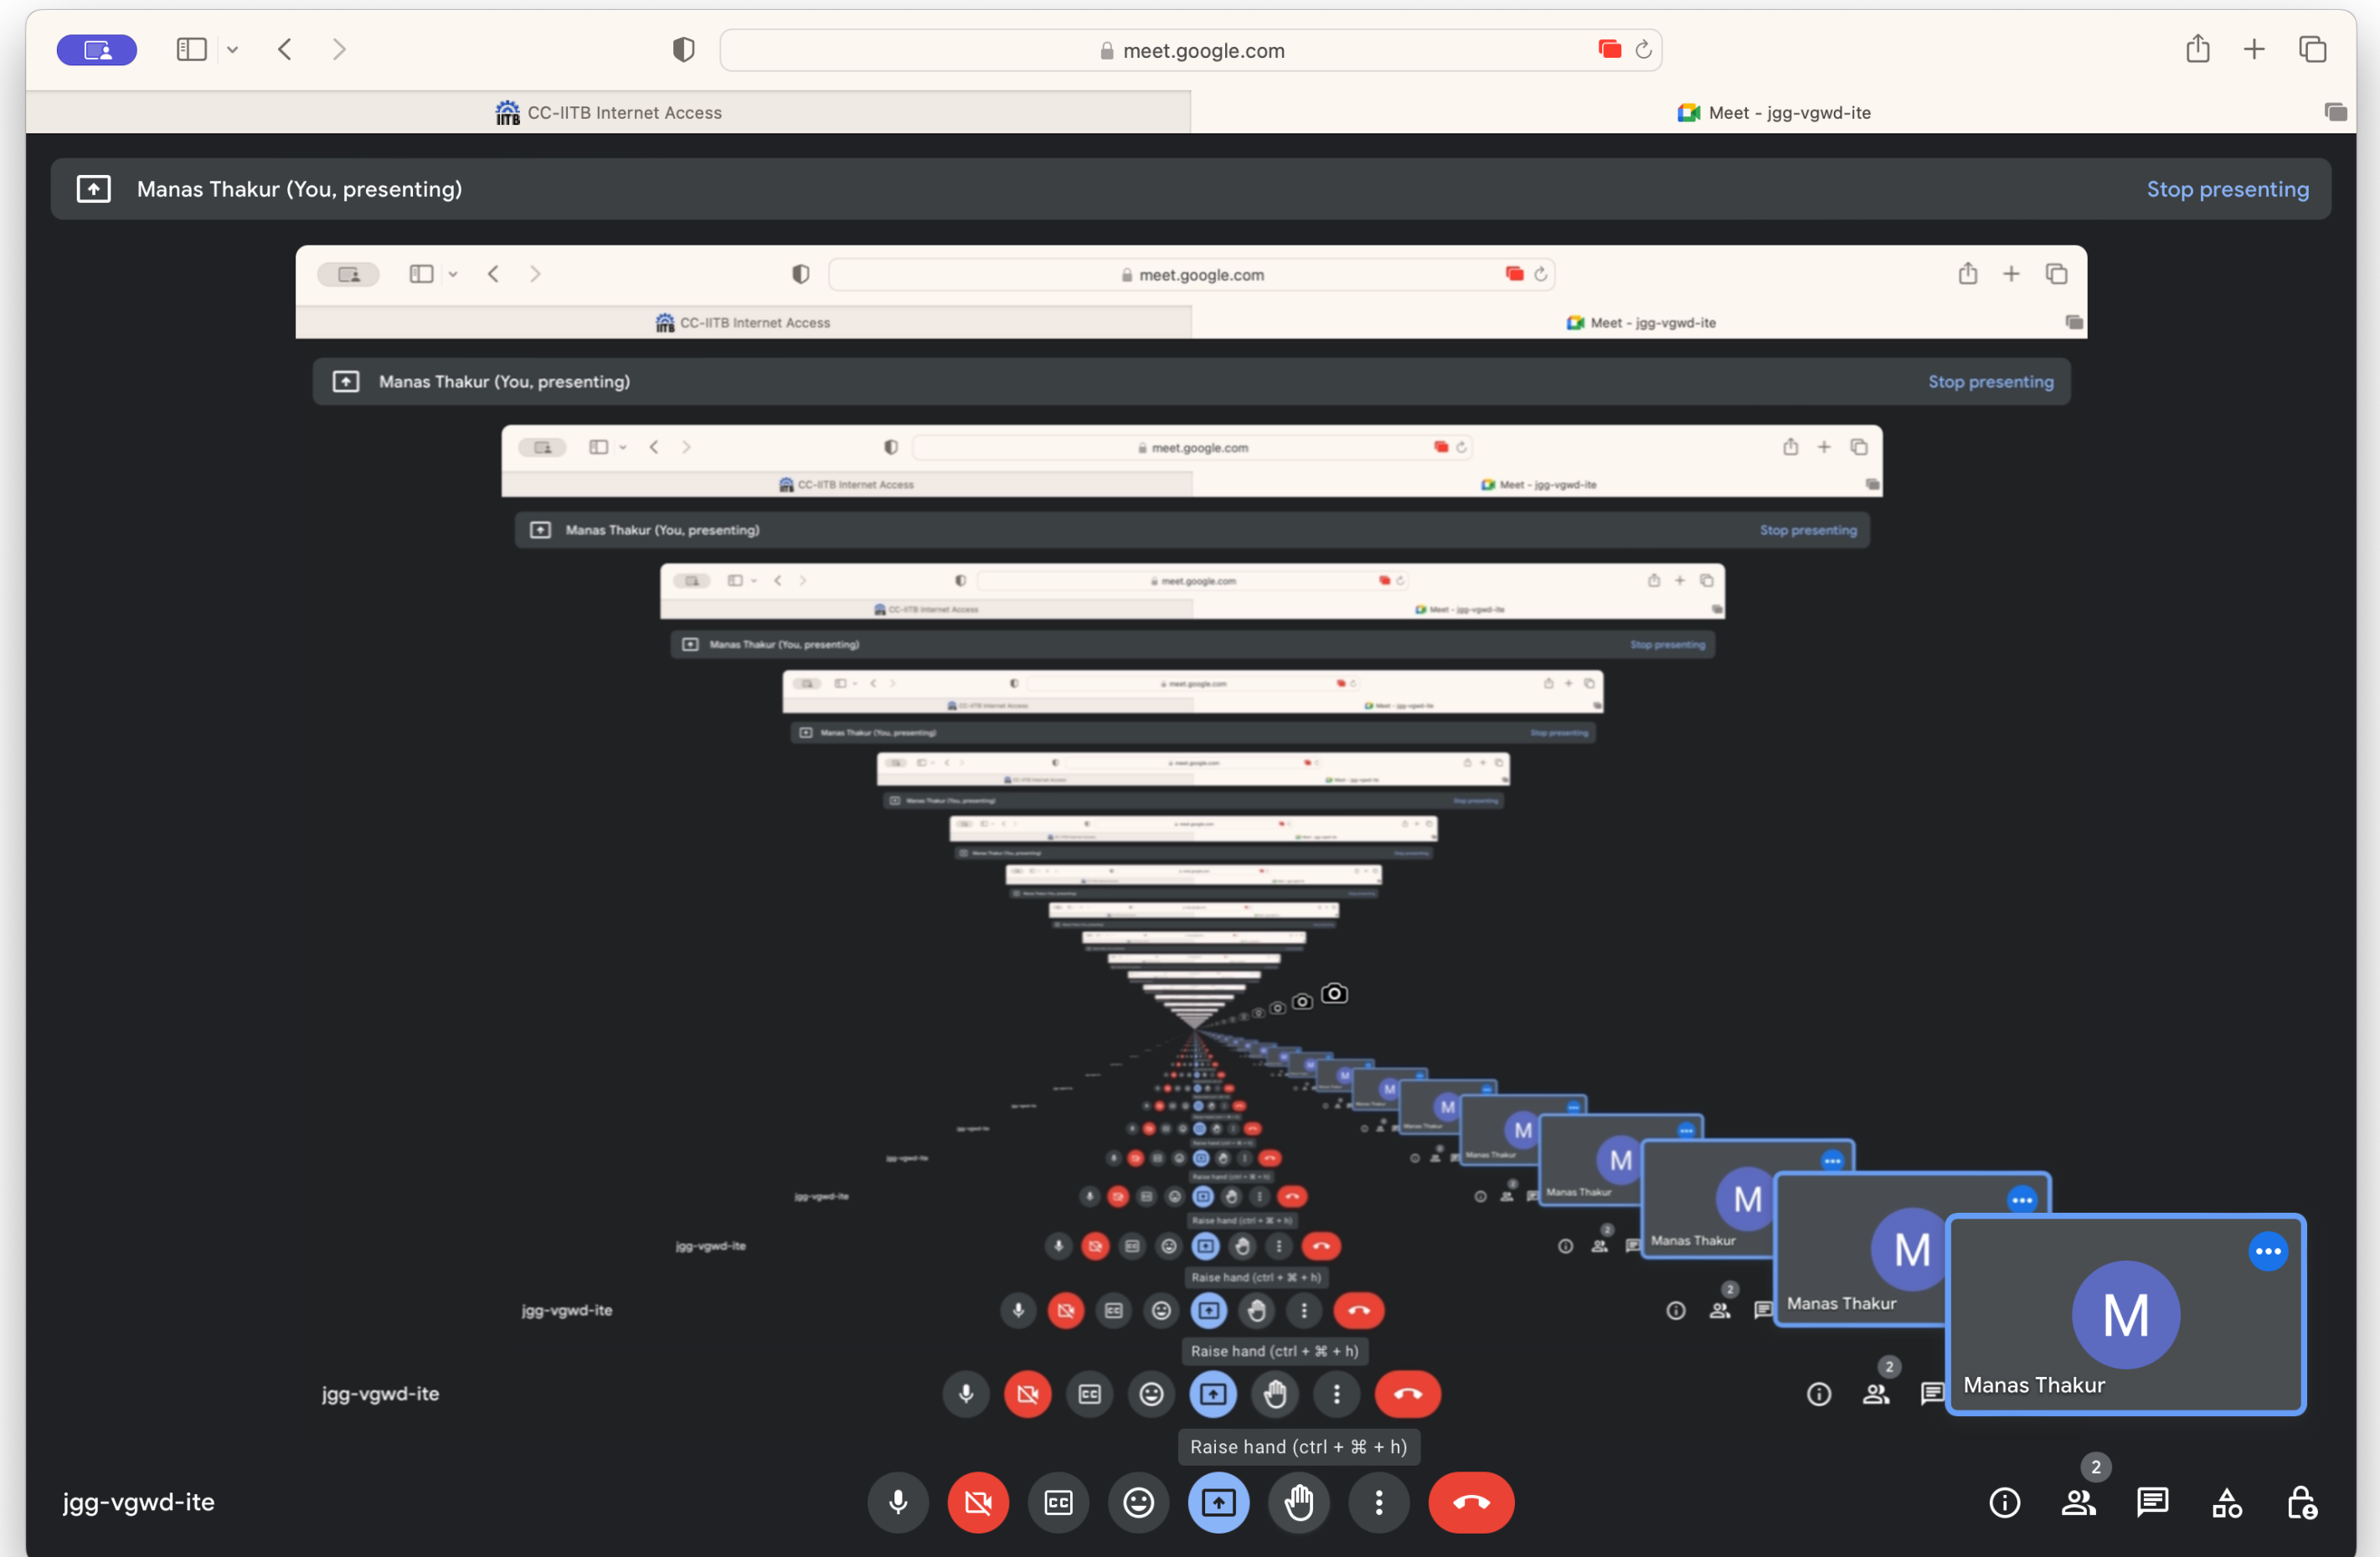
- Later, let the other optimization passes run over the modified CFG of the caller.



Can we always inline a function at its call site?

- Not if the calls are recursive

```
int foo(int z) {  
    ...  
    x = 20;  
    r = bar(p, q);  
    y = z + r;  
    return y;  
}  
int bar(int x, int y) {  
    if (x < y) {  
        x += foo(x);  
    }  
    return x;  
}
```



Can we always inline a function at its call site?

- Not if the call is virtual

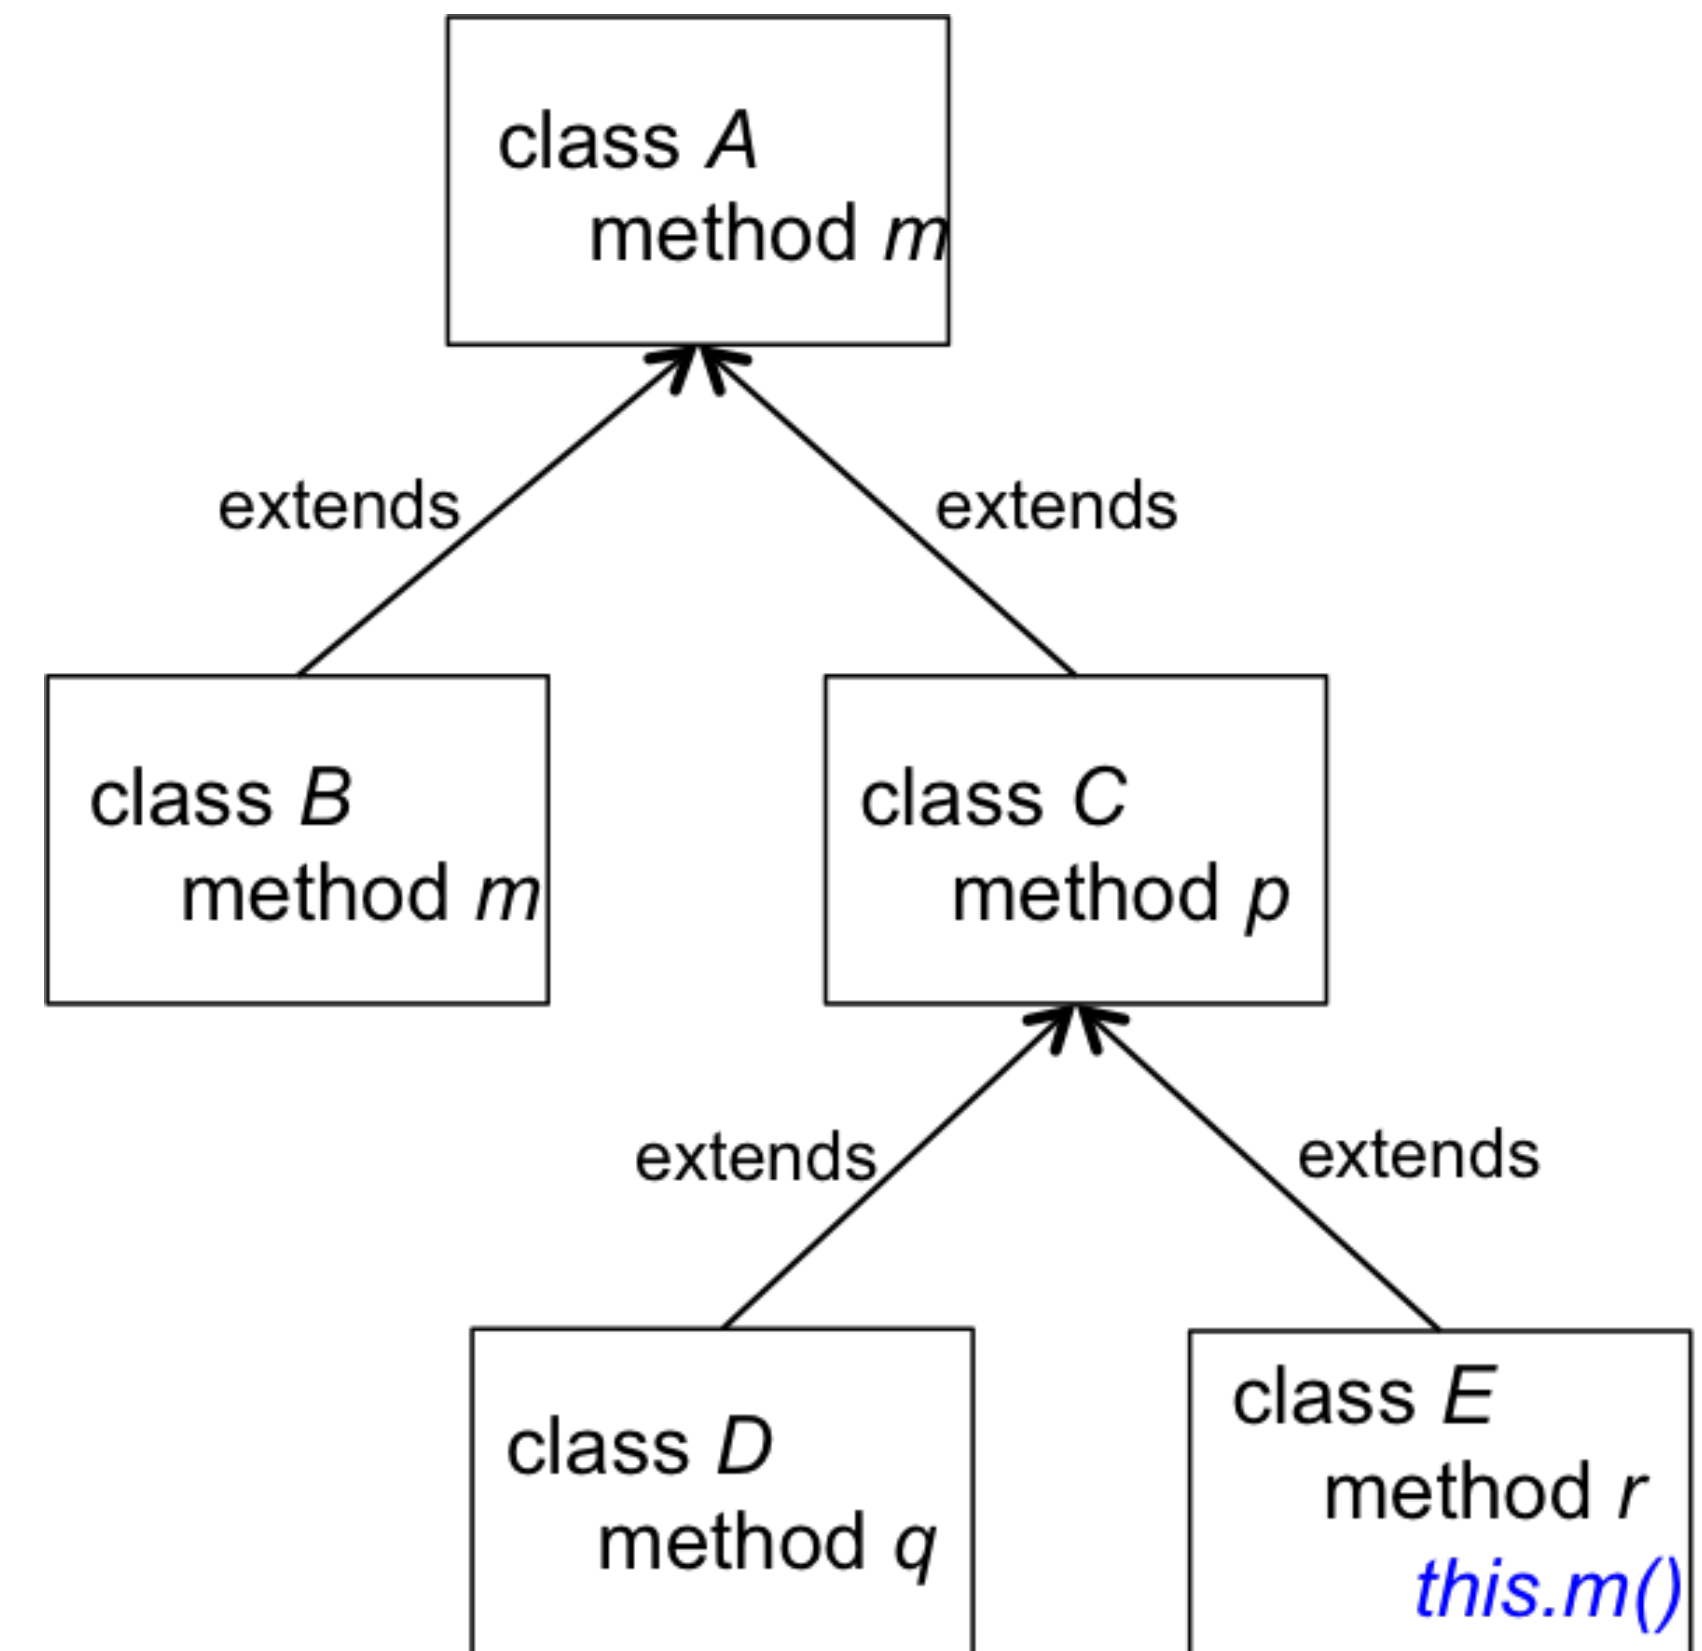
```
void foo(T x) {  
    U z;  
c1:int y = x.bar();  
    if (y > 40)  
        z = new V();  
    else  
        z = new W();  
c2:z.zap();  
}
```

- Which method(s) can be called at c1 and c2?
- If only one can be called (i.e. if we can devirtualize c1/c2), we can inline.



Devirtualization with Class Hierarchy Analysis

- Look at the **class hierarchy** of the type of the receiver to determine what methods can be called using the same.
- If a class re-defines a method **foo** defined in its parent class:
 - Only the redefined **foo** can be called using the objects of the extended class.
 - Methods not redefined are still accessible from the parent class.
- **CHA** helps in determining that only one implementation of **m** can be called from the blue call.



Devirtualization with Rapid Type Analysis

- Improves CHA with extra information:
 - Find if a class is ever instantiated in a program
 - If not, then remove it from the sets obtained using CHA
- Catch:
 - Assumption that the whole program is available
 - What about code that cannot be “seen”?
 - e.g., dynamically linked libraries
 - Even more interestingly done in the JVM (later).

```
class A {  
    A foo(A x) {...}  
}  
class B extends A {  
    A foo(A x) {...}  
}  
class D extends A {  
    A foo(A x) {...}  
}  
... new D() ...  
...  
void bar(A y) {  
    y.foo();  
    // CHA: {A.foo, B.foo, D.foo}  
    // RTA: {D.foo}  
}
```



Devirtualization with Control-Flow Analysis

- Tries to find which classes of objects can flow to the receiver reference variables.
- Higher precision of devirtualization than CHA and RTA.

```
class A { A foo(A x) { return x; } }  
class B extends A { A foo(A x) { return new D(); } }  
class D extends A { A foo(A x) { return new A(); } }  
class C extends A { A foo(A x) { return this; } }
```

```
void main() {  
    A x = new A();  
    while (...)  
c1:    x = x.foo(new B());  
    A y = new C();  
c2:    y.foo(x);  
}
```

- **CHA/RTA:** Any of the foos can be called at both call sites **c1** and **c2**.
- **Control-flow analysis:**
 - **c1:** {A.foo, B.foo, D.foo}
 - **c2:** {C.foo}



Should we always inline an **inlinable** function at **all its call sites**?

- Inlining a function increases code size.

- This may also increase cache misses.

- A phenomenon very similar to *loop unrolling*.

Instead of inlining, compilers may also replace “monomorphic” indirect calls with direct calls.

- Thus, **most compilers are very careful while making inlining decisions**:

- Inline only if the callee is $<K$ bytes.

- Inline only if the caller does not grow beyond $>K$ bytes.

- Perform nested inlining only up to a threshold.

- Still, method inlining is among the largest sources of optimization in OO language compilers, and developing heuristics for the same is an **interesting research problem**.



Improving the Precision of Control-Flow Analysis

- The CFA we saw earlier does not distinguish the calls made to the current method from different call sites.
 - It is *context-insensitive*.
- We can improve precision using a *context-sensitive* analysis (k -CFA),
 - where k is the length of the current method's call chain.

```
class A {  
  fb() { ... }  
class B extends A {  
  fb() { ... }  
}
```

```
class C {  
  void foo() {  
    A a1, a2;  
    a1 = new A(); //l1  
c1: bar(a1);  
    a2 = new B(); //l2  
c2: bar(a2);  
  }  
  void bar(A p1) {  
    p1.fb();  
  }  
}
```

➤ Context insensitive:

➤ $p1 \rightarrow \{l1, l2\}$

➤ 1-CFA:

➤ At $c1$: $p1 \rightarrow \{l1\}$

➤ At $c2$: $p1 \rightarrow \{l2\}$



k -CFA (Cont.)

```
class C {  
  void main() {  
    foo();  
    ...  
    foo();  
  }  
  
  void foo() {  
    bar();  
  }  
  
  void bar() {  
    fb();  
  }  
}
```

- 1-CFA:
 - 1 context for fb
- 2-CFA:
 - 2 contexts for fb
- 3-CFA?

The Opposite of *Polymorphism*

- **Monomorphic call:**

- Only one method can be called
- Target can either be bound statically, and sometimes inlined

- **Which calls in Java are always monomorphic?**

- Calls to static methods
- Calls to final methods
- Calls to methods of final classes

- We increase the size of the above set using the different analyses seen in the previous slides.

- **How could we compare the precision of two CG construction algorithms empirically?**

- Count the identified number of monomorphic call-sites.



What we have seen today

- **CHA:** only look at inheritance relations
 - **RTA:** also look at code
 - **0-CFA:** also analyze code
 - **k -CFA:** contextualize the analysis
-
- **Some points to note:**
 - k -CFA is a context-sensitive analysis
 - The length k is the length of the call-stack
 - Here the last k call-sites form our context



What we haven't seen today

- There are multiple ways we can define what is a context
 - Called the **context abstraction** of a given context-sensitive analysis

- **Object-sensitive contexts**

Ana Milanova, Atanas Rountev, Barbara G. Ryder. TOSEM '05.

- **Value contexts**

Uday P. Khedker and Bageshri Karkare. CC'08.

- **Type-sensitive contexts**


Yannis Smaragdakis, Martin Bravenboer, Ondřej Lhoták. POPL '11.

- **LSRV contexts**

Manas Thakur and V. Krishna Nandivada. CC'19.

- Many more **non-context-sensitive** algorithms:

Frank Tip and Jens Palsberg. OOPSLA '00.



No time for these yet!