

CS614: Advanced Compilers

Cache Optimizations

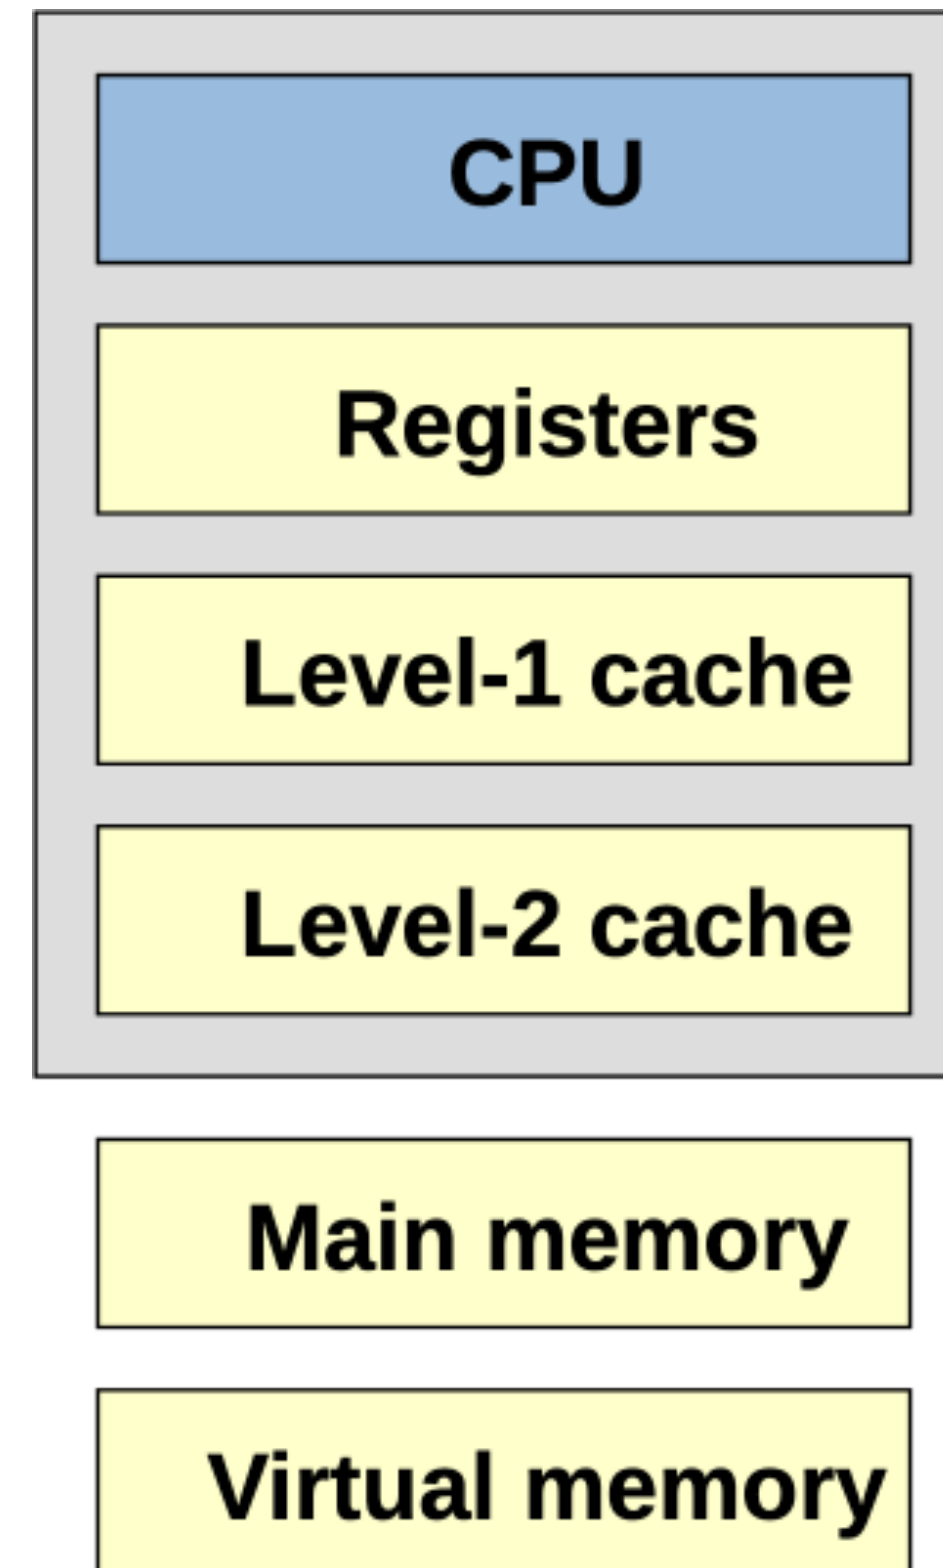

Manas Thakur
CSE, IIT Bombay



Spring 2025

Remember the hierarchy?

*Farther away,
larger,
slower*



<i>Pentium 4 3.2 Ghz</i>	<i>Core 2 Duo</i>	<i>Athlon 64</i>
1 cycle	1 cycle	1 cycle
2 cycles (16 KB)	3 cycles (64 KB)	3 cycles (128 KB)
19 cycles (2 MB)	14 cycles (2 MB)	13 cycles (1 MB)
204 cycles	180 cycles	125 cycles
millions of cycles	millions of cycles	millions of cycles

- We have worked with registers; now is the time to move down (**cache memory**).

Data Prefetching

- We have cache memory for *temporal* and *spatial* locality.
 - **Temporal**: Same data may be accessed **again**.
 - **Spatial**: Nearby data may be accessed **soon**.
- **Data prefetching** aims to utilize temporal and spatial locality (through the compiler).

prefetch is a hardware-supported instruction, and needs to be preserved across compiler transformations.

```
for (i = 0; i < N; i++) {  
    sum = sum + A[i];  
}
```



```
for (i = 0; i < N; i++) {  
    prefetch(A[i+pd]);  
    sum = sum + A[i];  
}
```

- Prefetch distance (pd) specifies the number of iterations before prefetching a data element and its access.
 - Needs to be determined carefully based on access pattern and latency.



Accuracy and Timeliness

- Is there a guarantee that the prefetched data will be accessed in future?
- What if the prefetch distance crosses the data structure boundary?
- Should we always attempt to compute the prefetch distance?

```
for (i = 0; i < N; i++) {  
    if (A[foo(i)] > 100)  
        sum = sum + B[i];  
}
```

- foo may have side effects
- Call to foo may be expensive

- What if the prefetch is done so late that its benefits do not even matter?
 - Most trivial (but highly likely useless) pd value is 0!
- What if prefetch is done so early that the prefetched data gets evicted before use?
 - But we should be able to finish the prefetch before starting with the iteration that needs the prefetched value.



Profitability

- No. of cache misses without prefetching = $MAX/64$
- Prefetching overhead = MAX cycles
- Even if all the misses are prefetched,
 - Cycles saved by prefetching = $MAX/64 * c$
- For prefetching to be beneficial,
 - $MAX/64 * c > MAX \implies c > 64$
- What if we *unroll the loop* body 64 times?
 - Prefetching overhead reduces to $MAX/64$ cycles!
- Thus, prefetching is beneficial when,
 - $MAX/64 * c > MAX/64 \implies c > 1$
 - Thus, even a cache-miss cost of >1 cycle would imply that prefetching is beneficial.

L1 cache line size: 64 Bytes
Issue width: 1
Cache miss cost: c cycles

```
char A[MAX];  
for (i = 0; i < MAX; i++) {  
    prefetch(A[i+pd]);  
    sum += A[i];  
}
```



Prefetching *affine* array accesses

1. Perform locality analysis.

- A is *reused* in each iteration with an increment of 1.

```
for (i = 0; i < N; i++) {  
    sum = sum + A[i];  
}
```

2. Identify accesses requiring prefetch.

- Assuming L1 cache line is 64 Bytes and A is stored in row-major order, we may prefetch every 64 iterations.
- Identify *leading references*: e.g. 0, 64, 128, etc.

3. Perform *loop decomposition* to insert prefetch instructions at the right place.

Avoids unnecessary prefetch every original iteration.

```
for (i = 0; i < N; i += 64){  
    // prefetch here  
    sum += A[i];  
    for(ii = i+1; ii < min(N, i+64); ii++){  
        sum += A[ii];  
    }  
}
```



Prefetching non-array accesses

- Arrays in loops often have regular access patterns (i.e., constant *strides*).
- How about arrays with non-affine indices?
- Traversal of a recursive data structure such as a linked list?
- Compilers may need *profiling feedback* to insert suitable prefetch instructions.
- If we know accesses are regular but the stride keeps changing after some time, we can compute *pd* dynamically:

```
while (ptr != NULL) {  
    sum += *ptr;  
    ptr = ptr->next;  
}
```



```
prev = ptr;  
while (ptr != NULL) {  
    sum += *ptr;  
    stride = ptr - prev;  
    prefetch(ptr+K*stride);  
    prev = ptr;  
    ptr = ptr->next;  
}
```

K needs to be chosen carefully based on the profile.

Data Layout Transformations

- Data prefetching aims to bring future accesses to the cache ahead of time.
- What is the **underlying idea**?
 - Bring in *frequently-accessed-together* data items in the cache at the same time.
 - Needs to be done because such data items may not be allocated together in memory.
- How about already keeping such data items **together in the memory itself**?
- Leads way to compilers changing the layout of data structures at the source-code level itself.



Field layout

- Group **hot fields** together and separate them from **cold fields**.
 - Can be found statically, dynamically, or with a hybrid strategy.
- Optimization called **structure splitting**:

```
struct S1 {  
    char hot[4];  
    char cold[60];  
};  
S1 s1[512];
```



```
struct S1_cold {  
    char cold[60];  
};  
struct S1 {  
    char hot[4];  
    struct S1_cold *cold_link;  
};  
S1 s1[512];
```

- Good: **s1** may now fit within the L1 cache.
- (Slightly) Bad: Additional level of indirection to access the cold fields.

Field layout (Cont.)

- We can even get rid of the indirection cost by *structure peeling*:

```
struct S1 {  
    char hot[4];  
    char cold[60];  
};  
S1 s1[512];
```



```
struct S1_cold {  
    char cold[60];  
};  
struct S1_hot {  
    char hot[4];  
};  
S1_hot s1_hot[512];  
S1_cold s1_cold[512];
```

- **Addendum:** Requires program refactoring to modify the usages of s1.
- Difficult when the structure type has a pointer to itself.



Field layout (More considerations)

- Structure peeling may have performance penalties when the size of the record structure is large:

```
struct S2 {  
    char hot1[32];  
    char hot2[32];  
    char not_so_hot[32];  
};  
struct S2 s2[128];
```

```
for (i = 0; i < 128; i++) // Invoked N times  
    x += foo(s2[i].hot1);  
...  
for (i = 0; i < 128; i++) // Invoked N times  
    x += foo(s2[i].hot2);  
...  
for (i = 0; i < 128; i++) // Invoked N/10 times  
    sum3 += foobar(s2[i].not_so_hot) + foo(s2[i].hot1);
```

- Fields hot1 and hot2 are accessed more frequently than the field not_so_hot, but not_so_hot is always accessed before hot1 (so, they should be placed together).
- Benefit of structure peeling depends on the cost of additional misses incurred in the third for loop.



Data object layout

- Size of a single structure may rarely exceed the size of a cache block, but what about laying out (potentially large) objects themselves together?
- We need to be careful about introducing cache-line conflicts:
 - Two distinct cache lines containing two different data objects may conflict in the cache.
 - If those two data objects are accessed together frequently, a large number of conflict misses may result.
- **Example:** Different data-structure instances, if accessed frequently in individuality, may be *pooled* together in memory.


```
struct list;  
struct tree;  
struct linked_list{  
    int n;  
    struct linked_list *next;  
};  
struct tree{  
    int n;  
    struct tree *left;  
    struct tree *right;  
};  
struct linked_list *l1, *l2;  
struct tree *t;  
...
```

*l1, l2 and t may be pooled together,
based on access patterns.*




But we have both D-cache and I-cache

- Most processors have two separate L1 caches for data (**L1D**) and instructions (**L1I**).
- Anything if needs to be fetched from the main memory has higher cost than from the cache.
- Data prefetching could be done based on data-access patterns; how about keeping the right set of future instructions in the instruction cache?
- Requires finding access patterns in the **control flow**.



Code layout in procedures

- Consider the following “profiled” CFG:



- Often, B3 is executed after B1 (e.g., in loops the branch is taken most of the times).
- We may want to place the instructions in B1, B3 (and if fits, even the ones in B4) together in the virtual-address space.
- Optimization called *procedure splitting*.



Cache-line coloring

- Even with hot-code basic-block separation, we may incur cache misses if a caller and callee are mapped to the same cache line.
- foobar calls foo (and transitively, bar) in a hot loop.
- If foo (or bar) are mapped to the same cache line as foobar, we may incur an I-cache miss in each iteration.
- Cache-line coloring takes a program's "call graph" as input, and colors the same such that K levels of callers-callees do not get assigned the same colour.
- Differently colored procedures get different cache lines assigned during execution.

```
void foo(){  
    bar();  
    ...  
}  
int foobar(){  
    for (i=0; i<n; i++)  
        foo();  
    ...  
}
```



Optimizing cache behaviour without prefetching

- Almost all kinds of software (compiler) prefetching require hardware support.
- Layout transformations require changes in the data structures.
- Can we *change computations* instead, to take advantage of locality?
- Next class onwards: **Loop Transformations** (probably the most involved as well as powerful optimizations modern compilers perform).
- Techniques naturally lead way to loop *parallelization* and *vectorization*.

