

# CS614: Advanced Compilers

*Memory Models*

**Manas Thakur**  
CSE, IIT Bombay



Spring 2025

# Multithreaded Programs

---

Initially,  $A = B = 0$ .

## Thread 1

(1)  $A = 1$   
(2) print B

## Thread 2

(3)  $B = 1$   
(4) print A


- What all can be printed?
  - (1) —> (2) —> (3) —> (4): **01**
  - (3) —> (4) —> (1) —> (2): **01**
  - (1) —> (3) —> (2) —> (4): **11**
  - (1) —> (3) —> (4) —> (2): **11**
  - A few others with the same effect.



# Things that shouldn't happen

---

- It should not be possible to print 00.



- In order to print 00:
  - For line (2) to print 0, (2) should *happen before* (3) writes 1 to B.
  - For line (4) to print 0, (4) should happen before (1) writes 1 to A.
  - Each thread's events should be in order: (1) before (2) and (3) before (4).
  - This implies (1) should happen before (1)!
  - A contradiction — means **00 cannot be printed** by this multithreaded program.



# Sequential Consistency

---

1. All operations executed in some sequential order.
  - As if they were manipulating a *single shared memory*.
2. Each thread's operations happen in program order.

Initially,  $A = B = 0$ .

Thread 1

$A = 1$   
 $r0 = B$

Thread 2

$B = 1$   
 $r1 = A$

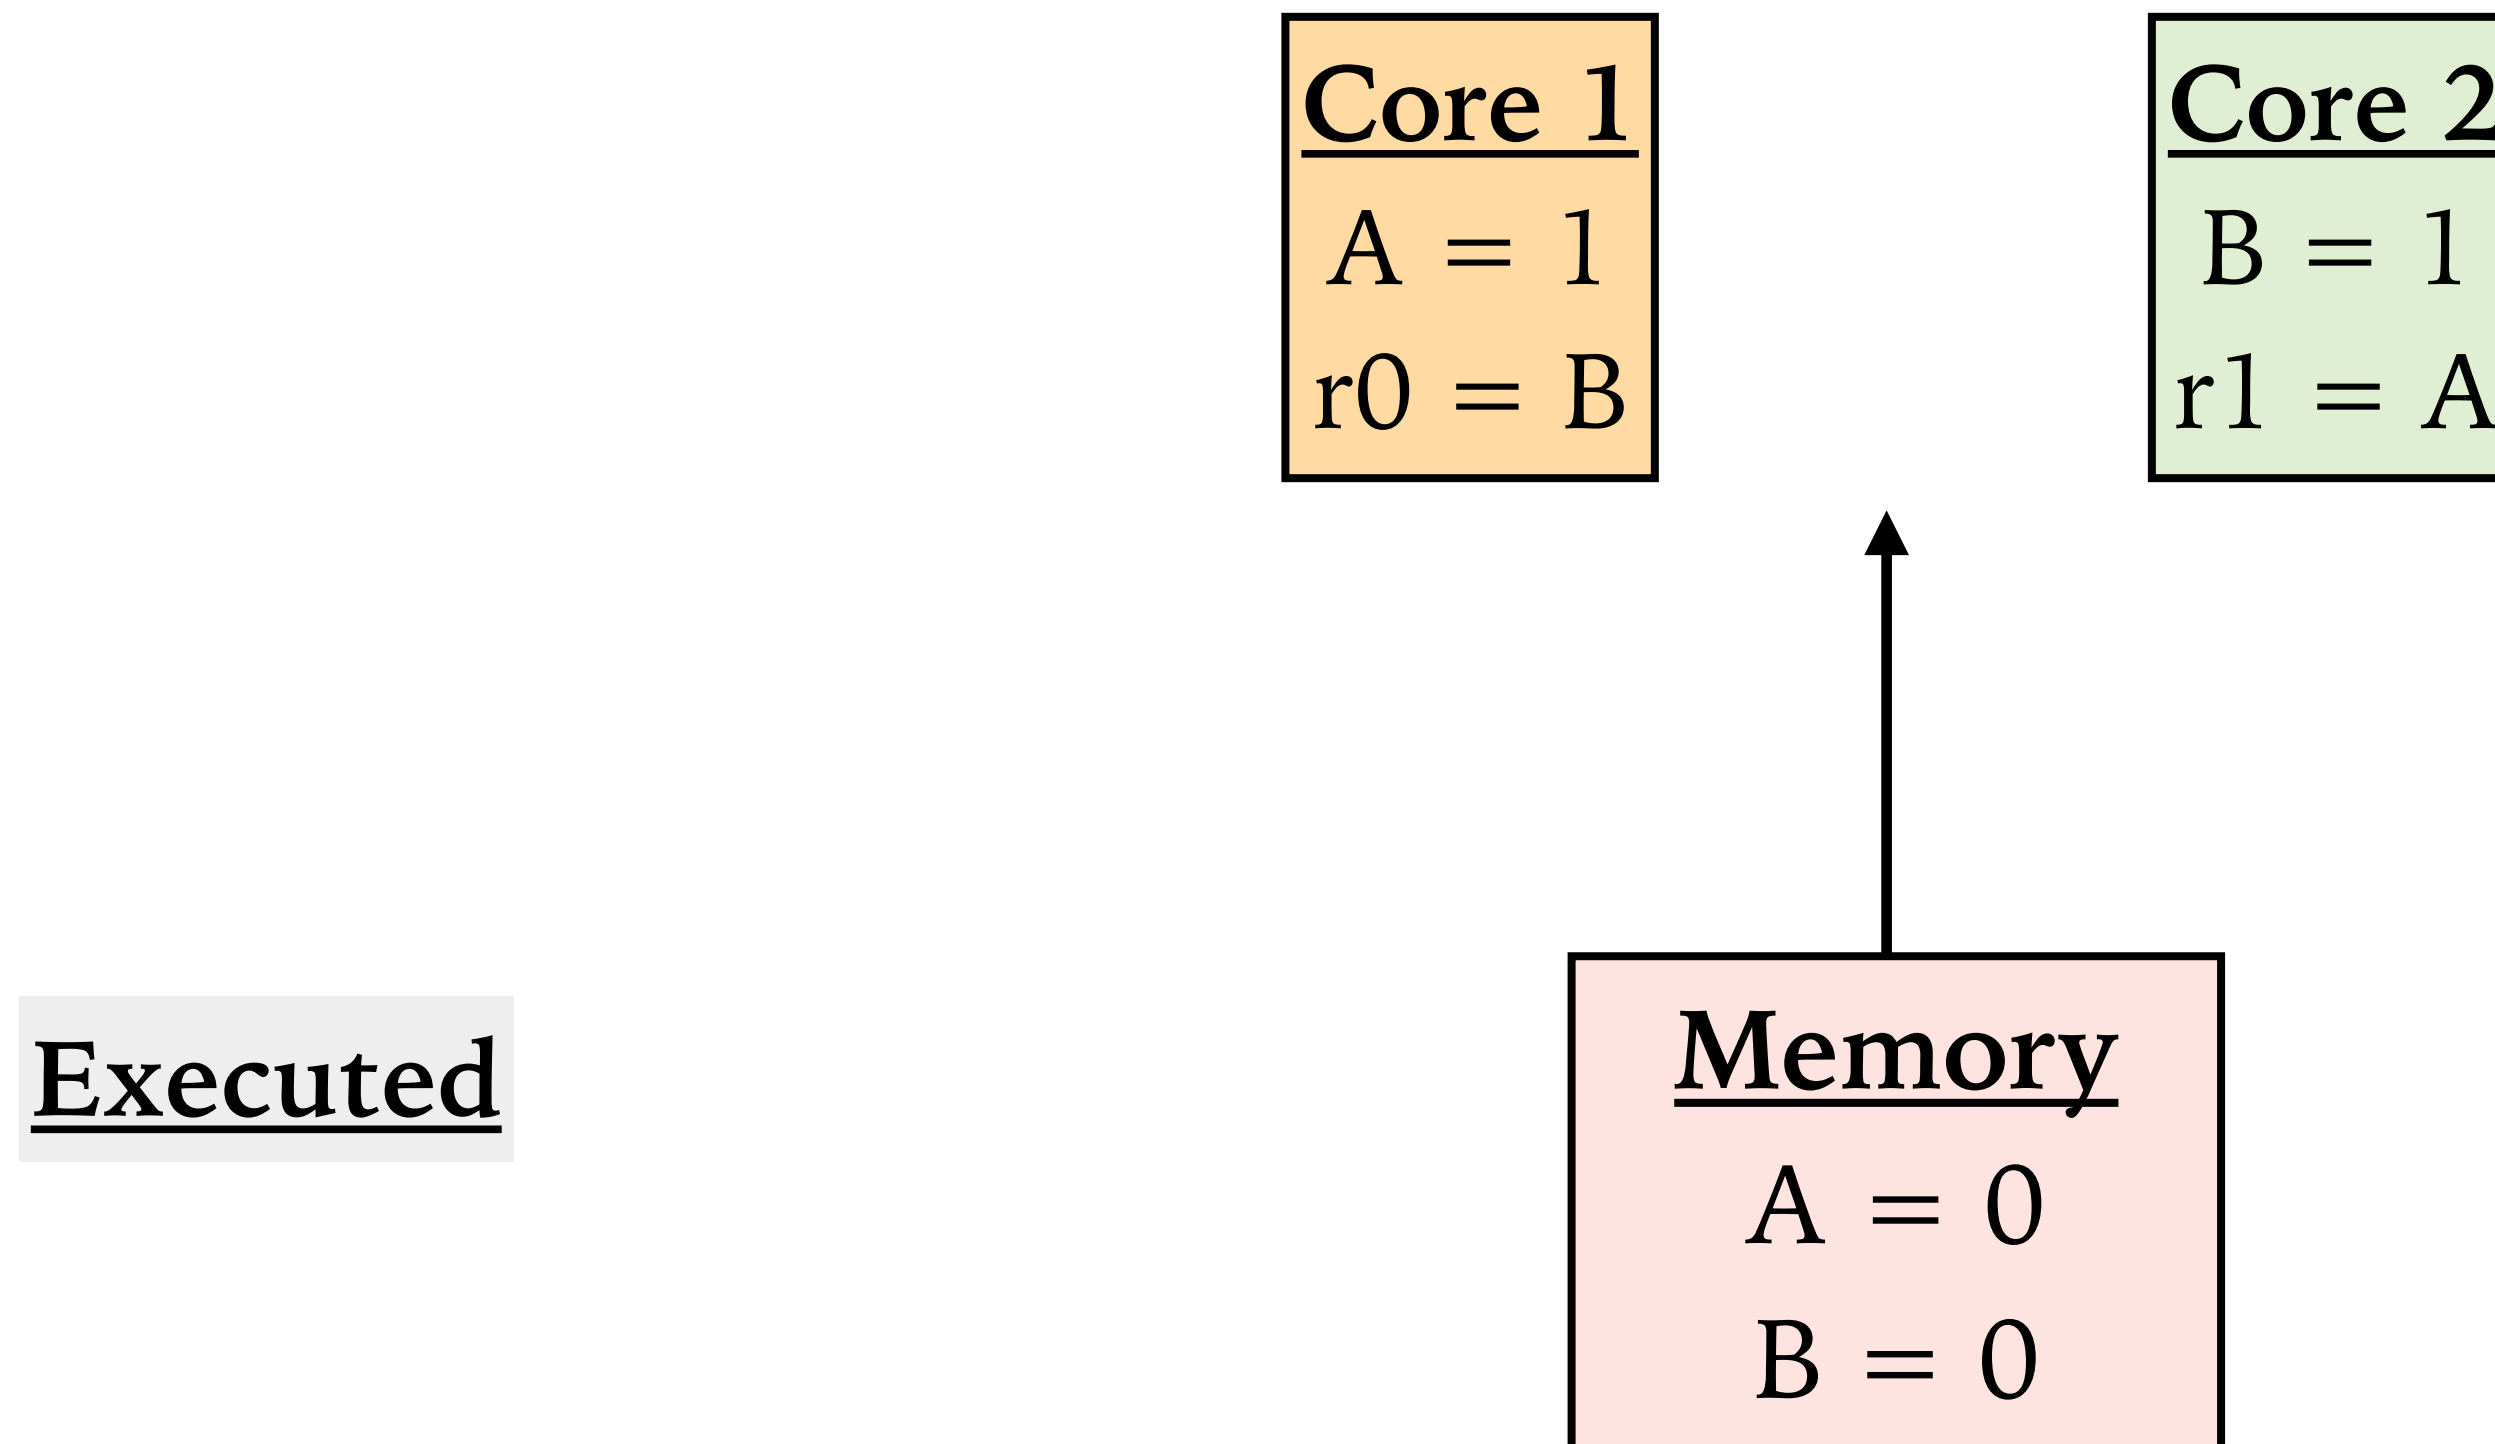
- Not allowed:  $r0 = 0$  and  $r1 = 0$ .



# Sequential Consistency

---

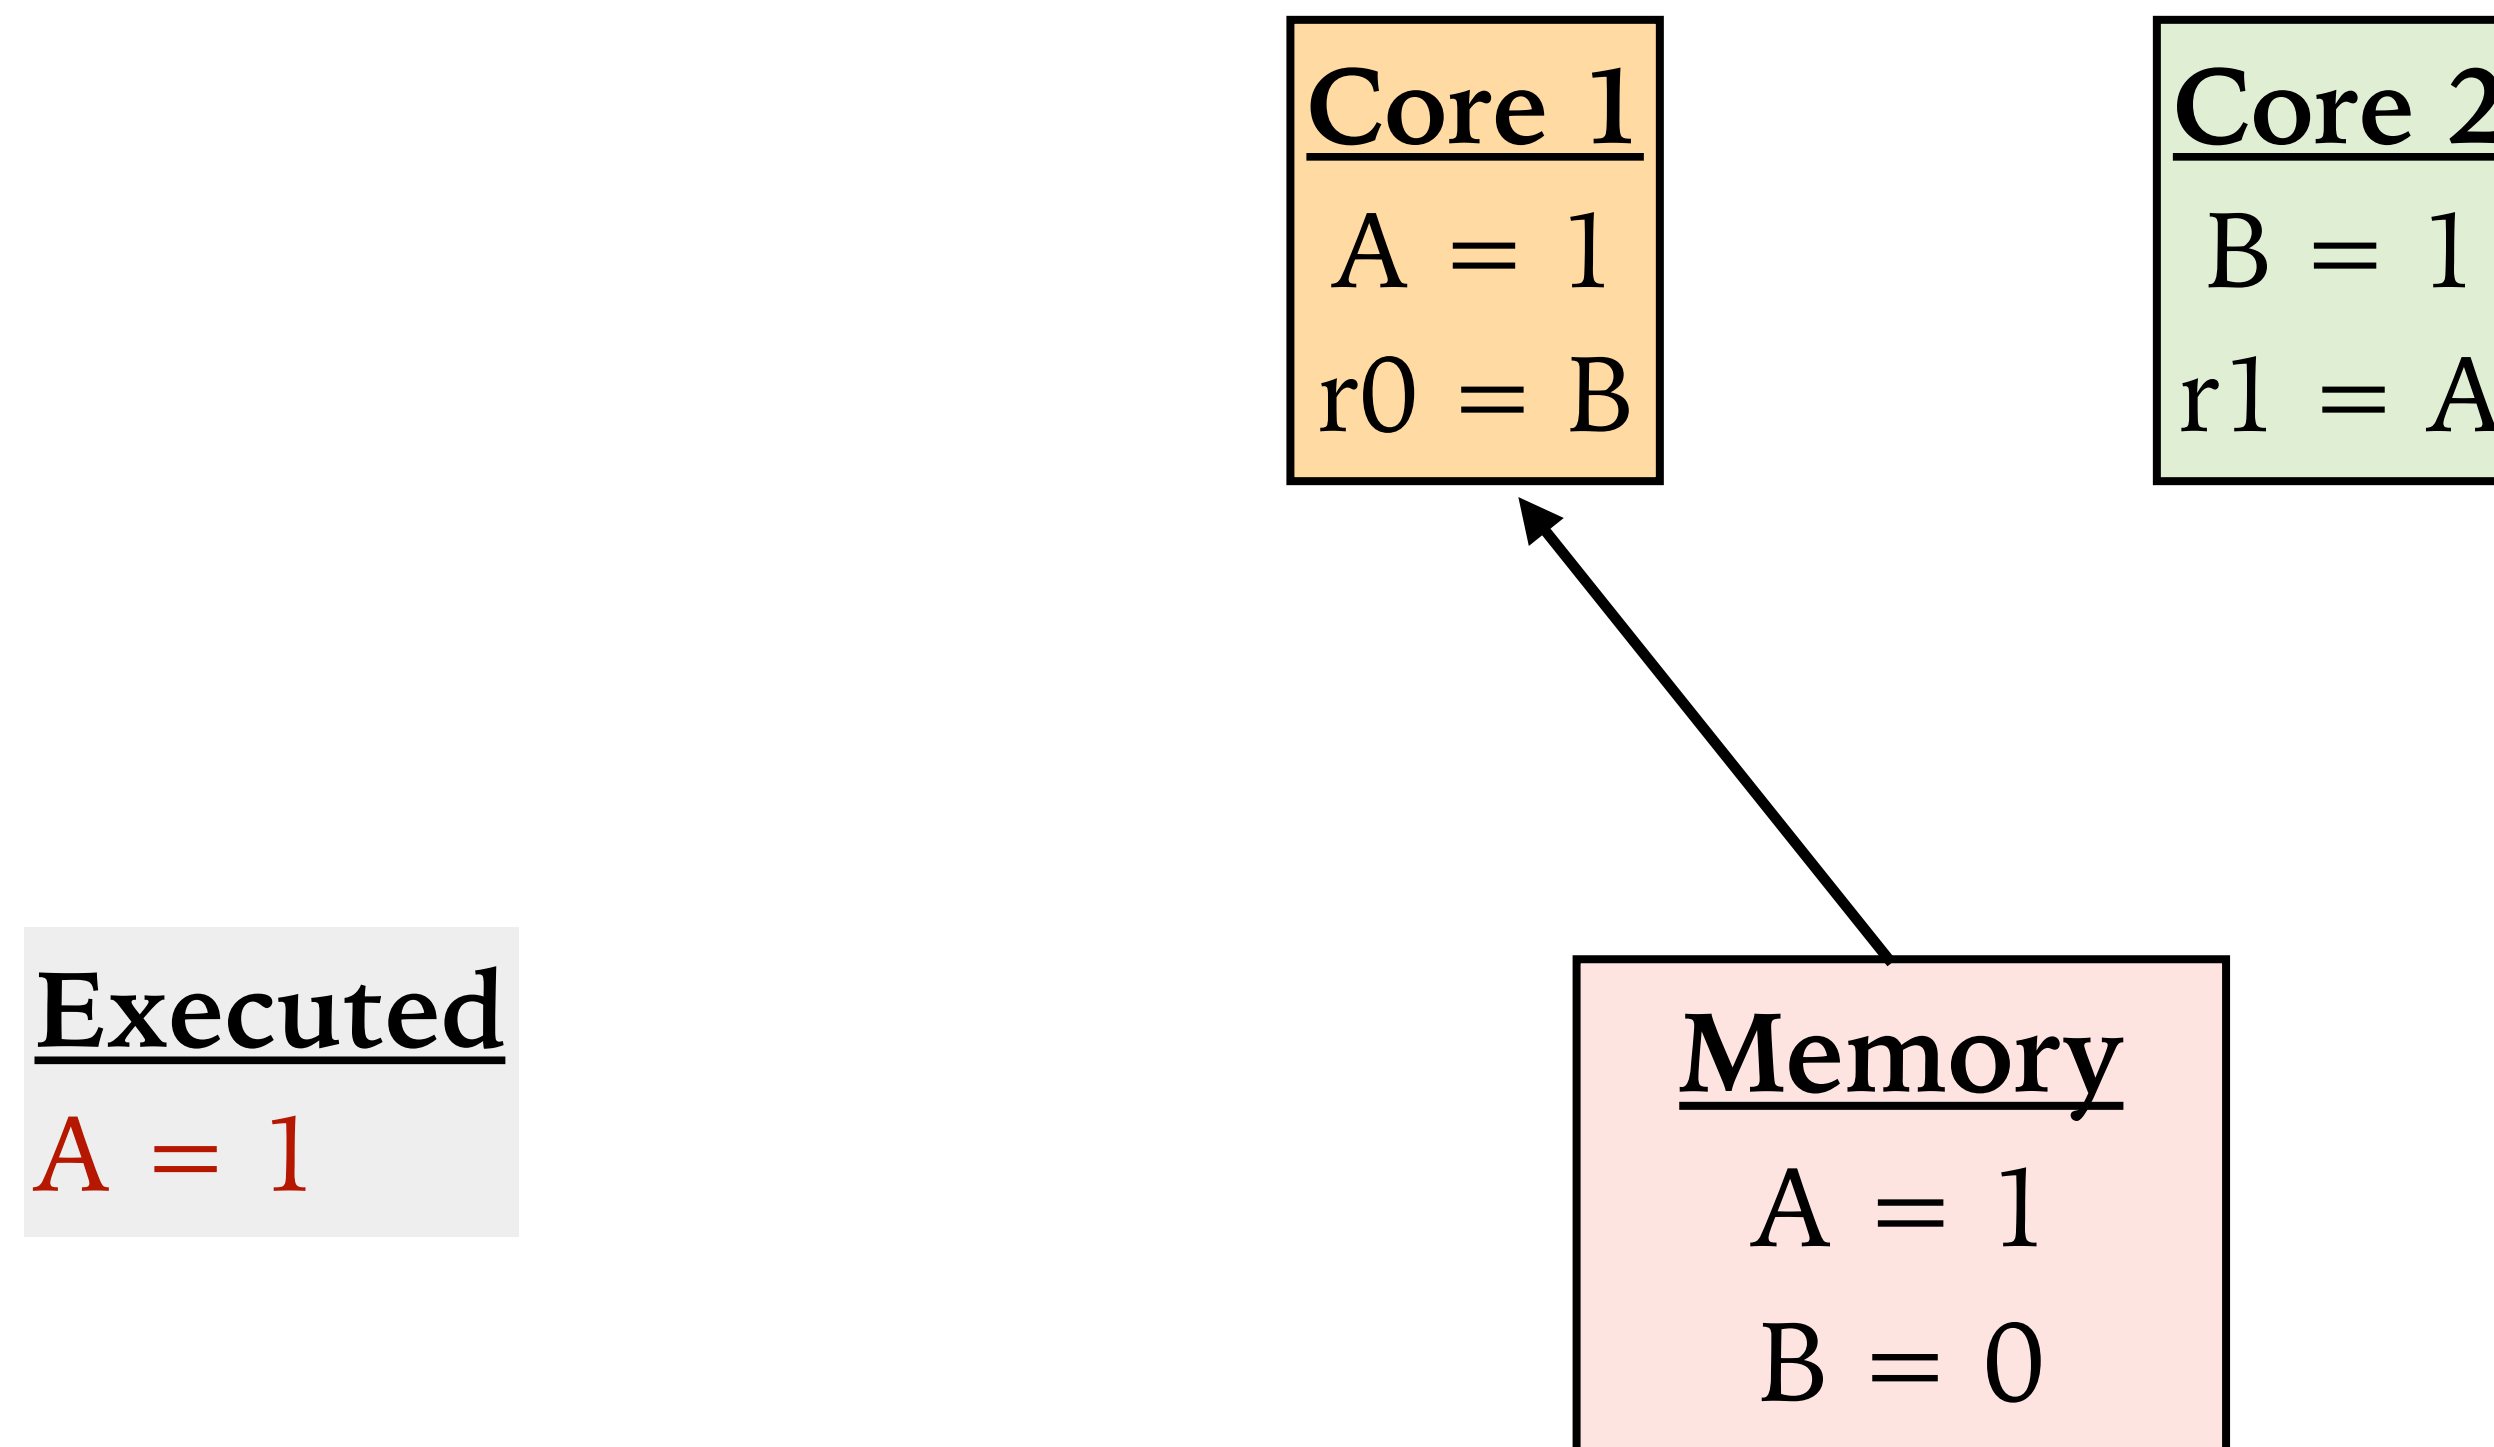
- Can be seen as a “switch” running one instruction at a time from different threads.



# Sequential Consistency

---

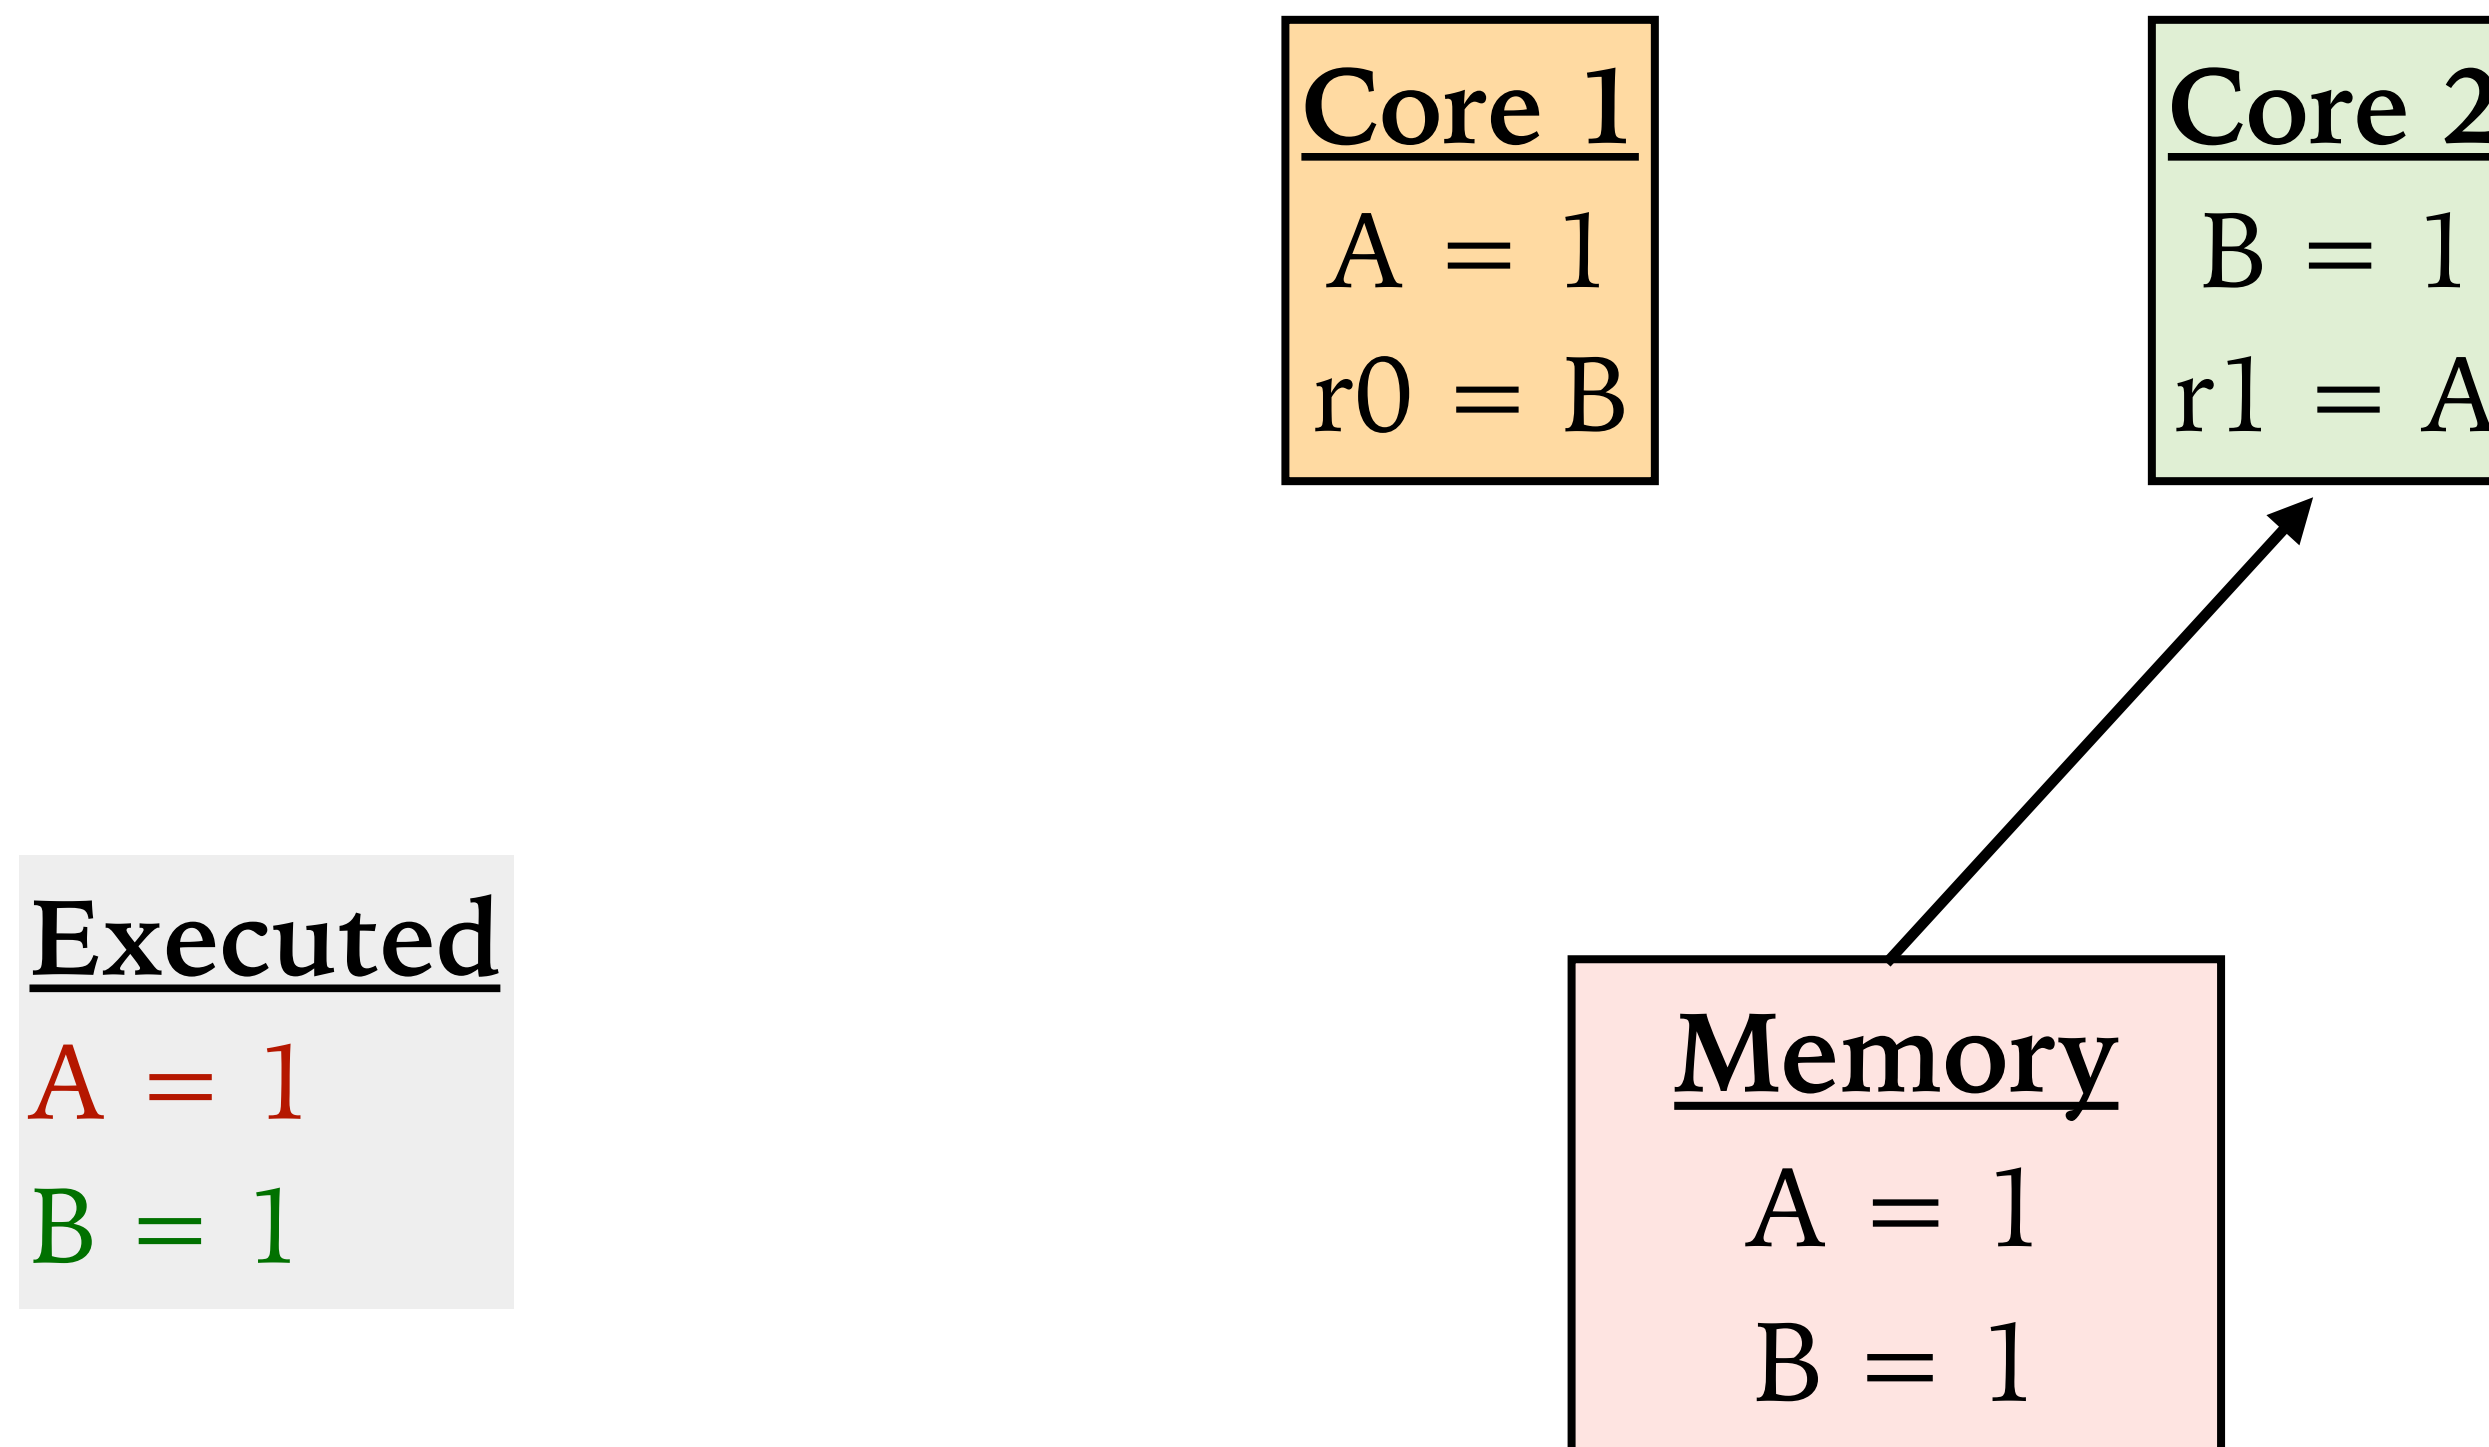
- Can be seen as a “switch” running one instruction at a time from different threads.



# Sequential Consistency

---

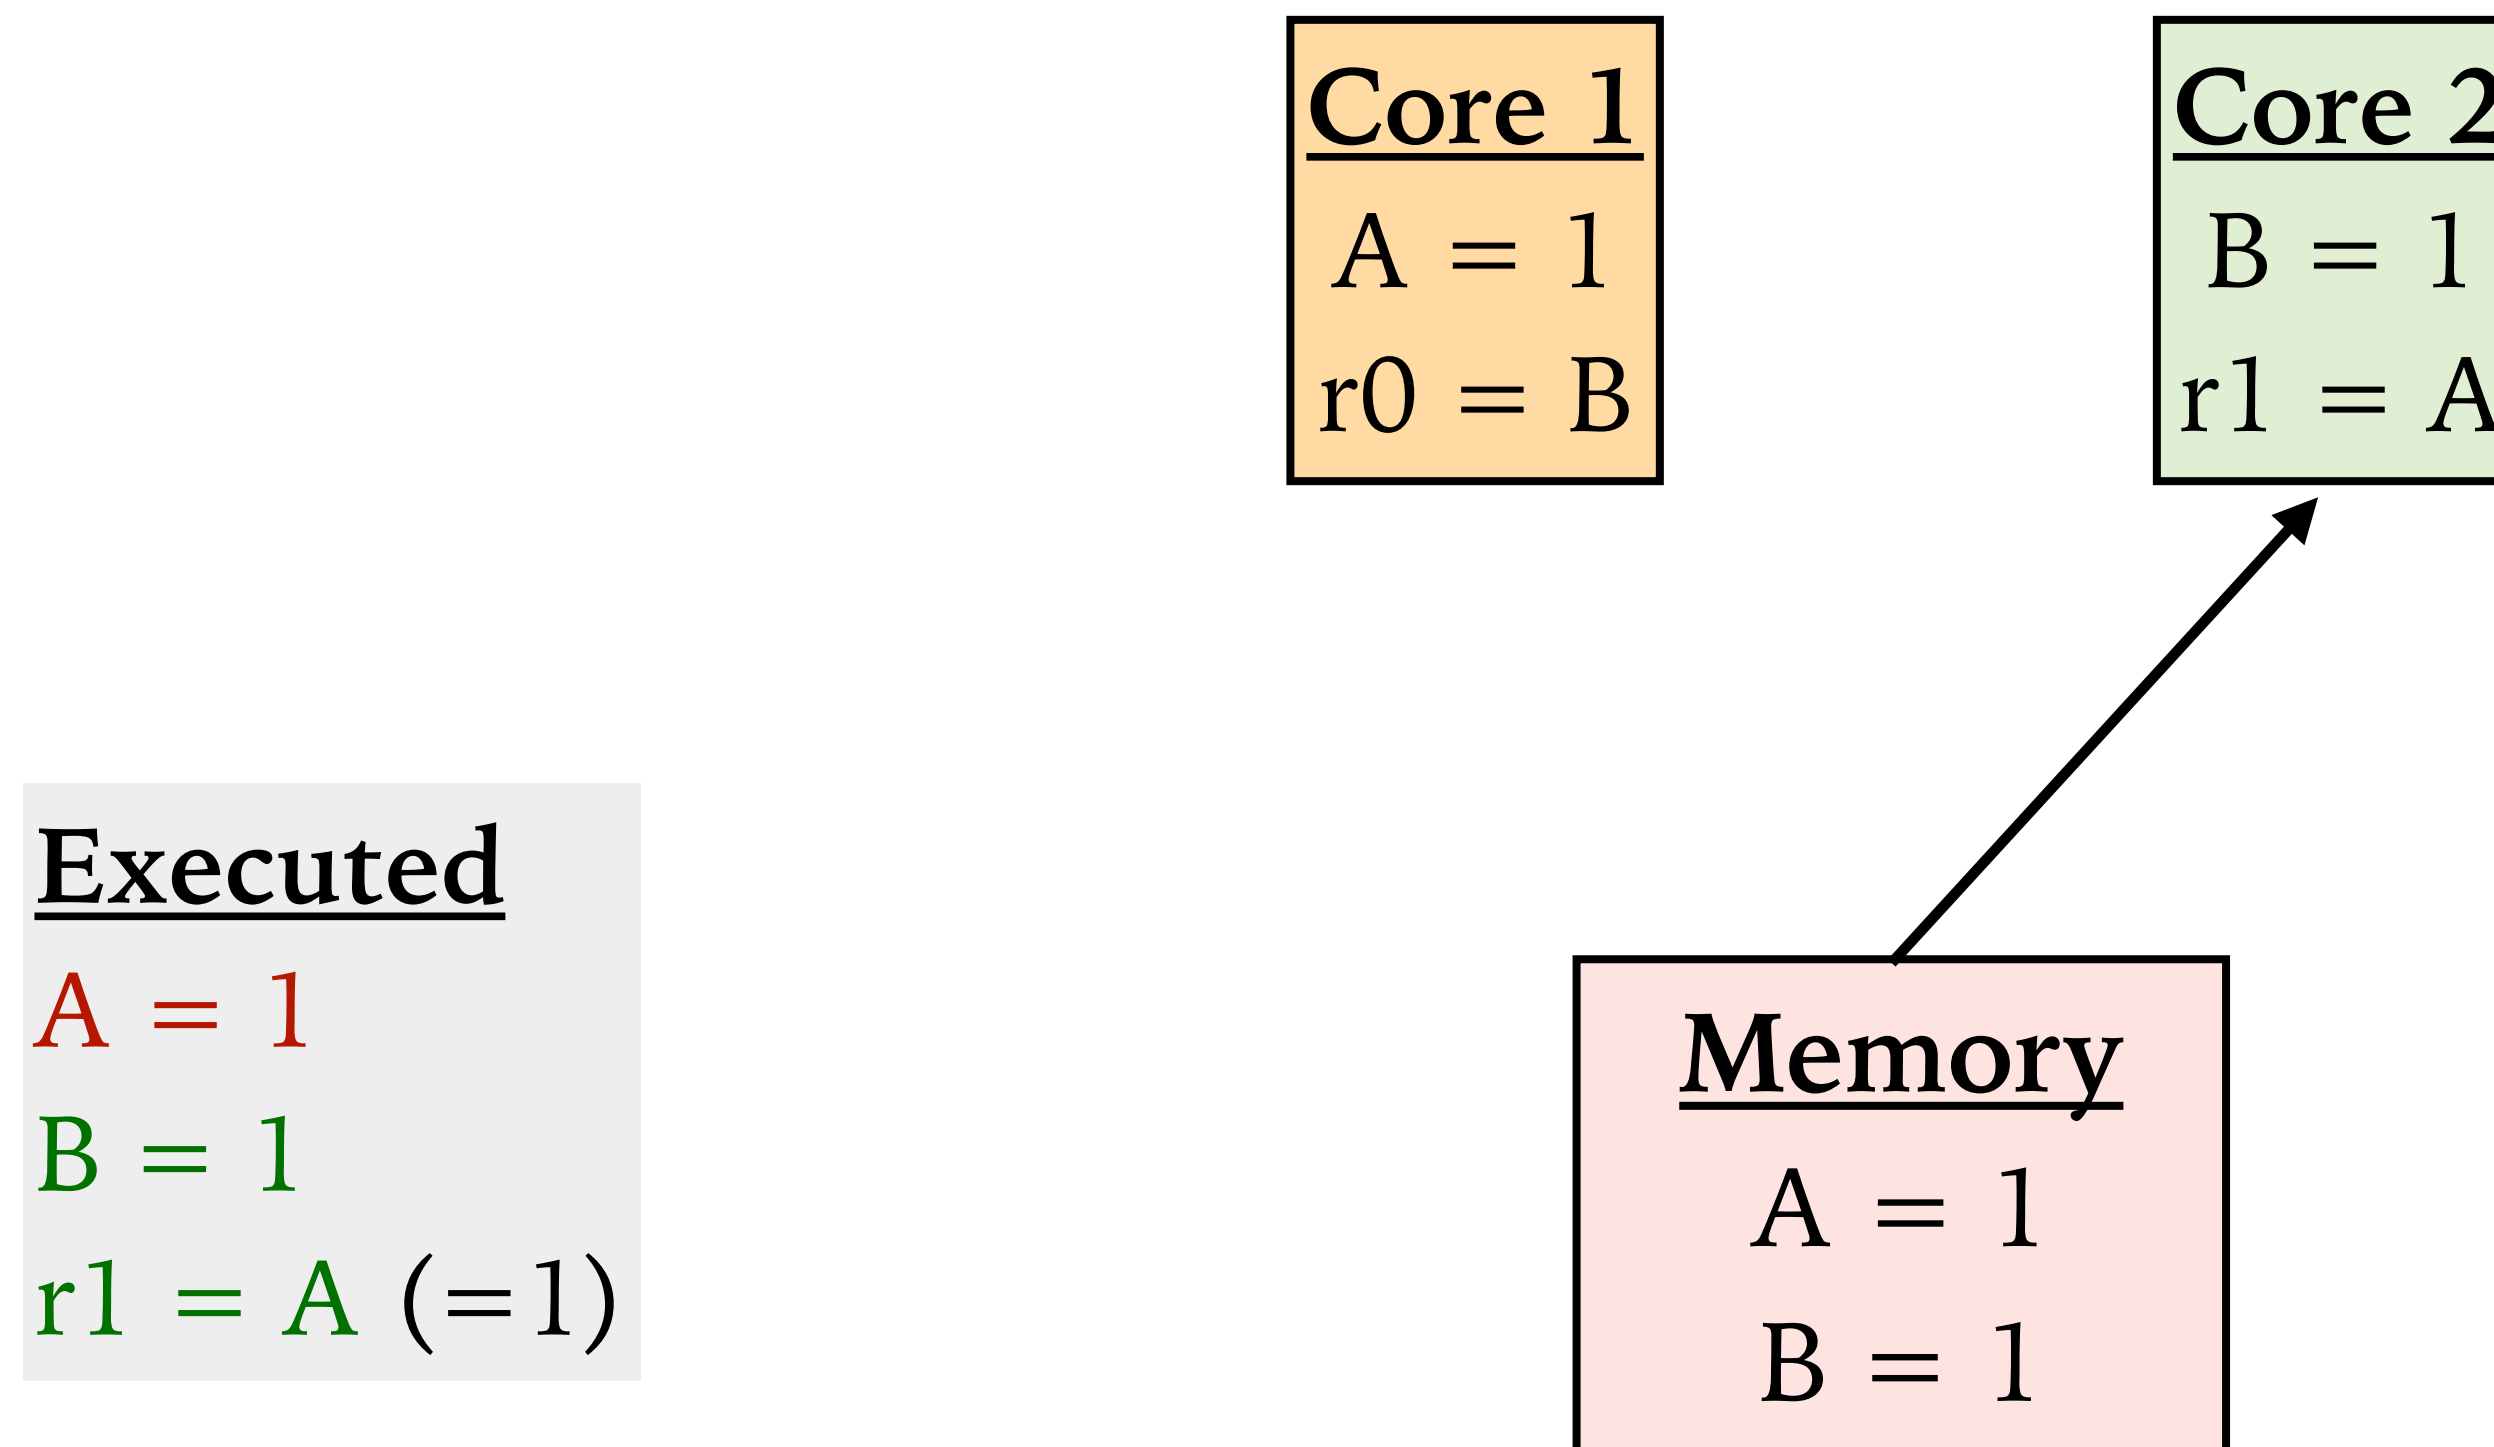
- Can be seen as a “switch” running one instruction at a time from different threads.



# Sequential Consistency

---

- Can be seen as a “switch” running one instruction at a time from different threads.

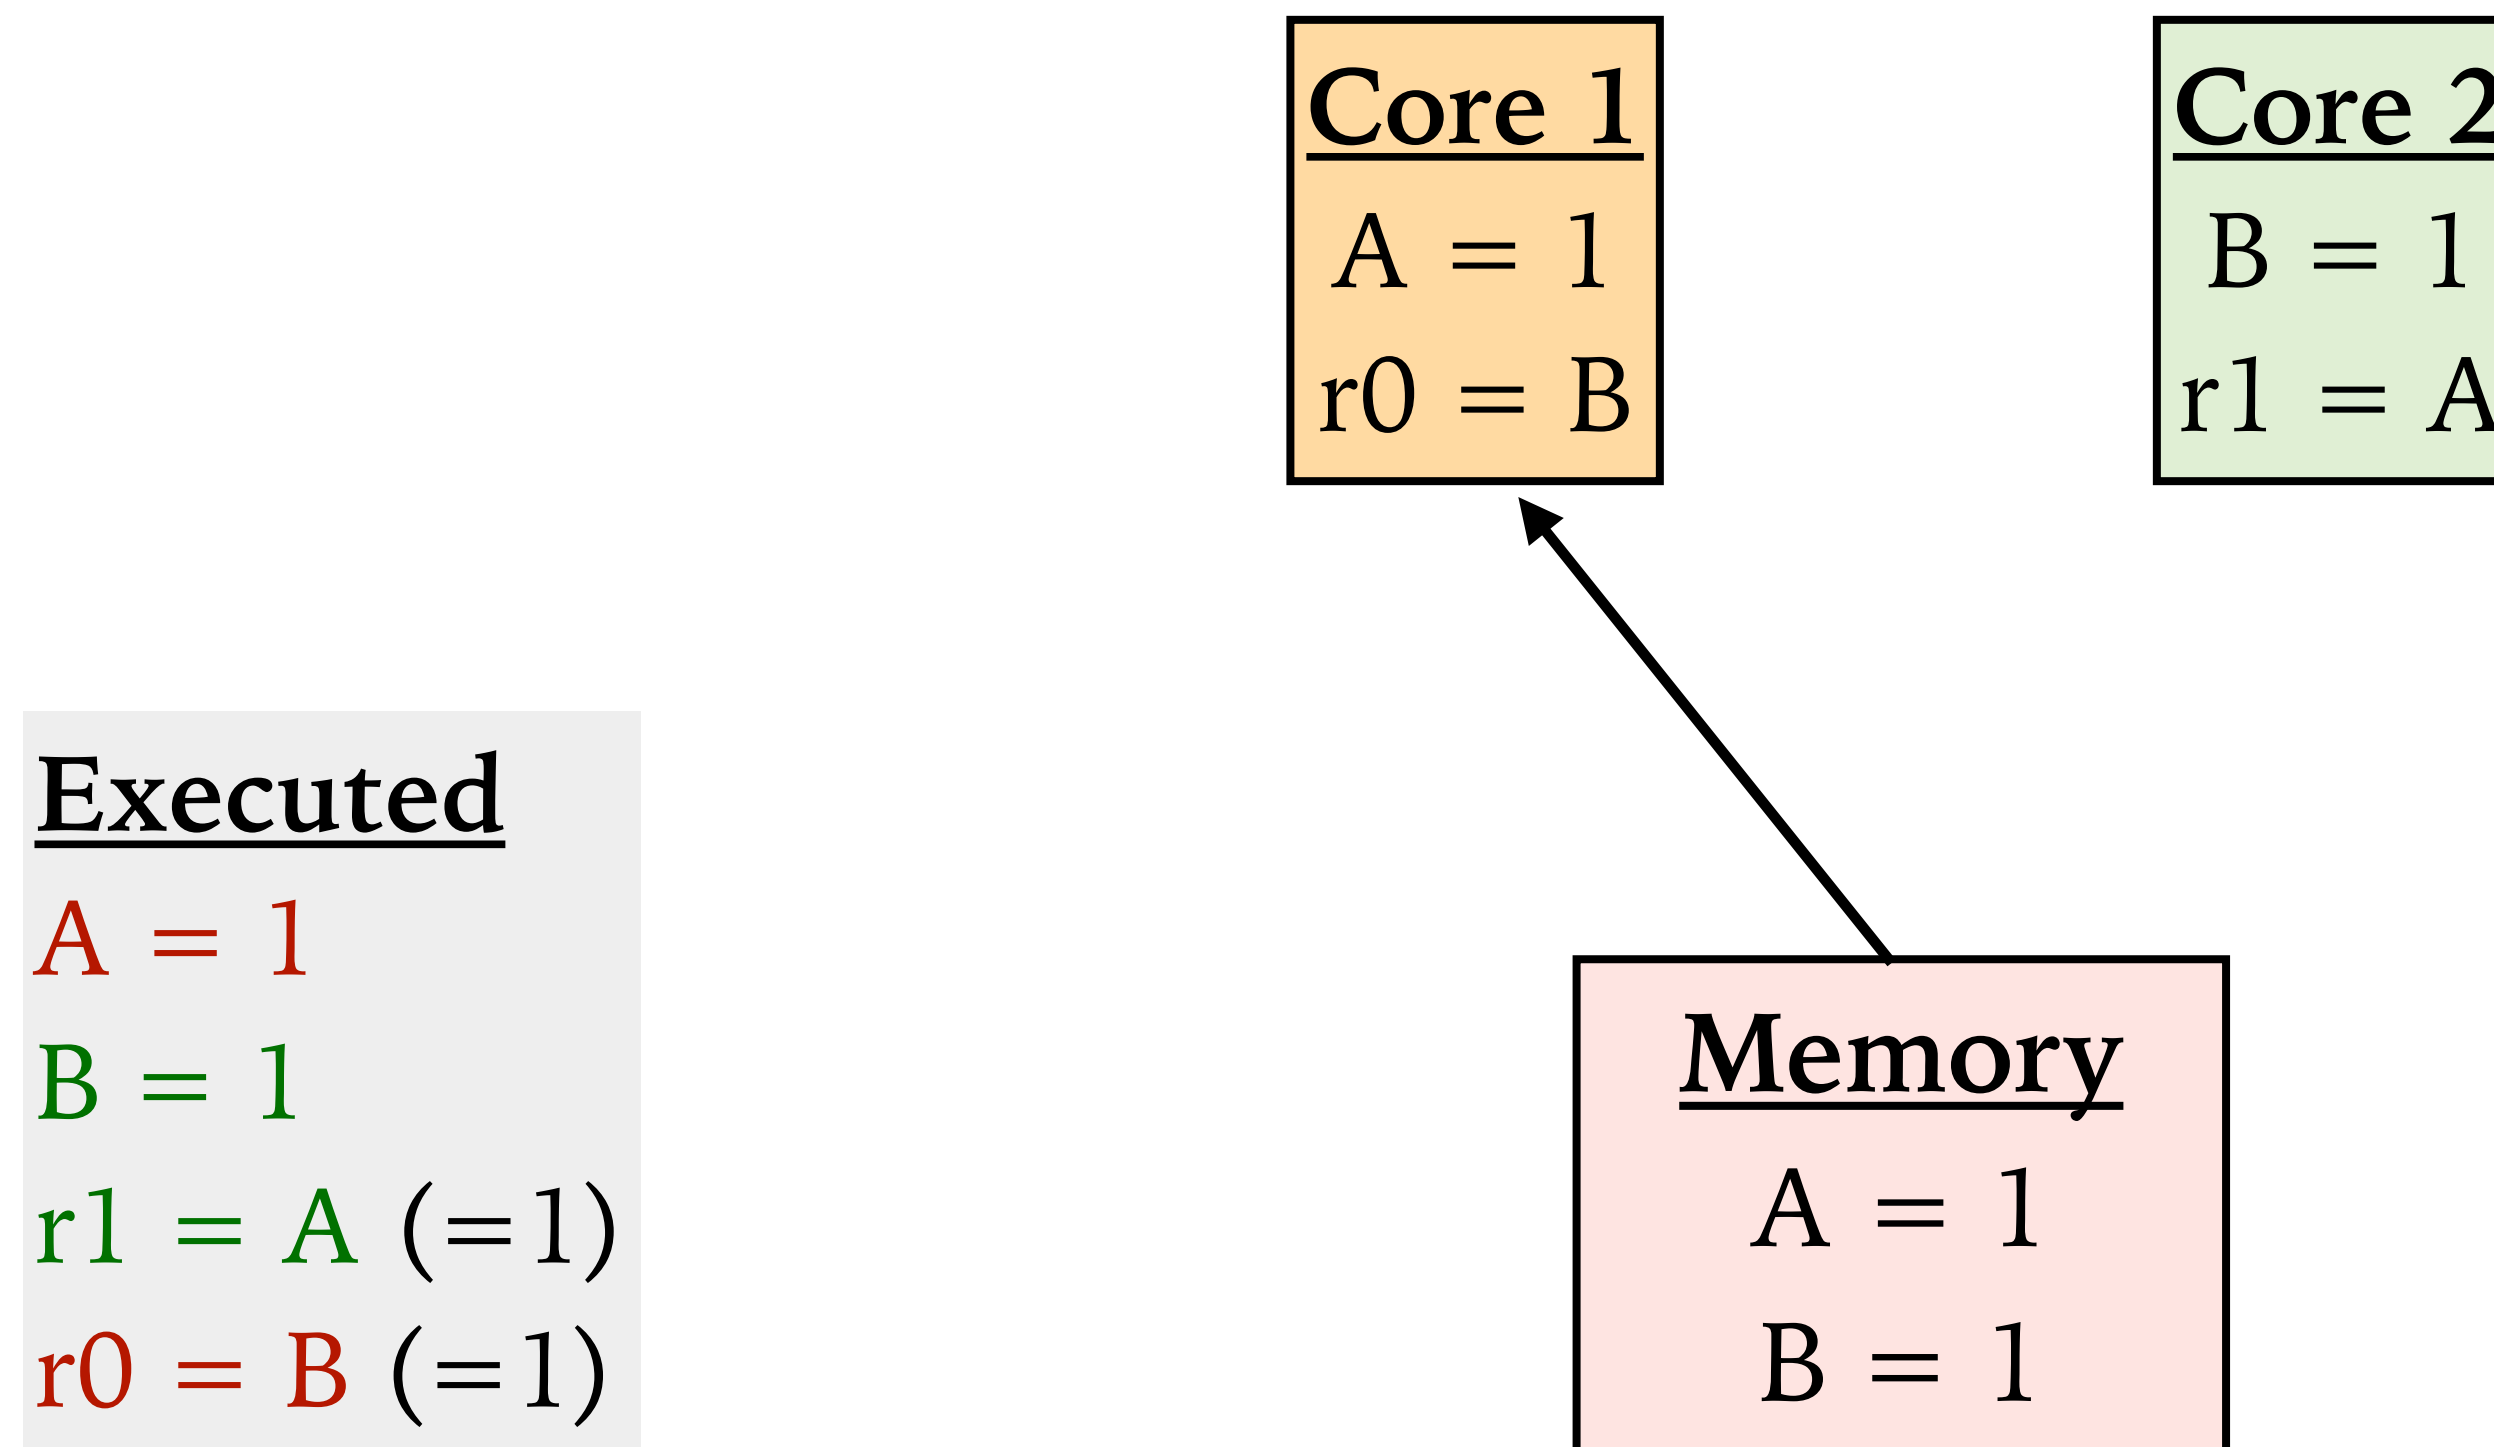




# Sequential Consistency

---

- Can be seen as a “switch” running one instruction at a time from different threads.



# Sequential Consistency

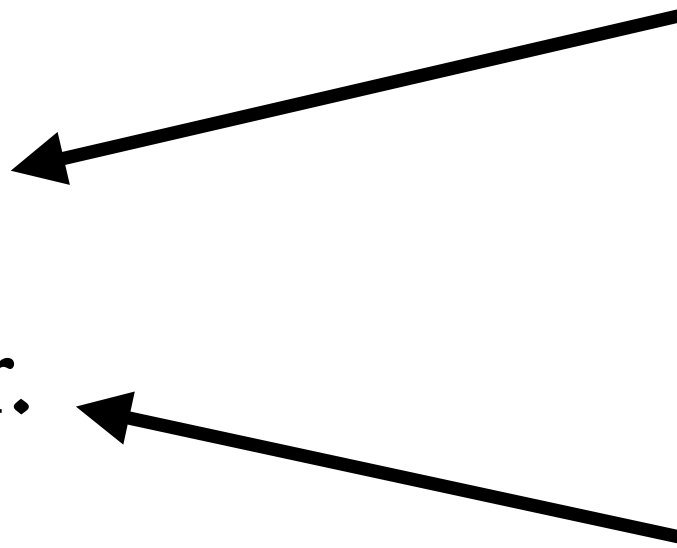
---

- Two invariants:

- All operations executed in some sequential order.

- Each thread's operations happen in program order.

Order across  
threads

Two arrows originate from the boxes on the right. One arrow points from the 'Order across threads' box to the first bullet point 'All operations executed in some sequential order.' The other arrow points from the 'Order within a thread' box to the second bullet point 'Each thread's operations happen in program order.'

Order within  
a thread

- Says nothing about **which** order all operations happen in

- Any interleaving of threads is allowed

- Developed by **Leslie Lamport** in 1979 (**Turing Award in 2013**).



# Memory (Consistency) Models

---

- A memory model defines the permitted reorderings of memory operations during execution.
- A **contract between hardware and software**: the hardware will only mess with your memory operations in these ways.
- Sequential consistency is the strongest memory (consistency) model; allows the fewest reorderings.



# Practice

---

- Assume sequential consistency, and all variables to be initially 0.

## Thread 1

(1)  $X = 1$   
(2)  $Y = 1$

## Thread 2

(3)  $r0 = Y$   
(4)  $r1 = X$

- |                           |   |
|---------------------------|---|
| ➤ $r0 = 0$ and $r1 = 0$ ? | (3) $\longrightarrow$ (4) $\longrightarrow$ (1) $\longrightarrow$ (2) |
| ➤ $r0 = 1$ and $r1 = 1$ ? | (1) $\longrightarrow$ (2) $\longrightarrow$ (3) $\longrightarrow$ (4) |
| ➤ $r0 = 0$ and $r1 = 1$ ? | (1) $\longrightarrow$ (3) $\longrightarrow$ (4) $\longrightarrow$ (2) |
| ➤ $r0 = 1$ and $r1 = 0$ ? | Not possible.   |



# Why SC?

---

- Agrees with programmer intuition!

# Why not SC?

---

- Extremely slow to guarantee in hardware!
  - The “switch” model is very conservative and does not allow several valid reorderings.



# The Problem with SC

---

These instructions don't conflict;  
we don't really need to  
wait for the first to finish.

Thread 1

A = 1  
r0 = B

Thread 2

B = 1  
r1 = A

Executed

A = 1

Memory

A = 1  
B = 0

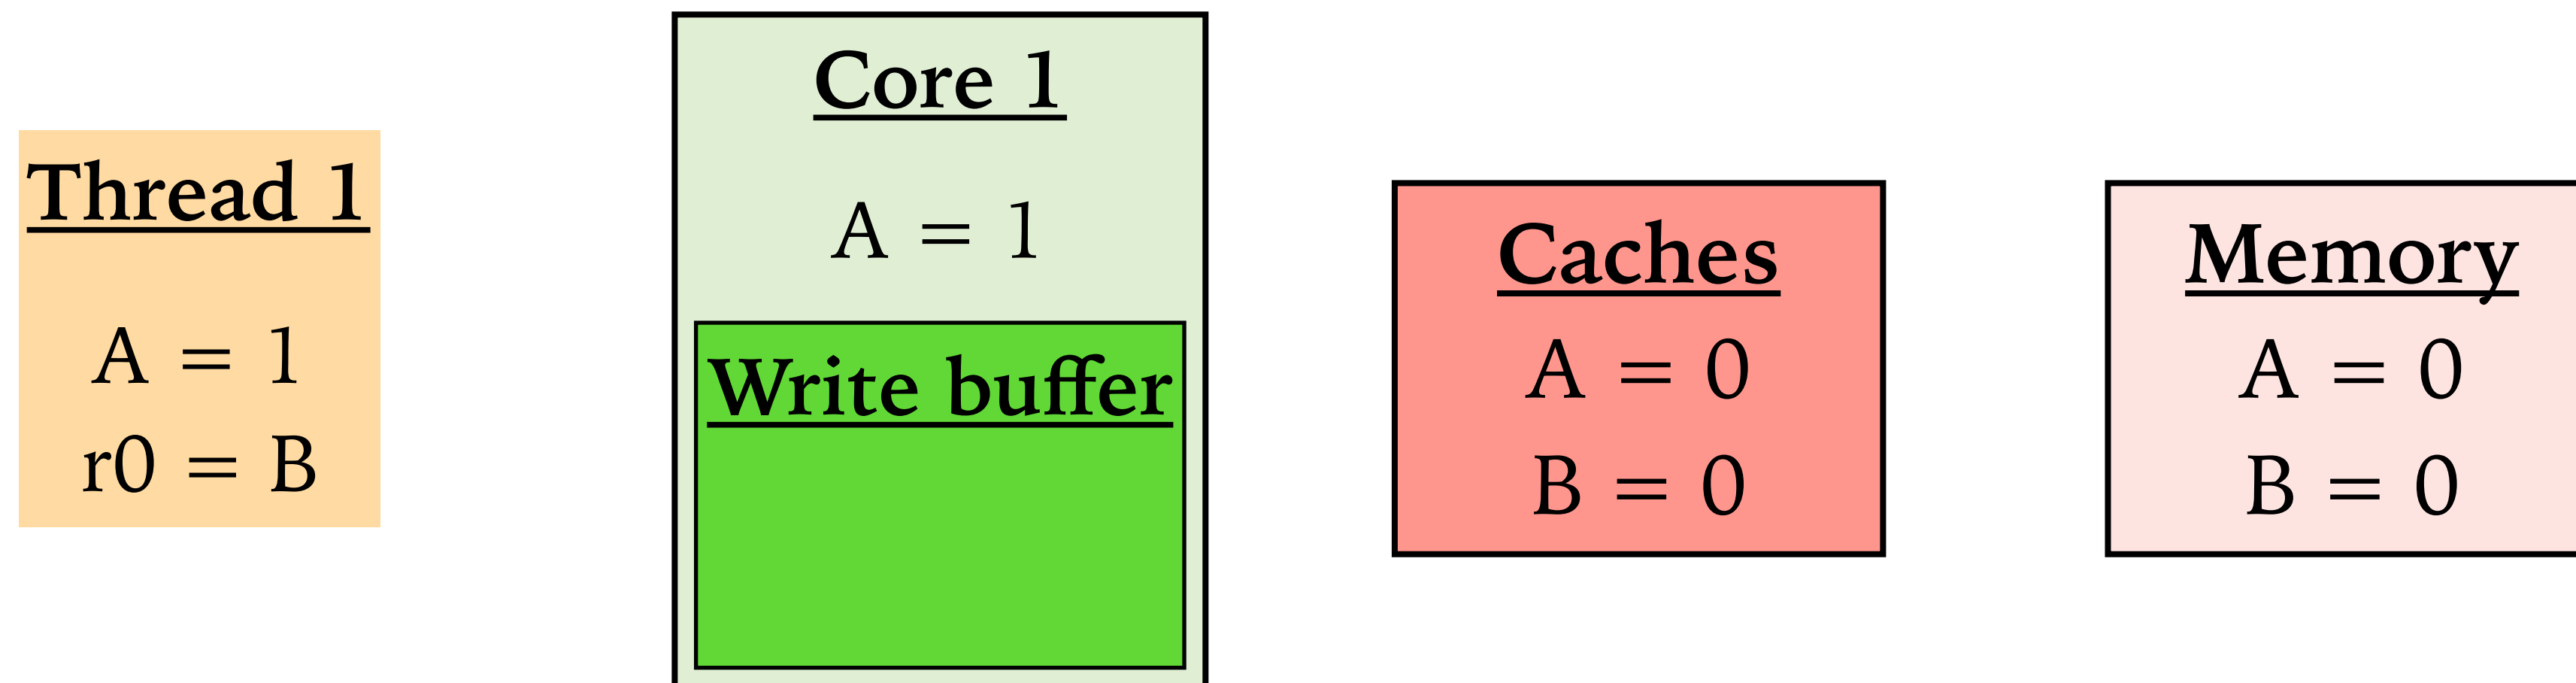
Also, writing to memory  
takes a lot of time.



# Hardware-Level Optimization: Write Buffers

---

- Store writes in a processor-local buffer and proceed to next instruction immediately.
- The cache will pull out writes out of the write buffer when it's ready.

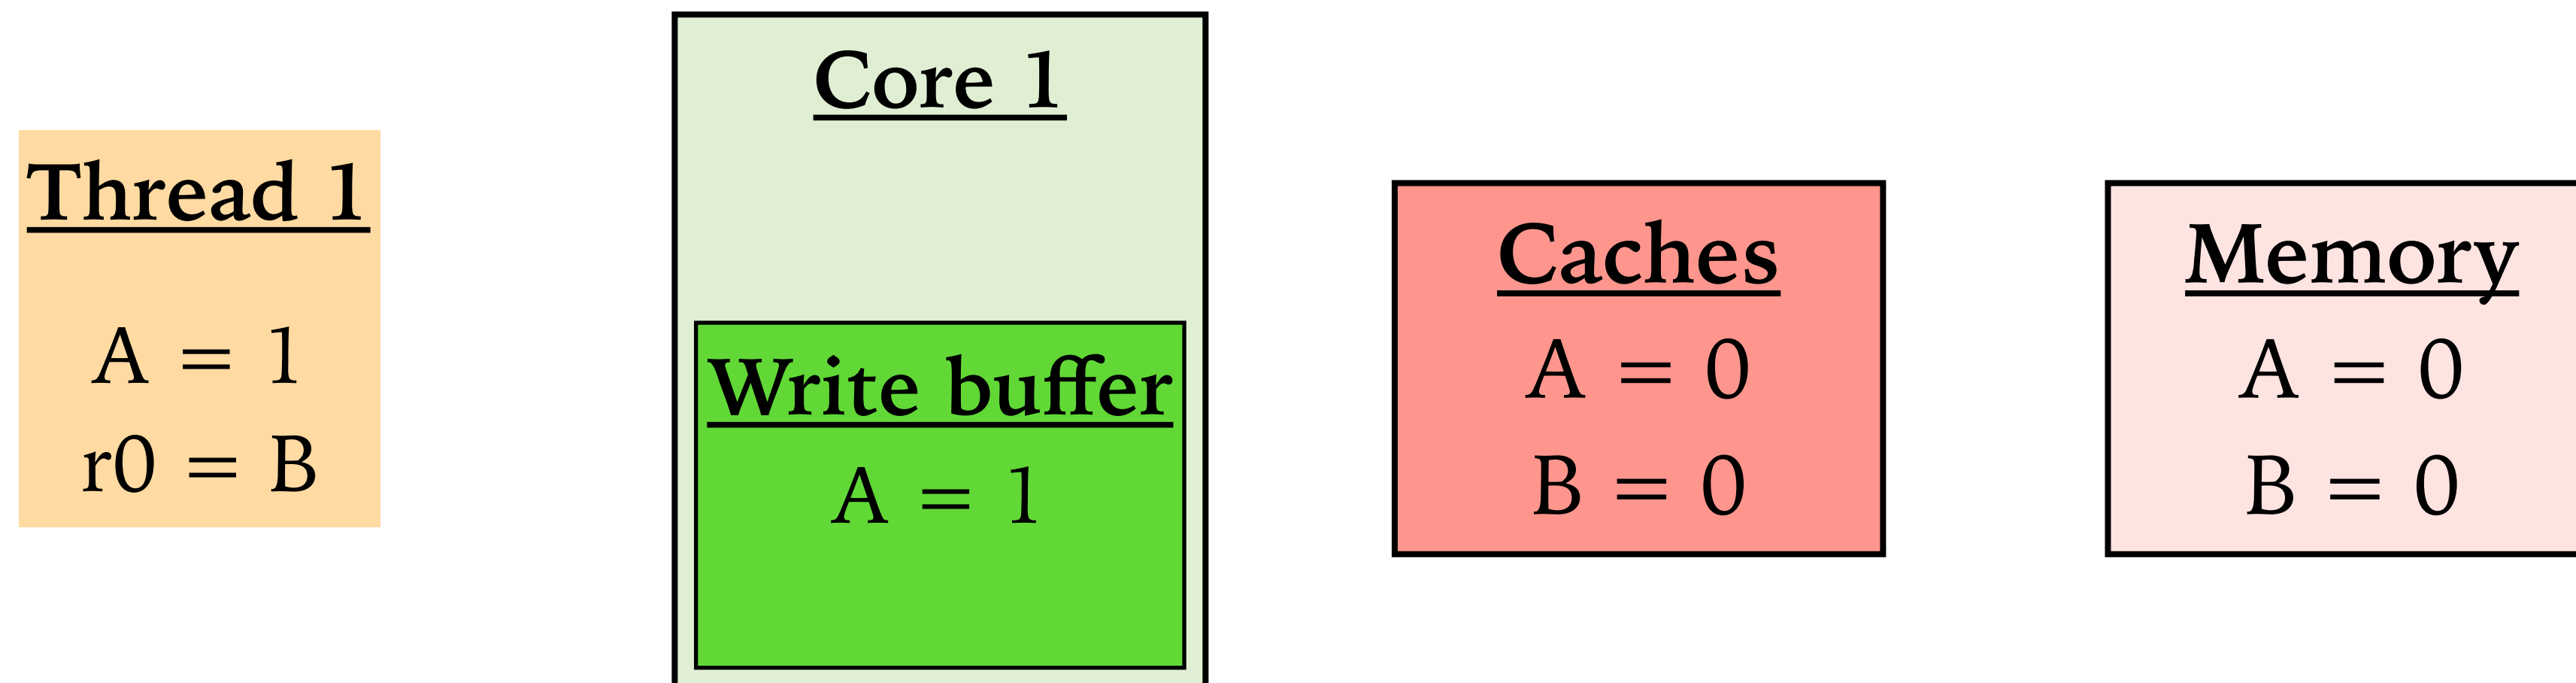




# Hardware-Level Optimization: Write Buffers

---

- Store writes in a processor-local buffer and proceed to next instruction immediately.
- The cache will pull out writes out of the write buffer when it's ready.

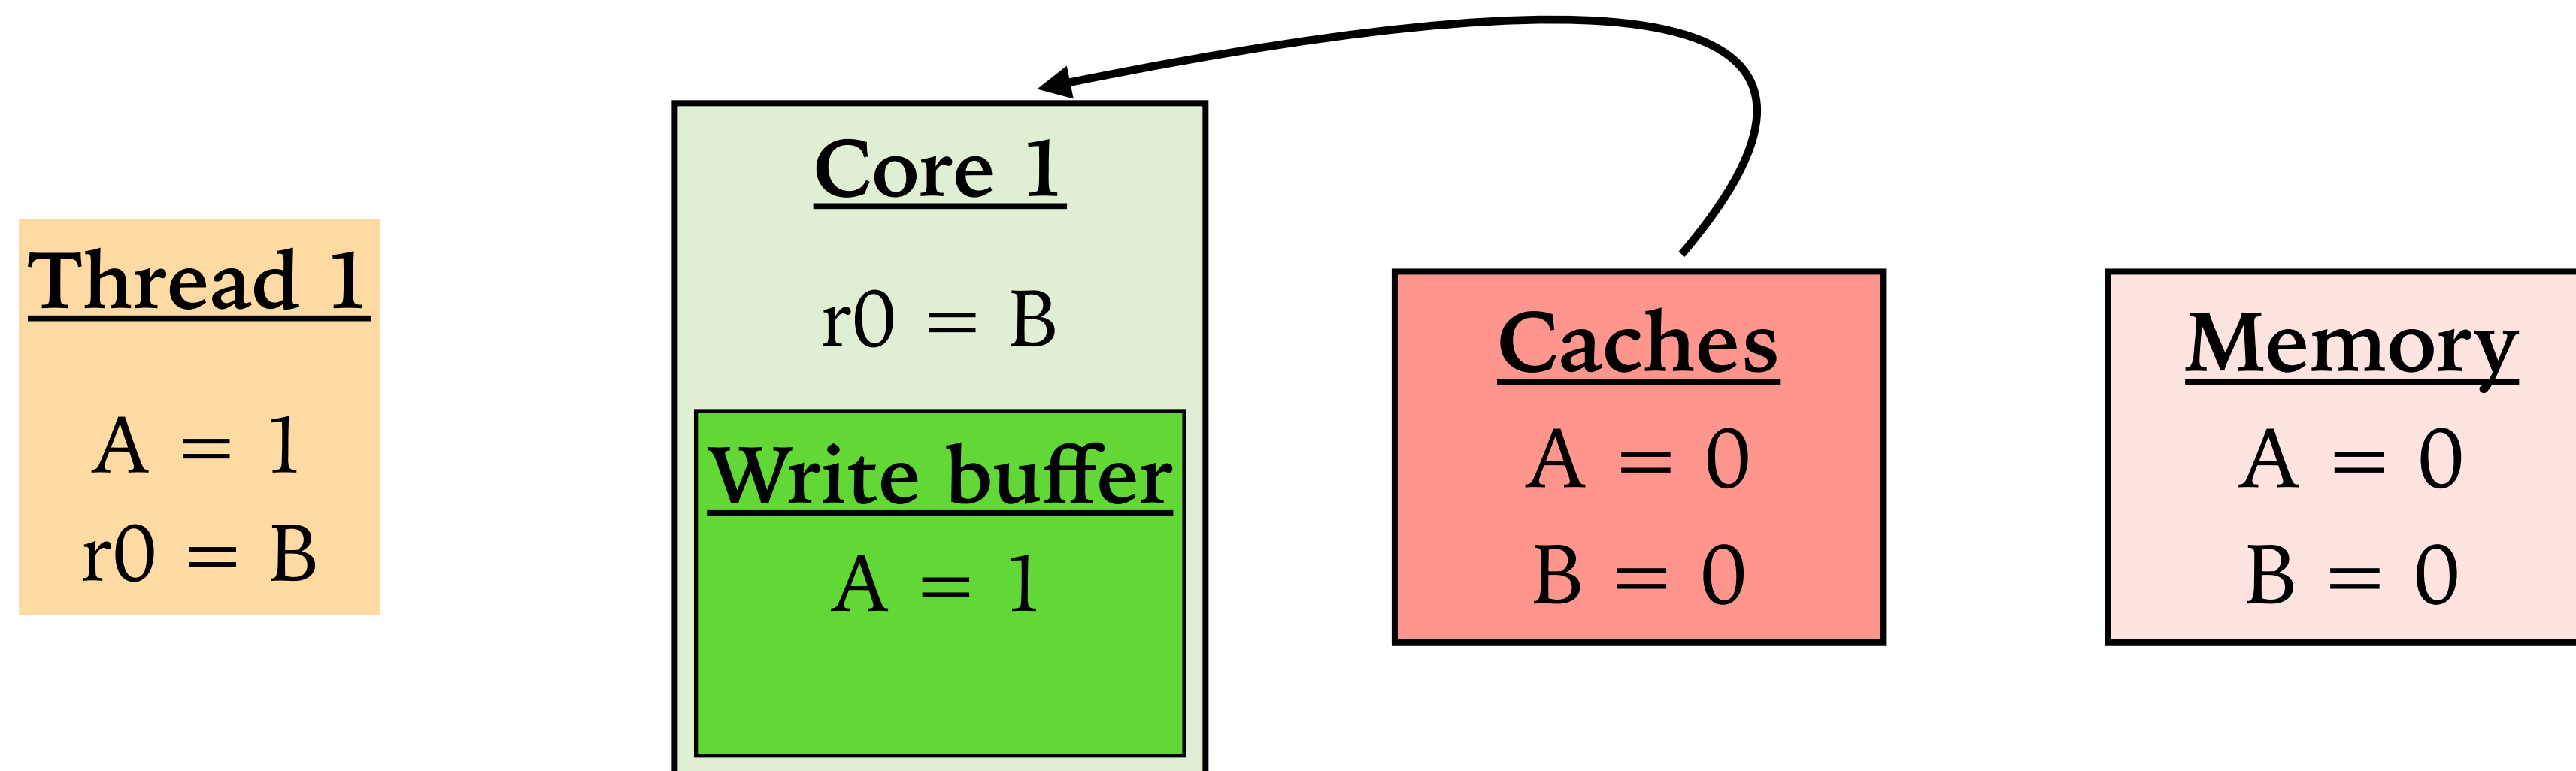




# Hardware-Level Optimization: Write Buffers

---

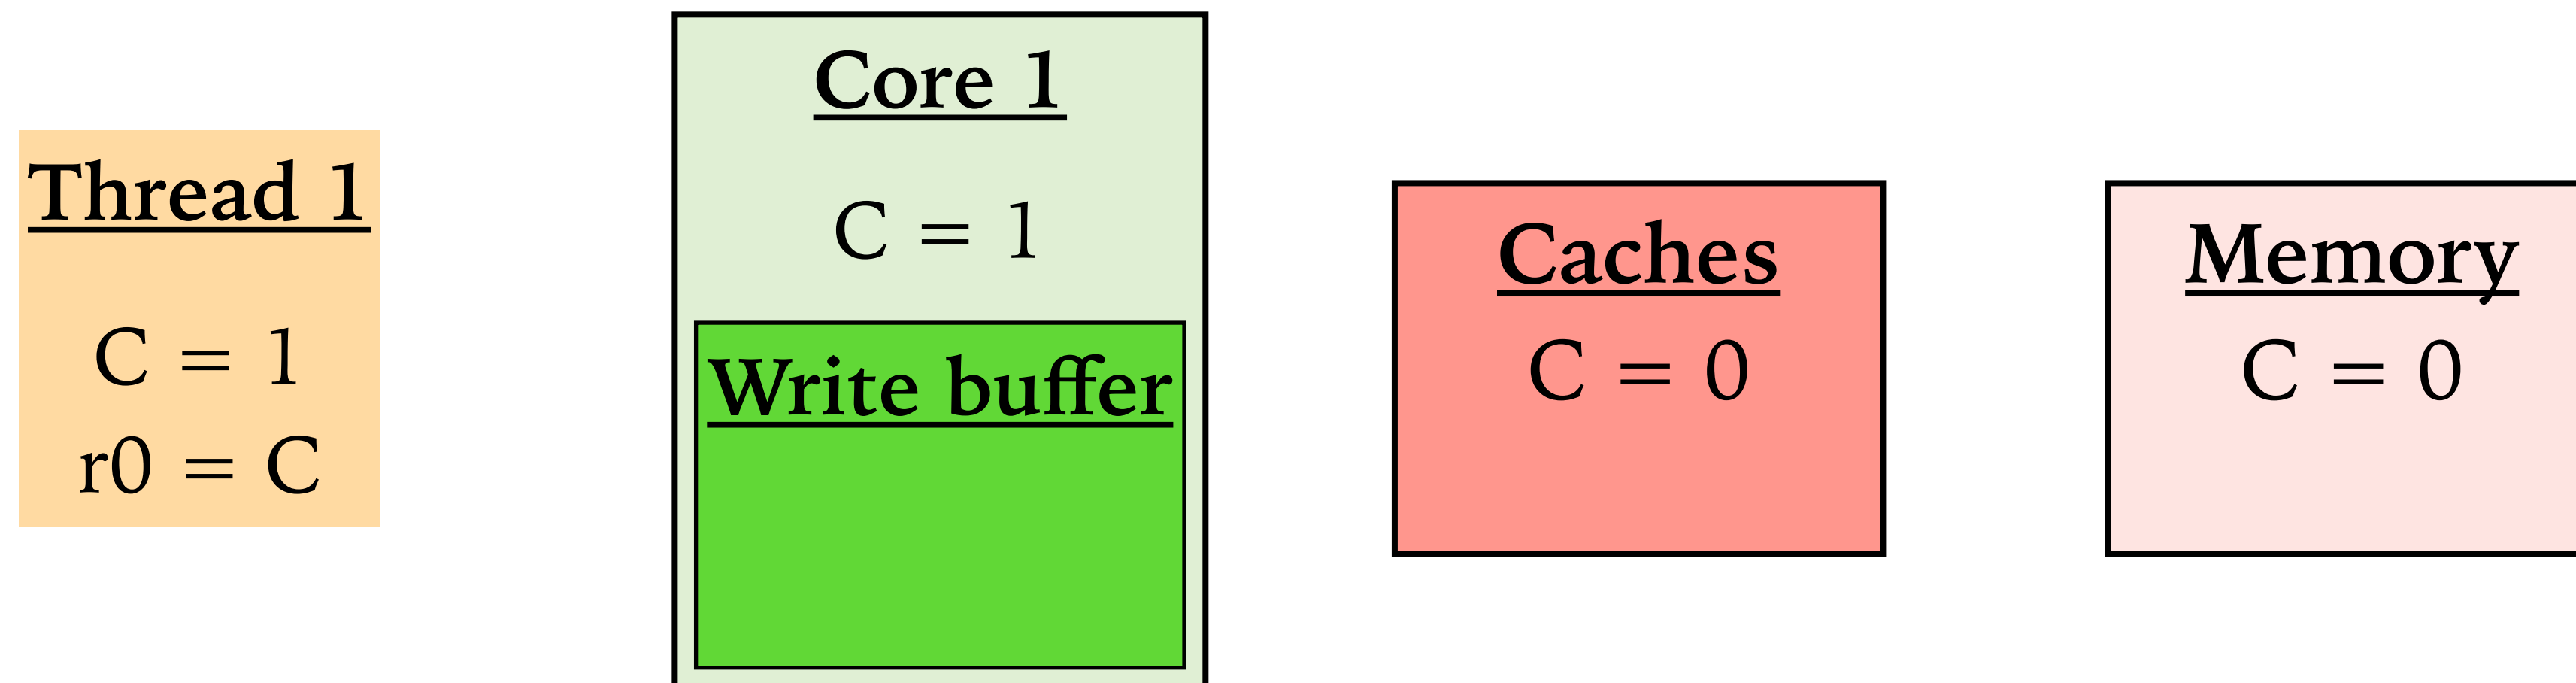
- Store writes in a processor-local buffer and proceed to next instruction immediately.
- The cache will pull out writes out of the write buffer when it's ready.



# Hardware-Level Optimization: Write Buffers

---

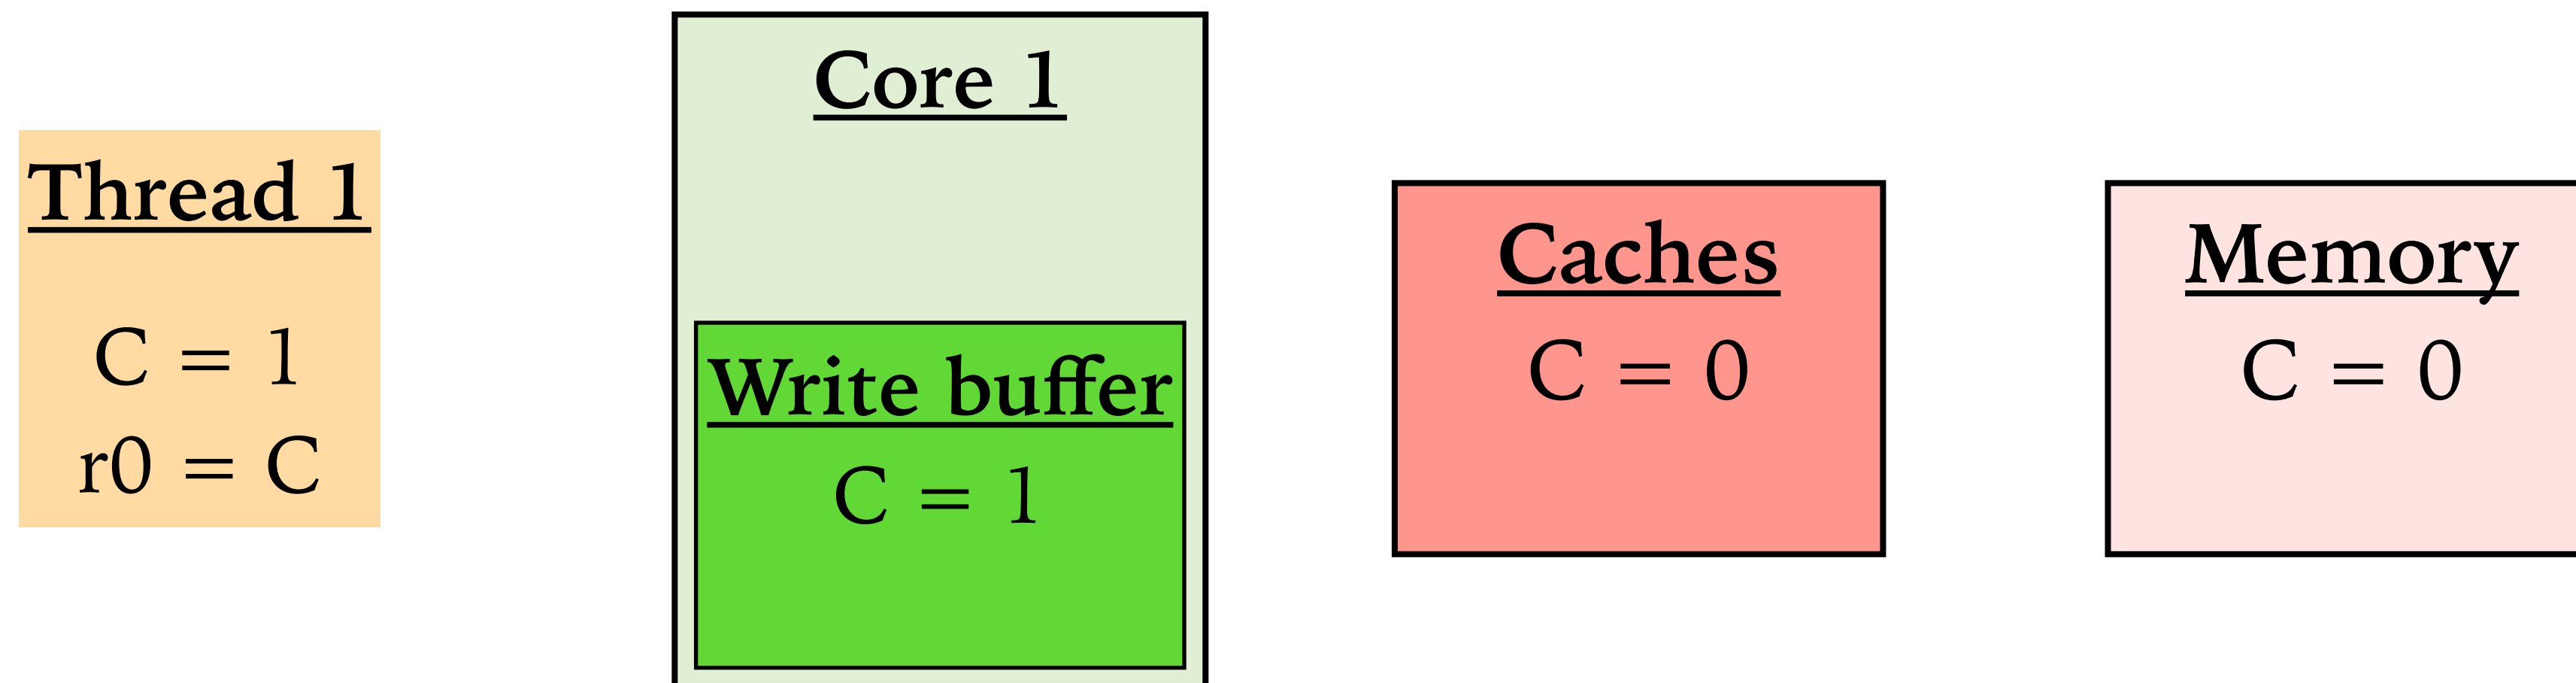
- Store writes in a processor-local buffer and proceed to next instruction immediately.
- The cache will pull out writes out of the write buffer when it's ready.



# Hardware-Level Optimization: Write Buffers

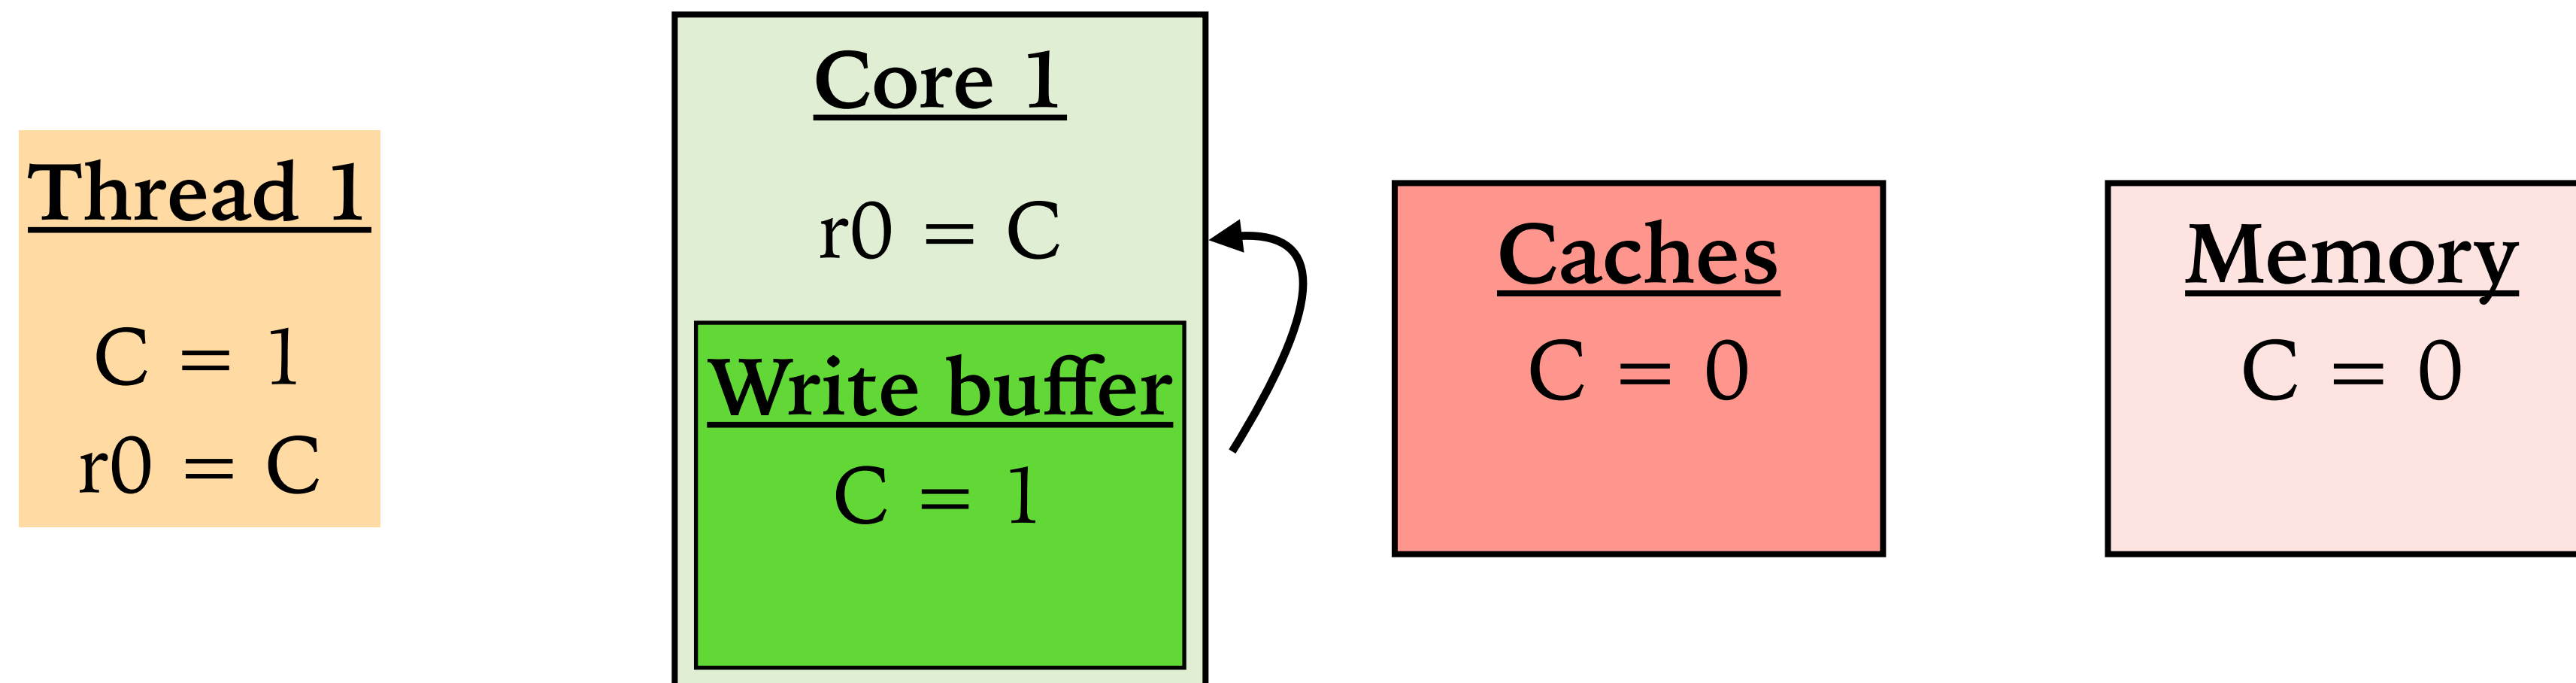
---

- Store writes in a processor-local buffer and proceed to next instruction immediately.
- The cache will pull out writes out of the write buffer when it's ready.



# Hardware-Level Optimization: Write Buffers

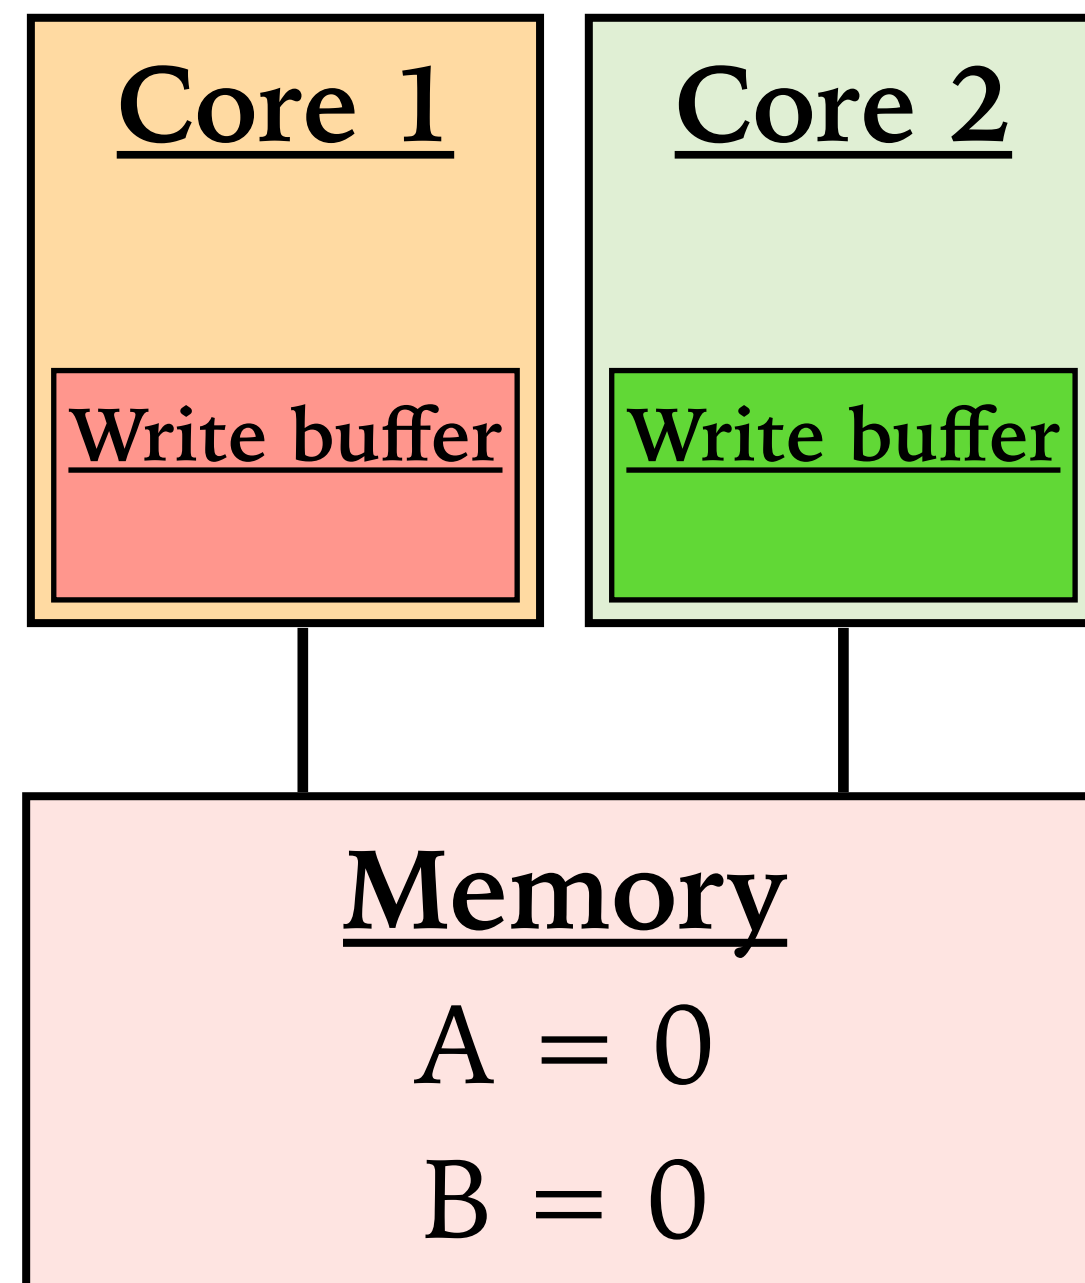
- Store writes in a processor-local buffer and proceed to next instruction immediately.
- The cache will pull out writes out of the write buffer when it's ready.



- Write buffers define a *new memory model*!

# Write buffers change memory behaviour

---



## Thread 1

(1)  $A = 1$   
(2)  $r0 = B$

## Thread 2


(3)  $B = 1$   
(4)  $r1 = A$

➤  $r0 = 0?$  and  $r1 = 0?$

➤ SC? No.

# Write buffers change memory behaviour

---



## Thread 1

(1)  $A = 1$   
(2)  $r0 = B$

## Thread 2

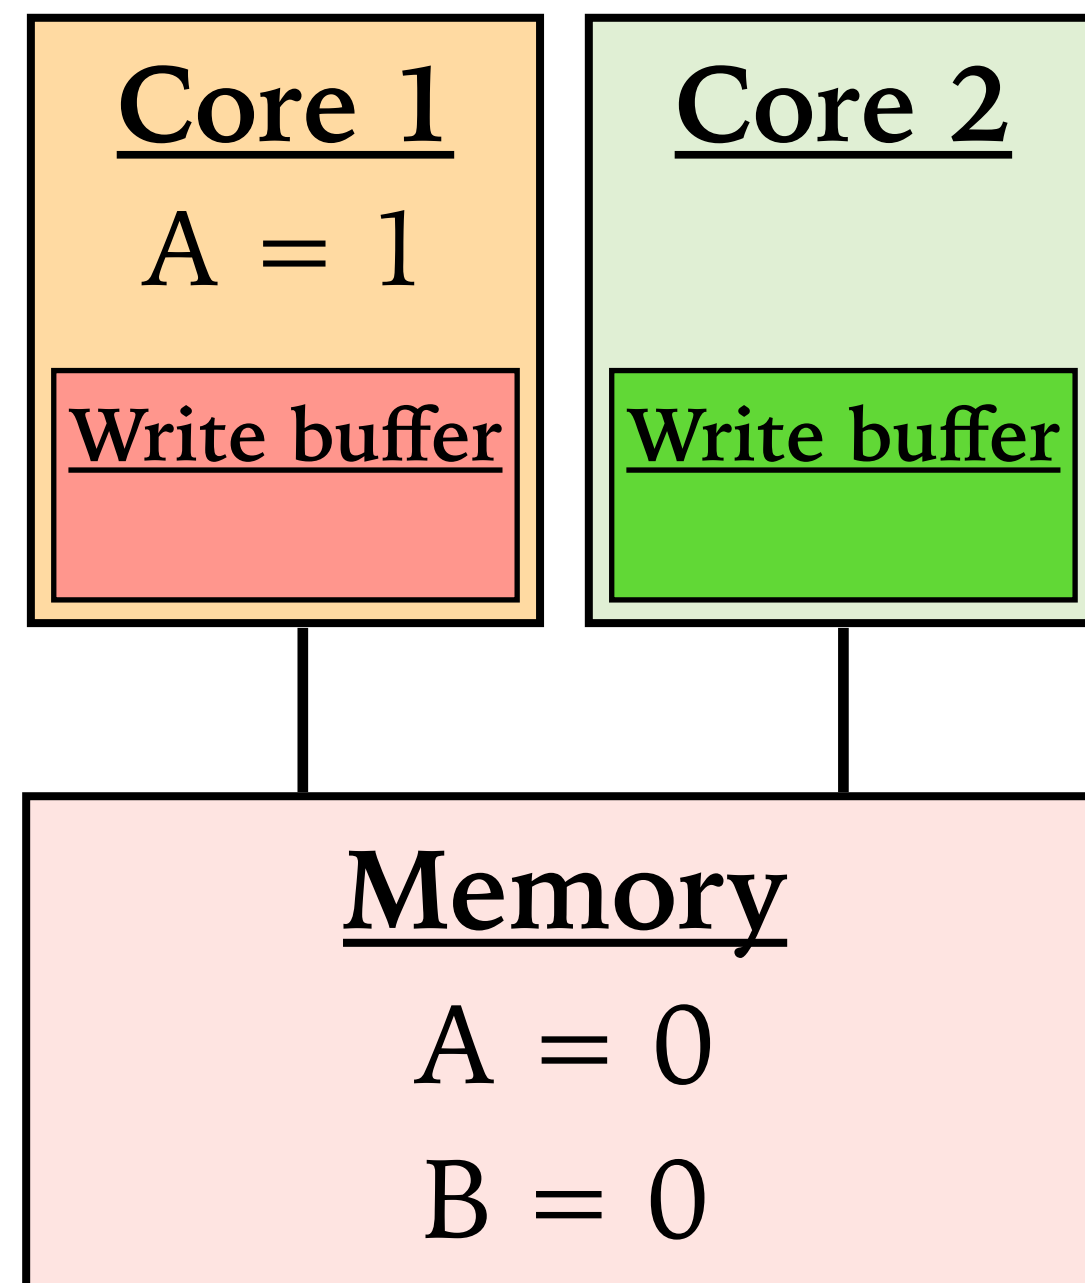
(3)  $B = 1$   
(4)  $r1 = A$

➤  $r0 = 0?$  and  $r1 = 0?$

➤ SC? No.

Executed

# Write buffers change memory behaviour



## Thread 1

(1)  $A = 1$   
(2)  $r0 = B$

## Thread 2

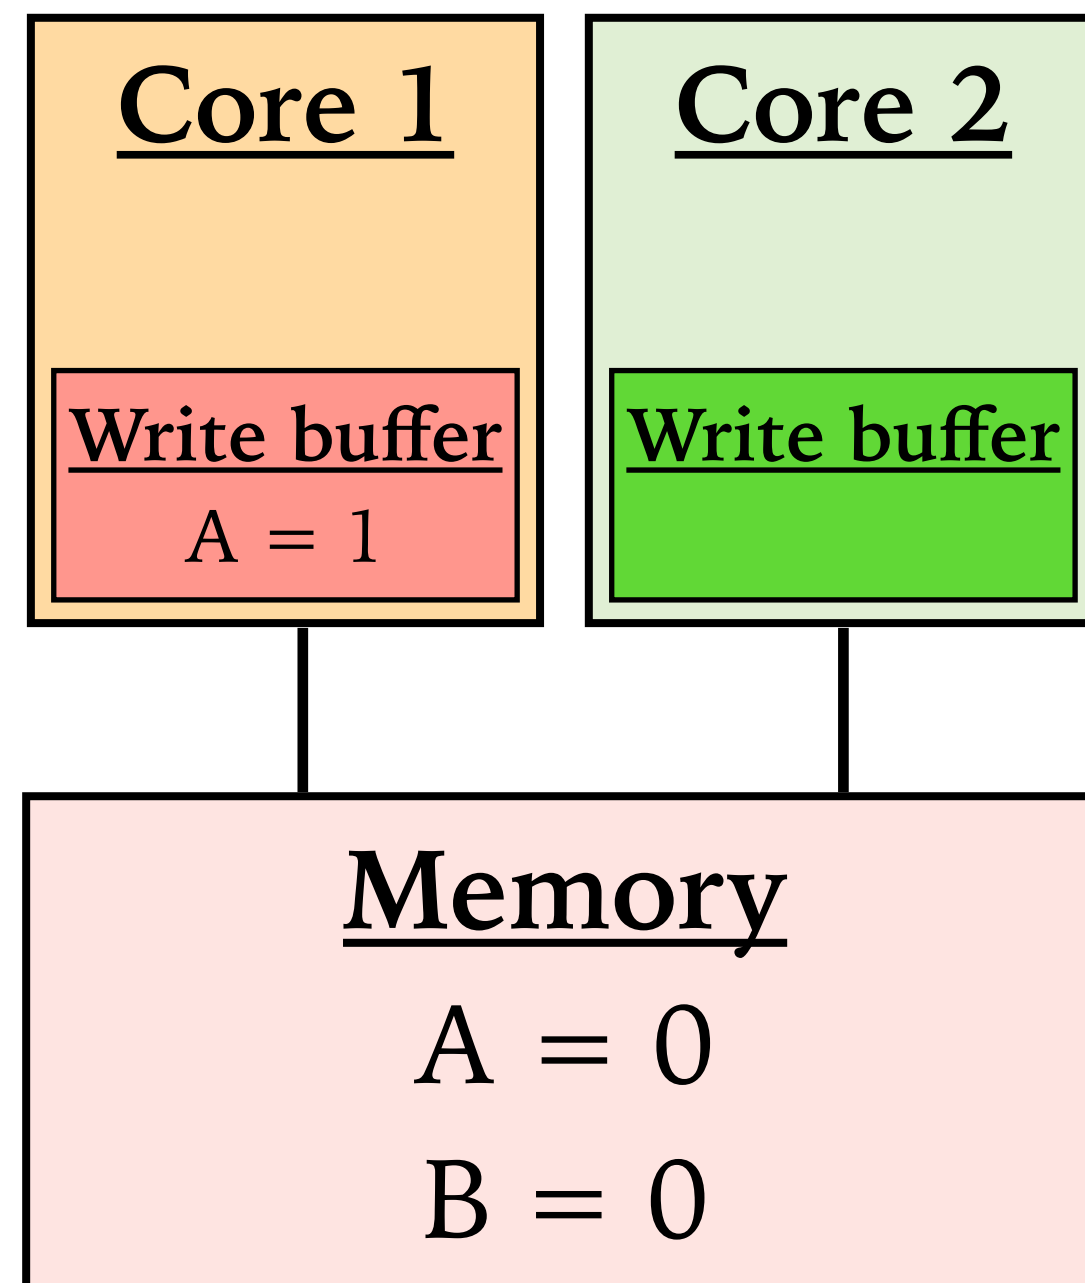
(3)  $B = 1$   
(4)  $r1 = A$

➤  $r0 = 0?$  and  $r1 = 0?$

➤ SC? No.

Executed

# Write buffers change memory behaviour



## Thread 1

(1)  $A = 1$   
(2)  $r0 = B$

## Thread 2

(3)  $B = 1$   
(4)  $r1 = A$

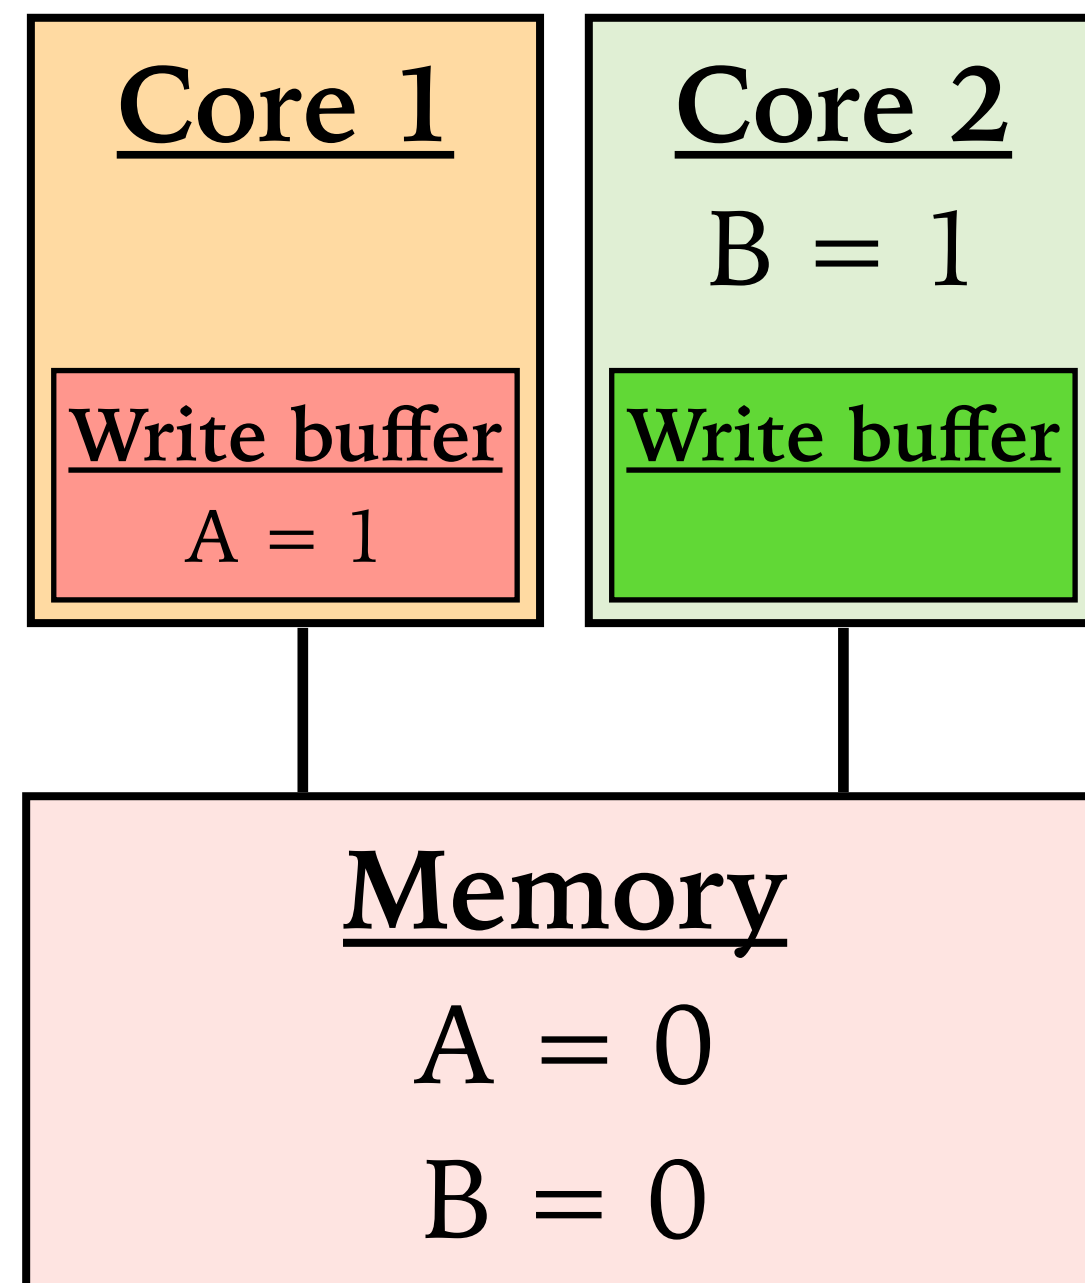
➤  $r0 = 0?$  and  $r1 = 0?$

➤ SC? No.

Executed



# Write buffers change memory behaviour



## Thread 1

(1)  $A = 1$   
(2)  $r0 = B$

## Thread 2

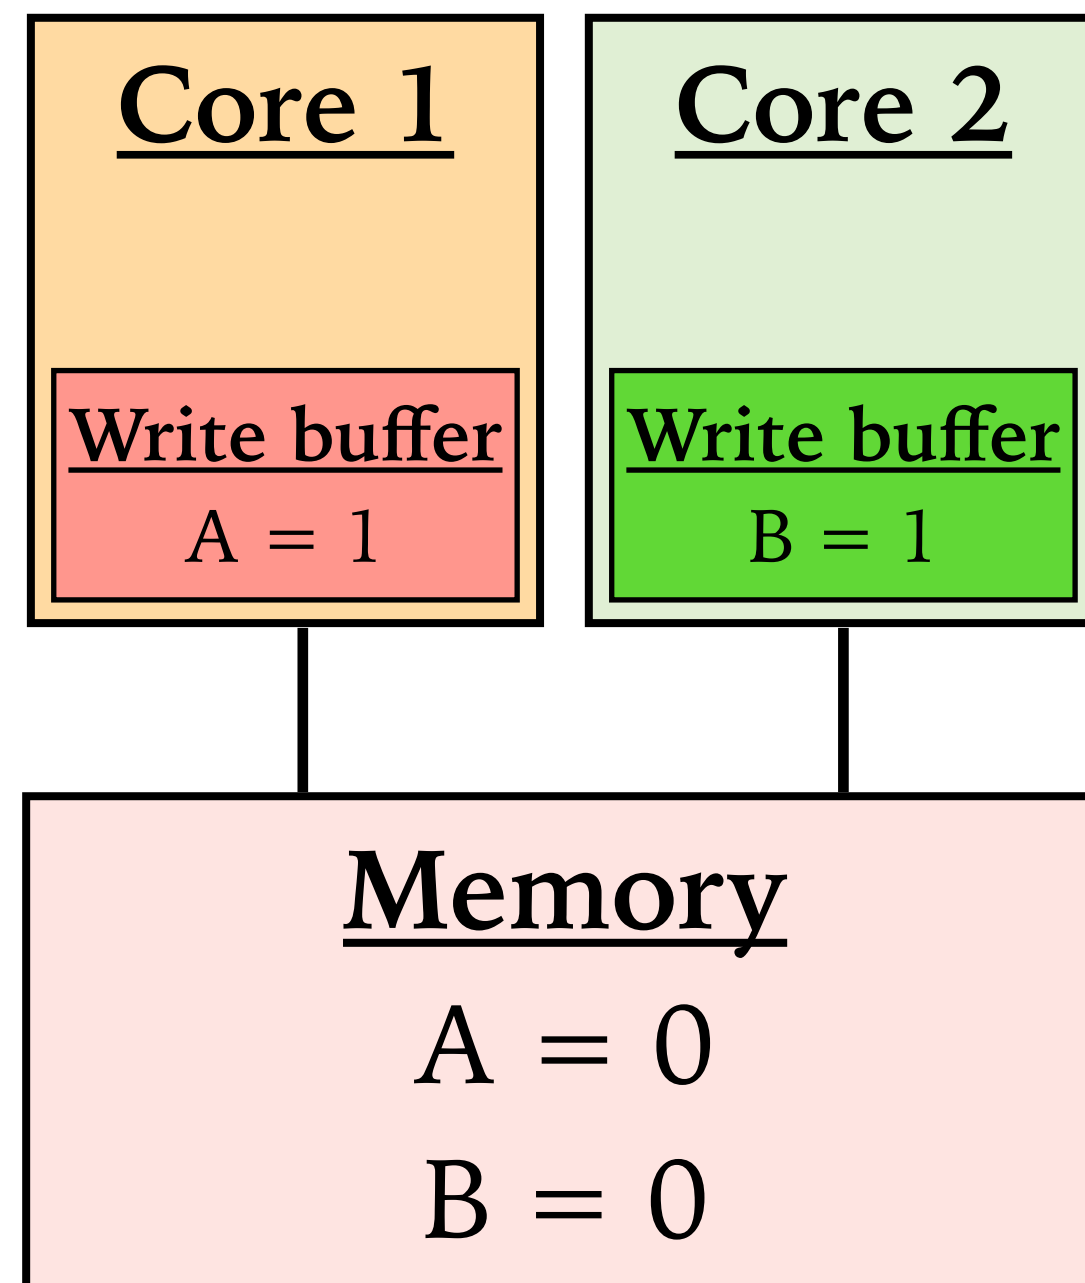
(3)  $B = 1$   
(4)  $r1 = A$

➤  $r0 = 0?$  and  $r1 = 0?$

➤ SC? No.

Executed

# Write buffers change memory behaviour



## Thread 1

(1)  $A = 1$   
(2)  $r0 = B$

## Thread 2

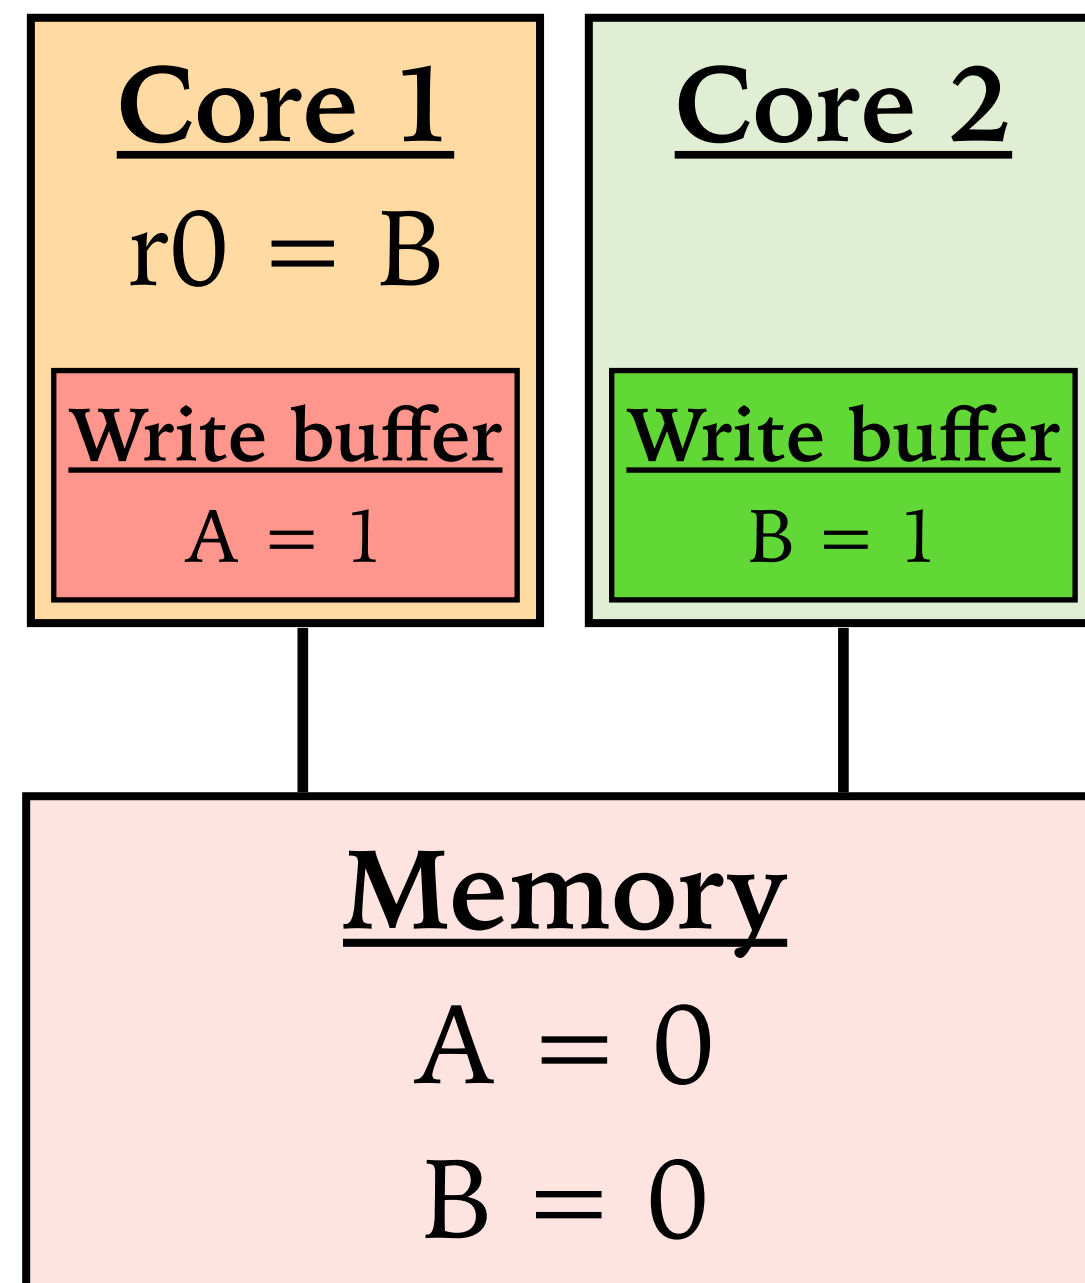
(3)  $B = 1$   
(4)  $r1 = A$

➤  $r0 = 0?$  and  $r1 = 0?$

➤ SC? No.

Executed

# Write buffers change memory behaviour



## Thread 1

(1)  $A = 1$   
(2)  $r0 = B$

## Thread 2

(3)  $B = 1$   
(4)  $r1 = A$

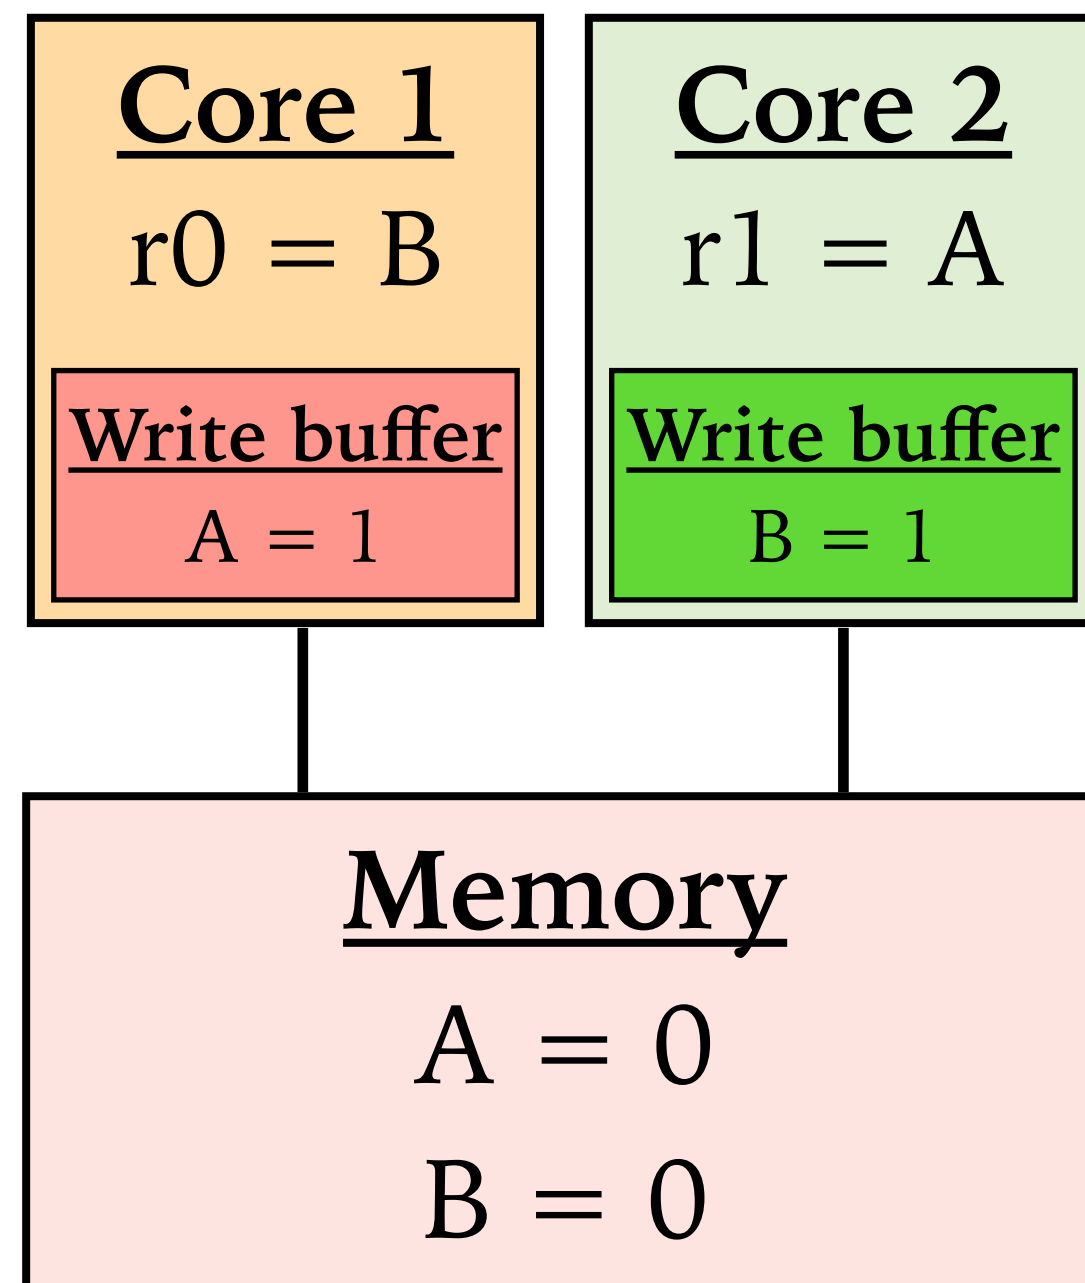
➤  $r0 = 0?$  and  $r1 = 0?$

➤ SC? No.

## Executed

$r0 = B (=0)$

# Write buffers change memory behaviour



## Thread 1

(1)  $A = 1$   
(2)  $r0 = B$

## Thread 2

(3)  $B = 1$   
(4)  $r1 = A$

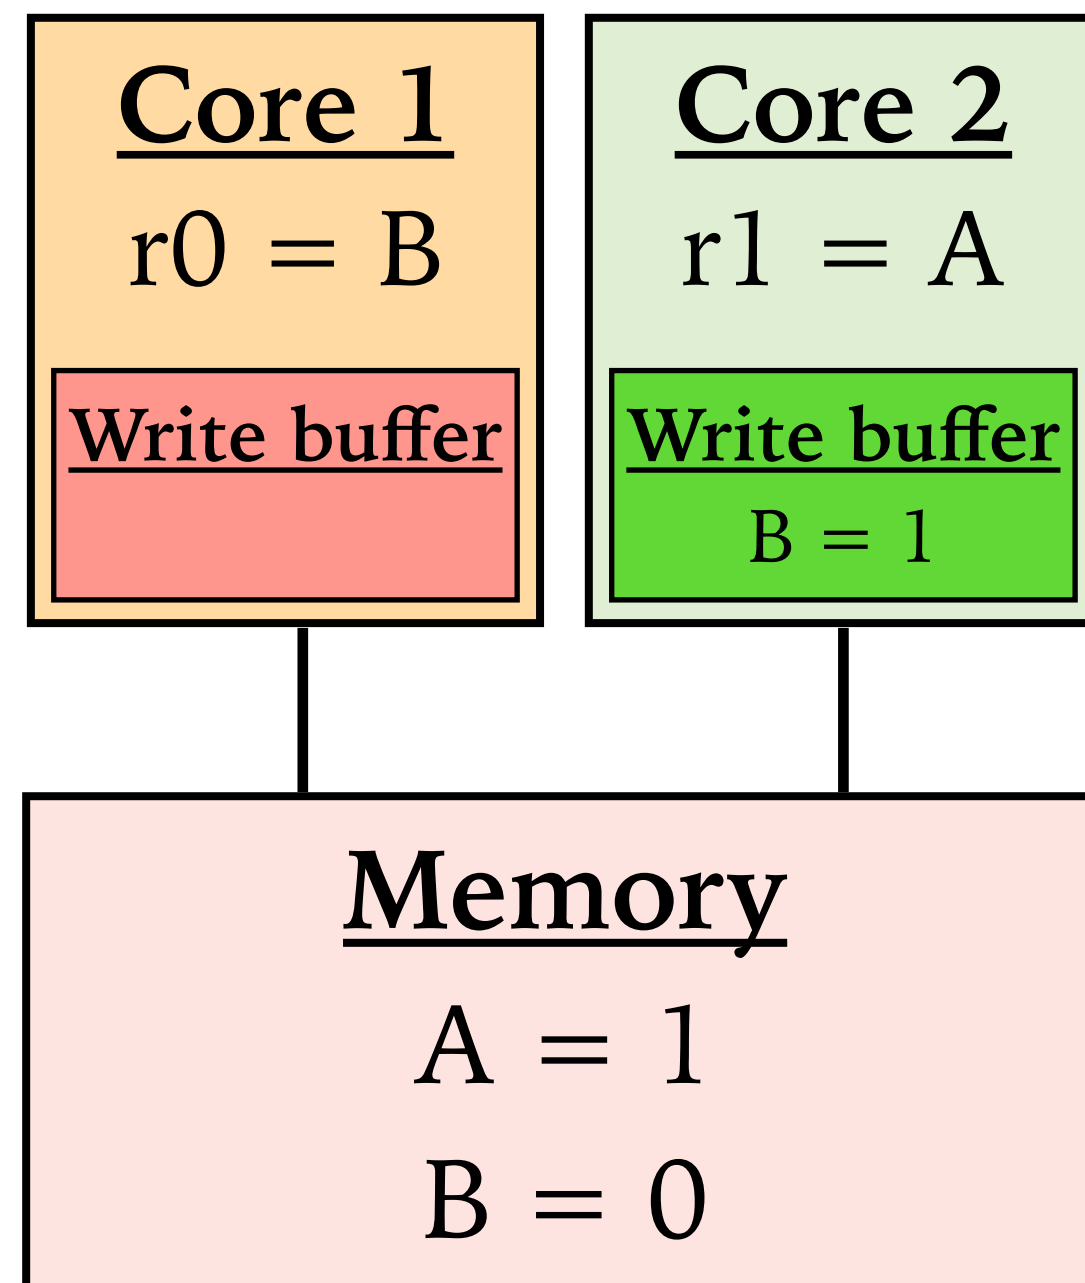
➤  $r0 = 0?$  and  $r1 = 0?$

➤ SC? No.

## Executed

$r0 = B (=0)$   
 $r1 = A (=0)$

# Write buffers change memory behaviour



## Thread 1

(1)  $A = 1$   
(2)  $r0 = B$

## Thread 2

(3)  $B = 1$   
(4)  $r1 = A$

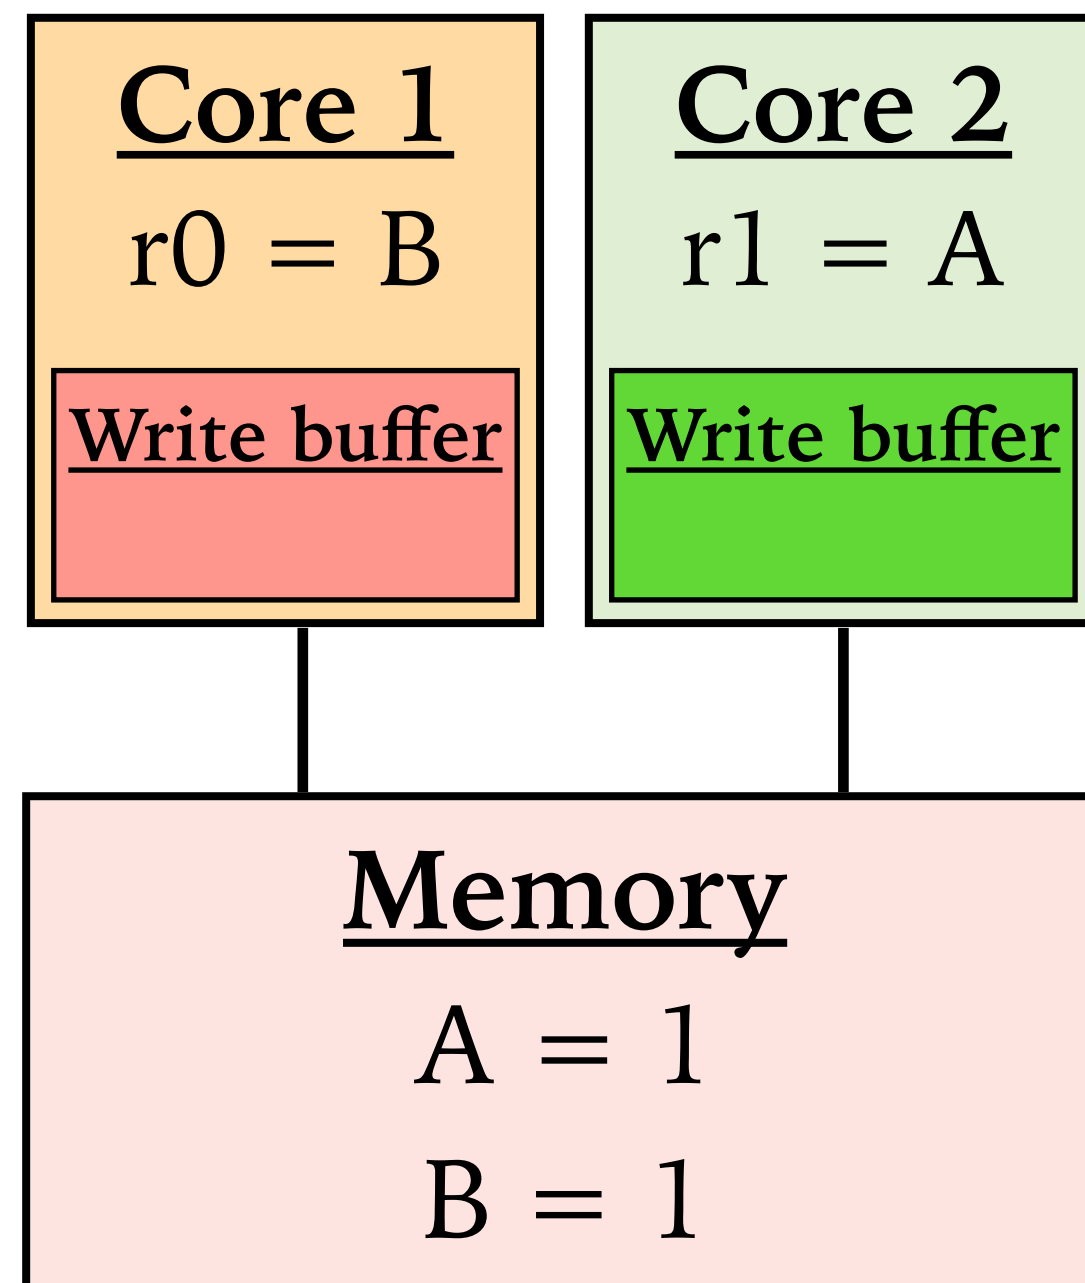
➤  $r0 = 0?$  and  $r1 = 0?$

➤ SC? No.

## Executed

$r0 = B (=0)$   
 $r1 = A (=0)$   
 $A = 1$

# Write buffers change memory behaviour



## Thread 1

(1)  $A = 1$   
(2)  $r0 = B$

## Thread 2

(3)  $B = 1$   
(4)  $r1 = A$

➤  $r0 = 0?$  and  $r1 = 0?$

➤ SC? No.

➤ Write Buffers? Yes!


## Executed

$r0 = B (=0)$   
 $r1 = A (=0)$   
 $A = 1$   
 $B = 1$

# But does anyone use write buffers?

- Every modern CPU!

- x86
- ARM
- PowerPC
- ...



- Write buffers define a *new memory model*: **Total Store Ordering.**

Performance evaluation of memory consistency models for shared-memory multiprocessors. Gharachorloo, Gupta, Hennessy. ASPLOS 1991.



# Total Store Ordering (TSO)

---

- TSO allows optimizations that can make the underlying program significantly faster.
- Essentially, SC + Write Buffers.
- But..
  - TSO allows more behaviours than SC.
  - Which means, programs are harder to reason about.
- Yet..
  - x86 specifies TSO as its memory model.





# More “relaxed” memory models

---

- Partial Store Ordering (used by SPARC)
  - Supports *write coalescing*: merge writes to the same cache line inside the write buffer to save memory bandwidth.
  - Allows writes to be reordered with other writes.
- Weak Ordering (ARM, PowerPC)
  - No guarantees about operations on data!
  - (Almost) Everything can be reordered.
- My head is spinning now!
  - How do we manage predictable outputs?
  - Most hardware provide fences, barriers, atomic instructions, and so on.



# So compilers can use fences while generating code

- But compilers too reorder instructions, heavily!!
- SC is too restrictive even for compiler optimizations.

## 1. Loop-Invariant Code Motion

### Thread 1

```
X = 0
for i = 0 to 100:
    X = 1
    print X
```

### Thread 2

```
X = 0
```

The set of valid program outputs has changed!

### Valid outputs:

```
1111111111...
1111101111...
```

### Thread 1

```
X = 1
for i = 0 to 100:
    print X
```

### Thread 2

```
X = 0
```

**So compiler can't  
perform LICM??**

### Valid outputs:

```
1111111111...
1111100000...
```



# So compilers can use fences while generating code

- And we know that compilers heavily reorder instructions!!
- SC is too restrictive even for compiler optimizations.

## 2. Common Subexpression Elimination

Thread 1  
 $A = 6$   
 $FLAG = 1$

Thread 2  
 $C = A - 1$   
 $\text{while } (FLAG == 0) \{ \}$   
 $B = A - 1$

$(A - 1)$  cannot be  
eliminated for  
assignment to B!

So compiler can't  
perform CSE??

# So compilers can use fences while generating code

- And we know that compilers heavily reorder instructions!!
- SC is too restrictive even for compiler optimizations.

## 3. Register Allocation

FLAG cannot be  
allocated a register!

Thread 1  
A = 6  
FLAG = 1

Thread 2  
while (FLAG == 0) {}  
B = A

Basically, PLs need  
memory models  
of their own.

Memory load  
in each iteration??

# Programming Language Memory Models

---


- Guess which was the first programming language to have a memory model?

- Of course:



- And which language followed suit?

- Of course:



- Java and C++ have the most extensive memory-model descriptions till date.
  - But Java has multithreading in-built, which is good.



# Java Memory Model

---

- SC is guaranteed for programs that are **data-race free (DRF)**.
  - Observe that all confusing programs that we saw had data-races!
  - **SC for DRF.**
- If the program has data-races:
  - the user is supposed to know what they were doing;
  - the compiler can reorder anything that is not synchronized.
- *p.s.* C++ memory model is quite similar, except that it has undergone revisions more number of times than Java's, because:
  - multi-threading is a library feature (different stakeholders, more confusion).



# JMM Guarantees

- Each action in a thread happens before every action in that thread that comes later in the program's order.
- An unlock on a monitor happens before every subsequent lock on that same monitor.
- A write to a volatile field happens before every subsequent read of that same volatile.
- A call to `start()` on a thread happens before any actions in the started thread.
- All actions in a thread happen before any other thread successfully returns from a `join()` on that thread.

**TLDR:** Synchronization and fork-join constructs, as well as writes to volatile fields, enforce our good-old friend: *happens-before relationship* :-)



# Key Takeaways

- A memory model defines what valid behaviours can a programmer expect, with a multi-threaded program written in a language  $L$  and executed on a machine  $M$ .
- Simpler to understand memory models (e.g. SC) may disallow impactful reordering transformations, both at compiler and hardware levels.
- Relaxed memory models (e.g. TSO) allow performing some thread-local optimizations.
- Java assures SC on DRF programs; that is, programmer has to prevent data races using synchronization constructs.
  - Observe: SC on DRF is not the same as determinism (various thread-local reorderings and cross-thread interleavings still allowed).
- But programmers are stupid (okay, sometimes)!
- Consequence: Data-race detection is a *very* active research area.

See you  
on Monday!

